

SQL Processing with SAS®

Course Notes

SQL Processing with SAS® Course Notes was developed by Jens Dahl Mikkelsen and Amy Peters. Additional contributions were made by Kathy Kiraly, Theresa Stemler, Rich Papel, and Roger Staum. Editing and Production support was provided by the Curriculum Development and Support Department.

SAS and all other SAS Institute Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc. in the USA and other countries. ® indicates USA registration.

Other brand and product names are trademarks of their respective companies.

SQL Processing with SAS® Course Notes

Copyright © 2006 by SAS Institute Inc., Cary, NC 27513, USA. All rights reserved. Printed in the United States of America. No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by any means, electronic, mechanical, photocopying, or otherwise, without the prior written permission of the publisher, SAS Institute Inc.

Book code E70256, course code SQL, prepared date 25Sep06.

SQL_003

ISBN 1-59994-251-8

Table of Contents

Course Description	vi
Prerequisites	vii
Chapter 1 Introduction	1-1
1.1 Structured Query Language	1-3
1.2 SAS Functions	1-8
1.3 Case Study	1-14
1.4 Chapter Summary	1-17
Chapter 2 Basic Queries	2-1
2.1 Overview of the SQL Procedure	2-3
2.2 Specifying Columns	2-11
2.3 Specifying Rows	2-17
2.4 Presenting Data	2-31
2.5 Summarizing Data	2-41
2.6 Subqueries	2-51
2.7 Chapter Summary	2-76
2.8 Solutions to Exercises	2-77
Chapter 3 Combining Tables	3-1
3.1 Overview	3-3
3.2 Joins	3-5
3.3 Complex Joins	3-27
3.4 Set Operators	3-46

3.5	Chapter Summary	3-77
3.6	Solutions to Exercises	3-79
Chapter 4	Creating and Modifying Tables and Views	4-1
4.1	Creating Tables	4-3
4.2	Creating Views	4-20
4.3	Creating Indexes	4-31
4.4	Maintaining Tables	4-45
4.5	Chapter Summary	4-60
4.6	Solutions to Exercises	4-63
Chapter 5	Additional SQL Features	5-1
5.1	Setting SQL Procedure Options	5-3
5.2	Dictionary Tables and Views	5-9
5.3	Interfacing PROC SQL with Macro Language (Optional)	5-18
5.4	Program Testing and Performance	5-29
5.5	Chapter Summary	5-34
Appendix A	Overview of Table and Column Names	A-1
A.1	Table and Column Names Sorted by Column Names	A-3
Appendix B	Overview of Table and Column Names	B-1
B.1	Table and Column Names Sorted by Table Name	B-3
Appendix C	Table Listings	C-1
C.1	Partial Table Listings	C-3

Appendix D	Index	D-1
-------------------	--------------------	------------

Course Description

This intermediate course focuses on using SQL as a data query and manipulation tool. You learn to use the SQL procedure as a data retrieval tool within SAS programs. Specifically, you learn how to perform queries on data; retrieve data from multiple tables; create views, indexes, and tables; and update or delete values in existing tables and views. Using features of the SQL procedure to debug, test, and optimize the performance of SQL queries is also discussed.

To learn more...



SAS Education

A full curriculum of general and statistical instructor-based training is available at any of the Institute's training facilities. Institute instructors can also provide on-site training.

For information on other courses in the curriculum, contact the SAS Education Division at 1-919-531-7321, or send e-mail to training@sas.com. You can also find this information on the Web at support.sas.com/training/ as well as in the Training Course Catalog.



SAS Publishing

For a list of other SAS books that relate to the topics covered in this Course Notes, USA customers can contact our SAS Publishing Department at 1-800-727-3228 or send e-mail to sasbook@sas.com. Customers outside the USA, please contact your local SAS office.

Also, see the Publications Catalog on the Web at support.sas.com/pubs for a complete list of books and a convenient order form.

Prerequisites

Before attending this course, you should be able to

- submit SAS programs on your operating system
- create and access SAS data sets
- use arithmetic, comparison, and logical operators
- invoke SAS procedures.

You can gain this experience from the *SAS Programming I: Essentials* course. No knowledge of SQL is necessary.

Chapter 1 Introduction

1.1	Structured Query Language	1-3
1.2	SAS Functions	1-8
1.3	Case Study	1-14
1.4	Chapter Summary.....	1-17

1.1 Structured Query Language

Objectives

- Understand the background and applications of Structured Query Language.

3

Structured Query Language

Structured Query Language (SQL)

- is a standardized language that is widely used to retrieve and update data in tables and in views based on those tables
- was originally designed as a query tool for relational databases, but it is now used by many software products.

4

Structured Query Language: Timeline

- 1970 Conceptualized and proposed by Dr. E. F. Codd at the IBM Research Laboratory, San Jose, CA
- 1970-1980 Developed by IBM
- 1981 First commercial SQL-based product, the IBM SQL/DS System
- 1989 Over 75 SQL database management systems exist, including SAS Release 6.06.

5

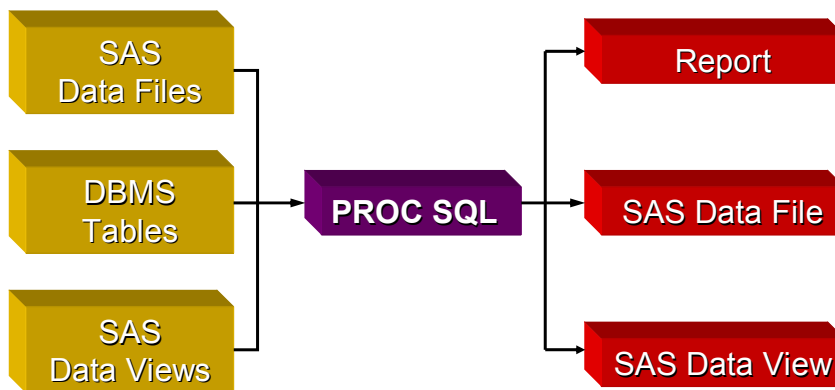
The SQL Procedure

The SQL procedure uses SQL to do the following:

- query SAS data sets
- generate reports from SAS data sets
- combine SAS data sets in many ways
- create and delete SAS data files, views, and indexes
- update existing SAS data sets

6

Structured Query Language



7

More About the SQL Procedure

The SQL procedure

- enables you to use SQL within SAS
- follows the guidelines set by the American National Standards Institute (ANSI)
- includes enhancements for compatibility with SAS software
- is part of Base SAS software
- can replace the need for multiple DATA and PROC steps with one query.

8

The SQL Procedure

IS NOT

- a replacement for the DATA step
- a custom reporting tool.

IS

- a tool for queries
- for data manipulation
- an augmentation to the DATA step.

9

SAS Data Sets

A SAS data set can be a:

- SAS data file that stores data descriptions and data values together
- PROC SQL view that stores a PROC SQL query that retrieves data stored in other files
- DATA step view that stores a DATA step that retrieves data stored in other files
- SAS/ACCESS view that stores information required to retrieve data stored in a DBMS

10

Terminology

Data Processing	SAS	SQL
File	SAS Data Set	Table
Record	Observation	Row
Field	Variable	Column

1.2 SAS Functions

Objectives

- Review the concatenation operator and various SAS functions.

13

SAS Numeric Functions

Example: Calculate age and create two new variables.

```
data new;  
  date=today();  
  birth='01jun1970'd;  
  bmonth=month(birth);  
  fullage=(date-birth)/365.25;  
  age=int(fullage);  
run;  
proc print data=new noobs;  
  format date birth date9.;  
run;
```

- The **TODAY** function returns today's date as a SAS date value.
- The **MONTH** function returns the month portion of a SAS date as an integer, 1-12.
- The **INT** function returns the integer portion of a numeric value.

14

c1s2d01

SAS Numeric Functions

Output

The SAS System				
date	birth	bmonth	fullage	age
12SEP2006	01JUN1970	6	36.2820	36

15



The results above were generated on September 22, 2006. Your values may differ.

SAS Character Functions

Use the SUBSTR function to extract individual characters from a character value.

General form of the SUBSTR function:

SUBSTR(*argument*,*position*<,*length*>)

argument can be a character constant, variable, or expression.

position specifies the starting position.

length specifies the number of characters to extract. If omitted, the substring consists of the remainder of *argument*.

16

SAS Character Functions

Example: Create a full name from a character string.

```
data new;
  name='Gomez, Gabriela  ';
  fname1=substr(name,8) || ' ' || substr(name,1,5);
  fname2=trim(substr(name,8)) || ' ' || substr(name,1,5);
run;
proc print data=new noobs;
run;
```

- **SUBSTR** extracts part of a character value. The length of the variable is determined by the first argument.
- **TRIM** removes trailing blanks from a value. The length of the variable is determined by the argument.
- **||** concatenates character values.

17

c1s2d02

SAS Character Functions

Output

The SAS System		
name	fname1	fname2
Gomez, Gabriela	Gabriela Gomez	Gabriela Gomez

18

SAS Character Functions

Use the SCAN function to extract the n^{th} word of a character value.

General form of the SCAN function:

SCAN(*argument*,*n*<,*delimiters*>)

- argument* can be a character constant, variable, or expression.
- n* specifies the n^{th} word to extract from the argument.
- delimiters* defines characters that delimit (separate) words.

19



If the third argument is omitted, the default delimiters are

ASCII (PC, UNIX)	blank . < (+ & ! \$ *) ; - / , % ^
EBCDIC (z/OS)	blank . < (+ & ! \$ *) ; - / , % ¢ ¬

SAS Character Functions

Example: Create a full name from a character string that contains a reversed name.

```
data new;
  name='Gomez, Gabriela ';
  first=scan(name,2,',');
  last=scan(name,1,',');
  fname='Ms.'||trim(first)||' '||last;
proc print data=new noobs;
run;
```

SCAN returns a specific word from a character string. The default length of the variable is 200.

20

c1s2d03

The above program creates FIRST and LAST columns using 200 bytes, which is the default length returned from the SCAN function. To override this default, use a LENGTH statement, as shown below:

```
data new;
  length first last $20;
  name='Gomez, Gabriela ';
  first=scan(name,2,',');
  last=scan(name,1,',');
  fname='Ms.'||trim(first)||' '||last;
run;
```

If the second argument of the SCAN function is a negative number, the function counts from the right side of the text string instead of the left.

SAS Character Functions

Output

The SAS System			
name	first	last	fname
Gomez, Gabriela	Gabriela	Gomez	Ms. Gabriela Gomez

If the length of the variables `first` and `last` is 200, where are the spaces?

1.3 Case Study

Objectives

- Explore the tables used in this course.

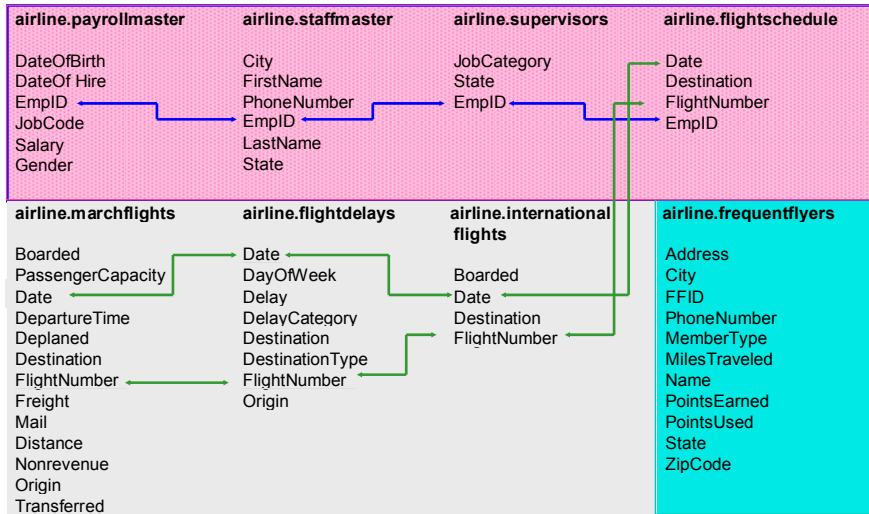
23

The examples and exercises in this course are based on the files of a fictitious airline company. These files include information on

- airline employees
- flights
- frequent flyers.

Airline Employee and Flight Tables

=airline employees
 =flights
 =frequent flyers



24



Not all tables used in this class are shown above.

Airline Destination Codes and Descriptions

Code	Description
CPH	Copenhagen
DFW	Dallas/Ft. Worth
FRA	Frankfurt
LAX	Los Angeles
LGA	New York
LHR	London
ORD	Chicago
CDG	Paris
WAS	Washington
YYZ	Toronto

25

Airline Job Codes and Descriptions

Code	Description
BCK	Baggage Check
FA	Flight Attendant
ME	Mechanic
NA	Navigator
PT	Pilot
SCP	Skycap
TA	Ticket Agent

1.4 Chapter Summary

Structured Query Language (SQL) is a standardized language that is widely used to retrieve and update data in tables and views based on those tables. The SQL procedure enables you to use SQL within SAS. You can use the SQL procedure to accomplish tasks such as querying SAS data sets, generating reports from SAS data sets, and combining SAS data sets.

The SQL procedure supports most of the functions available in the DATA step for data creation and manipulation. There are numeric functions to manipulate data values and character functions to manipulate character strings.

General form of the TODAY function:

TODAY()

General form of the MONTH function:

MONTH(*SAS date value*)

General form of the INT function:

INT(*numeric value*)

General form of the SUBSTR function:

SUBSTR(*argument,position<,length>*)

General form of the TRIM function:

TRIM(*argument*)

General form of the SCAN function:

SCAN(*argument,n<,delimiters>*)

Chapter 2 Basic Queries

2.1	Overview of the SQL Procedure.....	2-3
2.2	Specifying Columns	2-11
2.3	Specifying Rows	2-17
2.4	Presenting Data	2-31
2.5	Summarizing Data	2-41
2.6	Subqueries	2-51
2.7	Chapter Summary.....	2-76
2.8	Solutions to Exercises	2-77

2.1 Overview of the SQL Procedure

Objectives

- Understand SQL procedure syntax.

3

Features of PROC SQL

- The PROC SQL statement does not need to be repeated with each query.
- Each statement is processed individually.
- No PROC PRINT step is needed to view query results.
- No PROC SORT step is needed to order query results.
- No RUN statement is needed.
- Use a QUIT statement to terminate PROC SQL.

4

SQL is a modular language because queries (or statements) are composed of smaller building blocks (or clauses).

The SELECT Statement

A SELECT statement is used to query one or more SAS data sets.

```
proc sql;  
  select EmpID, JobCode, Salary  
    from airline.payrollmaster  
   where JobCode contains 'NA'  
   order by Salary desc;
```

5

c2s1d01



Use a comma to separate items in a list, such as column or table names. Place a single semicolon at the end of the last clause.

Features of the SELECT Statement

The features of the SELECT statement include the following:

- selects data that meets certain conditions
- groups data
- specifies an order for the data
- formats the data
- queries 1 to 32 tables

6

Table names can be 1 to 32 characters in length and are not case-sensitive.

Variable names can be 1 to 32 characters in length and are stored in mixed case but are normalized for lookups and comparisons. However, the first usage of the variable determines the capitalization pattern.

Librefs and filerefs are limited to 8 characters. Starting in SAS[®]9, format and informat names can be up to 32 characters in length.

SELECT Statement Syntax

General form of the SELECT statement:

```
SELECT column-1<, column-2>...
FROM table-1|view-1<, table-2|view-2>...
<WHERE expression>
<GROUP BY column-1<, column-2>...>
<HAVING expression>
<ORDER BY column-1<, column-2>... <DESC>>;
```

7

- SELECT** specifies the columns to be selected.
- FROM** specifies the table to be queried.
- WHERE** subsets the data based on a condition.
- GROUP BY** classifies the data into groups.
- HAVING** subsets groups of data based on a group condition.
- ORDER BY** sorts rows by the values of specific columns. By default, results are sorted in ascending order. Use the DESC keyword to sort in descending order.



The order of the above clauses within the SELECT statement **does** matter.

table is a SAS data set (data file or data view).

column is a column name, expression, or summary function.

The VALIDATE Keyword

Partial SAS Log

```
proc sql;  
  validate  
    select EmpID, JobCode, Salary  
      from airline.payrollmaster  
     where JobCode contains 'NA'  
     order by Salary desc;  
NOTE: PROC SQL statement has valid syntax.
```

8

c2s1d02

Features of the VALIDATE Keyword

The features of the VALIDATE keyword include the following:

- is used only in a SELECT statement
- tests the syntax of a query without executing the query
- checks column name validity
- prints error messages for invalid queries

9

The NOEXEC Option

The NOEXEC option can also be used for syntax checking.

Partial SAS Log

```
proc sql noexec;  
  select EmpID,JobCode,Salary  
    from airline.payrollmaster  
   where JobCode contains 'NA'  
   order by Salary desc;  
NOTE:Statement not executed due to NOEXEC option.
```

Additional PROC SQL Statements

PROC SQL supports many statements in addition to the SELECT statement.

PROC SQL *<option <option>...>*;

Chapter 4

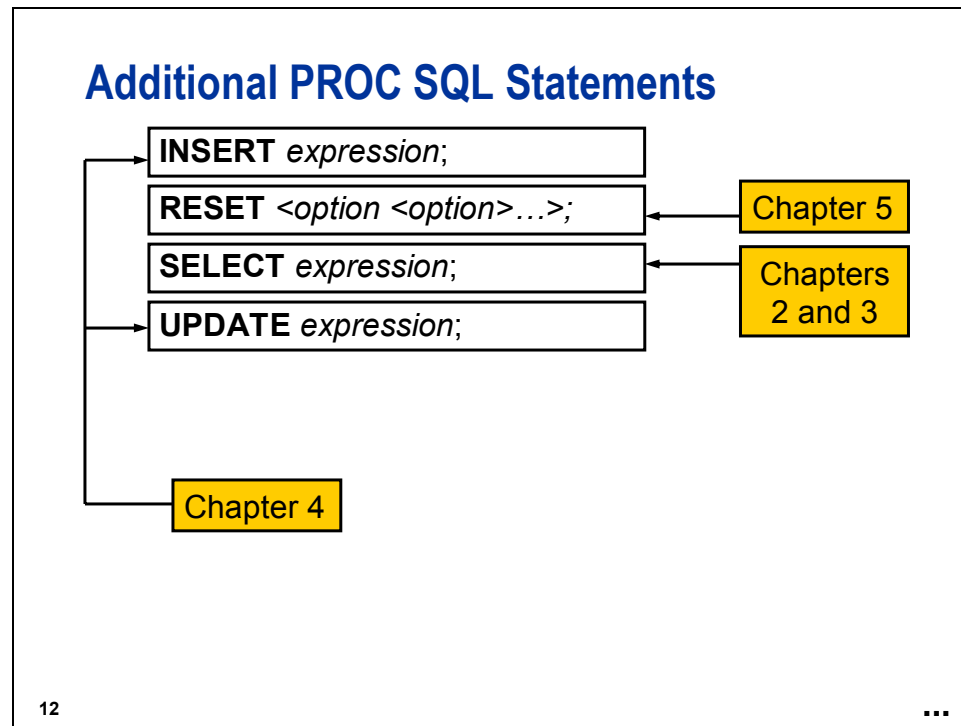
ALTER *expression*;
CREATE *expression*;
DELETE *expression*;
DESCRIBE *expression*;
DROP *expression*;

continued...

11

...

ALTER	adds, drops, and modifies columns in a table.
CREATE	builds new tables.
DELETE	eliminates unwanted rows from a table or view.
DESCRIBE	displays table attributes or view definitions.
DROP	eliminates entire tables, views, or indexes.



INSERT adds rows of data to tables.

RESET adds to or changes PROC SQL options without re-invoking the procedure.

SELECT specifies columns to be printed.

UPDATE modifies data values in existing rows of a table or view.



The NOEXEC option checks for invalid syntax in all the statements previously mentioned, but the VALIDATE option applies only to the SELECT statement.

2.2 Specifying Columns

Objectives

- Display columns directly from a table.
- Display columns calculated from other columns in a query.

Retrieving Data from a Table

If you are familiar with a table, you can specify column names to be printed in the SELECT statement.

Example: Print employee IDs, job codes, and salaries.

```
proc sql;  
    select EmpID, JobCode, Salary  
    from airline.payrollmaster;
```

15

c2s2d01



Starting with SAS®9, you can reference tables using their physical filename:

```
proc sql;  
    select *  
    from 'c:\workshop\winsas\sql\payrollmaster.sas7bdat';  
quit;
```

Employee IDs, Job Codes, and Salaries

Partial Output

The SAS System		
Emp ID	Job Code	Salary
1919	TA2	\$48,126
1653	ME2	\$49,151
1400	ME1	\$41,677
1350	FA3	\$46,040
1401	TA3	\$54,351
1499	ME3	\$60,235
1101	SCP	\$26,212
1333	PT2	\$124,048
1402	TA2	\$45,661

16

Retrieving Data from a Table

```
proc sql;
  select *
  from airline.payrollmaster;
```

If you are not familiar with a table, an asterisk in the SELECT statement prints all columns in their originally stored order.

Partial Output

The SAS System					
Emp ID	Gender	Job Code	Salary	DateOfBirth	DateOfHire
1919	M	TA2	\$48,126	16SEP1958	07JUN1985
1653	F	ME2	\$49,151	19OCT1962	12AUG1988
1400	M	ME1	\$41,677	08NOV1965	19OCT1988
1350	F	FA3	\$46,040	04SEP1963	01AUG1988
1401	M	TA3	\$54,351	16DEC1948	21NOV1983
1499	M	ME3	\$60,235	29APR1952	11JUN1978

17

c2s2d02

The FEEDBACK Option

Use the FEEDBACK option to write the expanded SELECT statement to the SAS log.

Partial SAS Log

```
proc sql feedback;
  select *
    from airline.payrollmaster;
NOTE: Statement transforms to
      select PAYROLLMASTER.EmpID,
             PAYROLLMASTER.Gender, PAYROLLMASTER.JobCode,
             PAYROLLMASTER.Salary,
             PAYROLLMASTER.DateOfBirth,
             PAYROLLMASTER.DateOfHire
      from AIRLINE.PAYROLLMASTER;
```

18

c2s2d03



This option expands any use of an asterisk into the list of qualified columns it represents. NOFEEDBACK is the default.

Expressions

Calculate new columns from existing columns, and name the new columns using the AS keyword.

Example: Calculate employee bonuses.

```
proc sql;
  select EmpID, JobCode, Salary,
         Salary * .10 as Bonus
  from airline.payrollmaster;
```

19

c2s2d04

The new column is called an *alias*. The AS keyword is required. Omission of the alias causes the column heading to be blank.

Employee Bonuses

Partial Output

The SAS System			
Emp ID	Job Code	Salary	Bonus
1919	TA2	\$48,126	4812.64
1653	ME2	\$49,151	4915.12
1400	ME1	\$41,677	4167.66
1350	FA3	\$46,040	4604.04
1401	TA3	\$54,351	5435.08
1499	ME3	\$60,235	6023.5
1101	SCP	\$26,212	2621.22
1333	PT2	\$124,048	12404.84
1402	TA2	\$45,661	4566.1
1479	TA3	\$54,299	5429.9
1403	ME1	\$39,301	3930.08
1739	PT1	\$93,124	9312.38
1658	SCP	\$25,120	2512.02

20

Expressions

Use SAS DATA step functions for calculating columns.

Example: Calculate the age of each employee.

```
proc sql;
  select EmpID, JobCode,
         int((today() - DateOfBirth) / 365.25)
         as Age
  from airline.payrollmaster;
```

21

c2s2d05

Employee Ages

Partial Output

The SAS System		
Emp ID	Job Code	Age
1919	TA2	47
1653	ME2	43
1400	ME1	40
1350	FA3	43
1401	TA3	57

22



All SAS DATA step functions are supported except LAG and DIF.

2.3 Specifying Rows

Objectives

- Eliminate duplicate rows in a query.
- Subset the data displayed in a query.

24

Specifying All Rows in a Table

By default, all rows in a table are returned in a query.

Example: Display all rows and columns of the **airline.internationalflights** table.

```
proc sql;  
  select *  
  from airline.internationalflights;
```

25

c2s3d01

All Rows in a Table

Partial Output

The SAS System			
FlightNumber	Date	Destination	Boarded
182	01MAR2000	YYZ	104
219	01MAR2000	LHR	198
387	01MAR2000	CPH	152
622	01MAR2000	FRA	207
821	01MAR2000	LHR	205
132	01MAR2000	YYZ	115
271	01MAR2000	CDG	138
182	02MAR2000	YYZ	116
219	02MAR2000	LHR	147
387	02MAR2000	CPH	105
622	02MAR2000	FRA	176
821	02MAR2000	LHR	201
132	02MAR2000	YYZ	106

26

Eliminating Duplicate Rows

Use the DISTINCT keyword to eliminate duplicate rows in query results.

Example: Determine the international flights that were flown during the month.

```
proc sql;
  select distinct FlightNumber,
                 Destination
  from airline.internationalflights;
```

27

c2s3d02



The DISTINCT keyword applies to all columns in the SELECT list. One row is displayed for each existing combination of values.

Eliminating Duplicate Rows

Output

The SAS System	
FlightNumber	Destination
132	YYZ
182	YYZ
219	LHR
271	CDG
387	CPH
622	FRA
821	LHR

28

Subsetting with the WHERE Clause

Use a WHERE clause to specify a condition that the data must satisfy before being selected.

Example: Display all employees that earn more than \$112,000.

```
proc sql;  
  select EmpID, JobCode, Salary  
    from airline.payrollmaster  
   where Salary > 112000;
```

29

c2s3d03

Subsetting with the WHERE Clause

Output

The SAS System		
Emp ID	Job Code	Salary
1333	PT2	\$124,048
1404	PT2	\$127,926
1118	PT3	\$155,931
1410	PT2	\$118,559
1777	PT3	\$153,482
1106	PT2	\$125,485
1442	PT2	\$118,350
1478	PT2	\$117,884
1890	PT2	\$120,254
1107	PT2	\$125,968
1830	PT2	\$118,259
1928	PT2	\$125,801

30

Subsetting with the WHERE Clause

You can use all common comparison operators in a WHERE clause.

Mnemonic	Symbol	Definition
LT	<	Less than
GT	>	Greater than
EQ	=	Equal to
LE	<=	Less than or equal to
GE	>=	Greater than or equal to
NE	≠	Not equal to (EBCDIC)
	^=	Not equal to (ASCII)

31

Subsetting with the WHERE Clause

You can use the IN operator to compare a value to a list of values. If the value matches at least one in the list, the expression is true; otherwise, the expression is false.

```
where JobCategory in ('PT','NA','FA')
```

```
where DayOfWeek in (2,4,6)
```

32

Subsetting with the WHERE Clause

You can specify multiple expressions in a WHERE clause by using logical operators.

Mnemonic	Symbol	Definition
OR		or, either
AND	&	and, both
NOT	¬	not, negation EBCDIC
NOT	^	not, negation ASCII

33

Subsetting with the WHERE Clause

Use either **CONTAINS** or **?** to select rows that include the substring specified.

```
where word ? 'LAM'
```

(BLAME, LAMENT, and BEDLAM are selected.)

Use either **IS NULL** or **IS MISSING** to select rows with missing values.

```
where FlightNumber is missing
```

34

Alternative statements are

- `where FlightNumber = ' '`
- `where FlightNumber = .`

With the `=` operator, you must know whether **FlightNumber** is character or numeric. However, if you use **IS MISSING**, you do not need advance knowledge of the column type.

Subsetting with the WHERE Clause

Use **BETWEEN-AND** to select rows containing ranges of values, inclusively.

```
where Date between '01mar2000'd  
and '07mar2000'd
```

```
where Salary between 70000 and 80000
```

35

Subsetting with the WHERE Clause

Use **LIKE** to select rows by comparing character values to specified patterns.

A % sign replaces any number of characters.

```
where LastName like 'H%'
```

(H plus any characters; for example, HENDRY, HANSON, and HALL are selected.)

36

Subsetting with the WHERE Clause

A single underscore ('_') replaces individual characters.

```
where JobCode like '___1'
```

captures any two characters and 1, for example, 'FA1', 'TA1', 'NA1'.

Two underscores,
followed by a 1

37

Subsetting with the WHERE Clause

Select the jobcodes that contain an underscore (_), followed by a number.

```
proc sql;
  select EmpID, Jobcode
  from airline.payrollmaster2
  where jobcode like 'FA_%';
quit;
```

Partial Output
(rows 7-13)

Emp ID	Jobcode
5098	FA_2
1124	FA_1
1422	FA_1
1094	FA_1
6345	FAN2
7109	FAN2
1113	FA_1

c2s3d04

38

ESCAPE Clause

The **ESCAPE** clause in the **LIKE** condition enables you to designate a single character string literal, known as an *escape character*, to indicate how PROC SQL should interpret the LIKE wildcards, percent (%) and underscore (_), if they are used within a character string.

39

ESCAPE Clause

```
proc sql;
  select EmpID, Jobcode
    from airline.payrollmaster2
   where jobcode like 'FA/_%' ESCAPE '/';
quit;
```

Partial Output
(rows 7-15)

Emp ID	Jobcode
5098	FA_2
1124	FA_1
1422	FA_1
1094	FA_1
1113	FA_1
8322	FA_2
1103	FA_1
1477	FA_2
1115	FA_3

c2s3d05

40

Subsetting with the WHERE Clause

The sounds-like (=*) operator selects rows containing a spelling variation of the specified word(s).

```
where LastName =* 'SMITH'
```

selects values SMITT, SMYTHE, and SMOTHE, in addition to SMITH.

41

Subsetting with Calculated Values

Because a WHERE clause is evaluated first, columns used in the WHERE clause must exist in the table or be derived from existing columns.

42

Subsetting with Calculated Values

Example: Display only the flights where the total number of passengers was fewer than 100 people.

```
proc sql;  
    select FlightNumber, Date, Destination,  
           Boarded + Transferred + Nonrevenue  
           as Total  
    from airline.marchflights  
   where Total < 100;
```

43

c2s3d06

Subsetting with Calculated Values

Partial Log

```
ERROR: The following columns were not  
found in the contributing tables: Total.
```

44

Subsetting with Calculated Values

One solution is to repeat the calculation in the WHERE clause.

```
proc sql;  
  select FlightNumber, Date, Destination,  
         Boarded+Transferred+Nonrevenue  
         as Total  
  from airline.marchflights  
 where Boarded+Transferred+Nonrevenue < 100;
```

45

c2s3d07

Subsetting with Calculated Values

A more efficient method is to use the CALCULATED keyword to refer to already calculated columns in the SELECT clause.

```
proc sql;  
  select FlightNumber, Date, Destination,  
         Boarded + Transferred + Nonrevenue  
         as Total  
  from airline.marchflights  
 where calculated Total < 100;
```

46

c2s3d08

Subsetting with Calculated Values

Partial Output

The SAS System			
FlightNumber	Date	Destination	Total
982	01MAR2000	DFW	70
416	01MAR2000	WAS	93
829	01MAR2000	WAS	96
416	02MAR2000	WAS	90
302	02MAR2000	WAS	93

47

Subsetting with Calculated Values

You can also use the CALCULATED keyword in other parts of a query, for example, in a SELECT clause.

```
proc sql;
  select FlightNumber, Date, Destination,
         Boarded + Transferred + Nonrevenue
         as Total,
         calculated Total/2 as half
  from airline.marchflights;
```

48

c2s3d09

Subsetting with Calculated Values

Partial Output

The SAS System				
FlightNumber	Date	Destination	Total	half
182	01MAR2000	YYZ	123	61.5
114	01MAR2000	LAX	196	98
202	01MAR2000	ORD	167	83.5
219	01MAR2000	LHR	222	111
439	01MAR2000	LAX	185	92.5
387	01MAR2000	CPH	163	81.5
290	01MAR2000	WAS	119	59.5

2.4 Presenting Data

Objectives

- Order the data displayed in a query.
- Use SAS formats, labels, and titles to enhance query output.

51

Ordering Data

Use the ORDER BY clause to sort query results in

- ascending order (the default)
- descending order by following the column name with the DESC keyword.

52

Ordering Data

```
proc sql;  
  select EmpID, JobCode, Salary  
    from airline.payrollmaster  
   where JobCode contains 'NA'  
   order by Salary desc;
```

You can specify the collating sequence by using the SORTSEQ= option in the PROC SQL statement. Use this option only if you want a collating sequence other than your system's or installation's default collating sequence. For additional information, see SORTSEQ under the SORT procedure in your online or written documentation.

PROC SQL uses information provided by a table's internal sort indicator (if available) to avoid performing unnecessary sorts.

Ordering Data

Output

Emp ID	Job Code	Salary
1352	NA2	\$75,317
1417	NA2	\$73,178
1935	NA2	\$71,513
1839	NA1	\$60,806
1443	NA1	\$59,184
1332	NA1	\$59,049
1269	NA1	\$58,366
1111	NA1	\$56,820

Notice the
descending
order of Salary.

54

Ordering Data

In an ORDER BY clause, you order query results by specifying the following:

- any column or expression (display or nondisplay)
- a column name or a number that represents the position of an item in the SELECT list
- multiple columns

55

Ordering Data

Example: Display the London flights in date order by descending total number of passengers.

```
proc sql;
  select FlightNumber, Date,
         Origin, Destination,
         Boarded+Transferred+Nonrevenue
  from airline.marchflights
  where Destination='LHR'
  order by Date,5 desc;
```

Mix and match

56

c2s4d02

Ordering Data

Partial Output

The SAS System				
FlightNumber	Date	Origin	Destination	
219	01MAR2000	LGA	LHR	222
821	01MAR2000	LGA	LHR	222
821	02MAR2000	LGA	LHR	210
219	02MAR2000	LGA	LHR	172
219	03MAR2000	LGA	LHR	211
821	03MAR2000	LGA	LHR	172
219	04MAR2000	LGA	LHR	250
821	04MAR2000	LGA	LHR	182
219	05MAR2000	LGA	LHR	167
821	06MAR2000	LGA	LHR	185
219	06MAR2000	LGA	LHR	183

57

Enhancing Query Output

You can use SAS formats and labels to customize PROC SQL output. After the column name in the SELECT list, you specify the following:

- LABEL= option to alter the column heading
- FORMAT= option to alter the appearance of the values in that column.

58



The LABEL, FORMAT, INFORMAT, and LENGTH options are not part of the ANSI standard, but are SAS enhancements.

Enhancing Query Output

Example: Enhance the report. Display the navigators and their salaries.

```
proc sql;  
    select EmpID label='Employee Identifier',  
           JobCode label='Job Code',  
           Salary label='Annual Salary'  
           format=dollar12.2  
    from airline.payrollmaster  
    where JobCode contains 'NA'  
    order by Salary desc;
```

59

c2s4d03

Enhanced Query Output

Output

The SAS System		
Employee Identifier	Job Code	Annual Salary
1352	NA2	\$75,317.20
1417	NA2	\$73,178.00
1935	NA2	\$71,513.40
1839	NA1	\$60,806.20
1443	NA1	\$59,183.60
1332	NA1	\$59,049.20
1269	NA1	\$58,366.00
1111	NA1	\$56,820.40

60

To force PROC SQL to ignore permanent labels in a table, specify the NOLABEL system option.

Enhancing Query Output

Here are examples of enhancing output:

- define a column containing a character constant by placing a text string in the SELECT list
- use SAS titles and footnotes to enhance the query's appearance

61

Enhancing Query Output

Example: Display bonus values for all flight engineers.

```
proc sql;
title 'Current Bonus Information';
title2 'Navigators - All Levels';
select EmpID
       label='Employee Identifier',
       'bonus is:',
       Salary *.05 format=dollar12.2
from airline.payrollmaster
where JobCode contains 'NA'
order by Salary desc;
```

62

c2s4d04

TITLE and FOOTNOTE statements must precede the SELECT statement.

Also, a column of numeric values can be defined in a way similar to the above character constant 'bonus is:'.
'bonus is:'.

Enhancing Query Output

Output

Current Bonus Information Navigators - All Levels		
Employee Identifier		
1352	bonus is:	\$3,765.86
1417	bonus is:	\$3,658.90
1935	bonus is:	\$3,575.67
1839	bonus is:	\$3,040.31
1443	bonus is:	\$2,959.18
1332	bonus is:	\$2,952.46
1269	bonus is:	\$2,918.30
1111	bonus is:	\$2,841.02

63



Exercises

Submit a LIBNAME statement to assign the libref **airline** to the SAS data library for this course.

TSO: libname airline '.sql.sasdata';

Directory-based systems: libname airline '.';

1. Querying a Table

- a. Submit a PROC SQL query that displays all rows and all columns of **airline.payrollmaster**.
- b. Recall the previous query and alter it so that only the columns for employee ID, gender, job code, and salary are displayed.
- c. Recall the previous query and alter it so that a new column is displayed as one third of the employee's salary. Name the new column **Tax**.
- d. Recall the previous query and alter it so that the Tax and Salary columns are displayed with commas and two decimal places.
- e. Recall the previous query and alter it so that only male employees are listed.
- f. Recall the previous query and alter it so that only male flight attendants are displayed.

2. Eliminating Duplicates

Use the **airline.staffmaster** table to create a report that displays the cities where airline employees reside. The report must contain only one row per city, be ordered by city, and have an appropriate title.

Cities Where Employees Live	
City	
BRIDGEPORT	
MT. VERNON	
NEW YORK	
PATERSON	
PRINCETON	
STAMFORD	
WHITE PLAINS	

3. Subsetting Data

Use the `airline.marchflights` table to create a report that shows all flights whose total number of passengers is less than one third of the airplane's capacity. Display the flights in descending number of total passengers. Create an appropriate title.



Total is the sum of **Boarded**, **Transferred** and **Nonrevenue**.

Flights Less Than One Third Full					
FlightNumber	Date	Destination	Total	Passenger Capacity	
290	19MAR2000	WAS	59	180	
523	05MAR2000	ORD	59	210	
290	05MAR2000	WAS	55	180	
183	19MAR2000	WAS	53	180	
982	12MAR2000	DFW	49	180	
183	25MAR2000	WAS	43	180	
302	31MAR2000	WAS	34	180	
302	22MAR2000	WAS	33	180	
416	05MAR2000	WAS	31	180	
872	21MAR2000	LAX	.	210	
921	27MAR2000	DFW	.	180	

4. Querying Data (Optional)

A customer service representative must contact a person in the frequent flyer table, but the service representative only remembers that the person's first name begins with an N. Use the `airline.frequentflyers` table to list the names of all possible people.

Frequent Fliers with First Names Beginning with an 'N'		
Name	Frequent FlyerNumber	
CARAWAY, NEIL	WD4762	
CHAPMAN, NEIL	WD8968	
OVERBY, NADINE	WD5201	
WILDER, NEIL	WD6169	
JONES, NATHAN	WD1961	
TUCKER, NEIL	WD2719	
WELLS, NADINE	WD6504	
SANDERSON, NATHAN	WD7916	

5. Using SAS Functions (Optional)

Query the **airline.payrollmaster** table to determine how old each employee was when the employee was hired. Display the employee's ID, birth date, hire date, and age at time of employment. Format the two dates with the MMDDYY10. format and label each column appropriately.

Partial Output

Employee Age Information			
Employee ID	Birth Date	Hire Date	Age At Employment
1919	09/16/1958	06/07/1985	26
1653	10/19/1962	08/12/1988	25
1400	11/08/1965	10/19/1988	22
1350	09/04/1963	08/01/1988	24
1401	12/16/1948	11/21/1983	34
1499	04/29/1952	06/11/1978	26
1101	06/09/1960	10/04/1988	28
1333	04/03/1959	02/14/1979	19

2.5 Summarizing Data

Objectives

- Use functions to summarize data in a query.

66

Summary Functions

Example: Find the total number of passengers for each flight in March.

```
proc sql;  
  select Date, FlightNumber, Boarded,  
         Transferred, Nonrevenue,  
         sum(Boarded, Transferred, Nonrevenue)  
         as Total  
  from airline.marchflights;
```

This calculation
is performed *across*
columns
for each row.

67

c2s5d01

Summary Function

The SAS System						
Date	Flight Number	Boarded	Transferred	Nonrevenue	Total	
01MAR2000	182	104	16	3	123	
01MAR2000	114	172	18	6	196	
01MAR2000	202	151	11	5	167	
01MAR2000	219	198	17	7	222	
01MAR2000	439	167	13	5	185	
01MAR2000	387	152	8	3	163	
01MAR2000	290	96	+	16	+	7 = 119

If you specify more than one column name in a summary function, the function acts like a DATA step function. The calculation is performed for each row.

68

Summary Functions

If you specify only one column name in a summary function, the statistic is calculated down the column.

Example: Determine the average salary for the company.

```
proc sql;
  select avg(Salary) as MeanSalary
  from airline.payrollmaster;
```

69

c2s5d02

Summary Function

Data set **airline.payrollmaster**

Employee Number	Gender	Job Code	Salary
1919	M	TA2	\$48,126
1653	F	ME2	\$49,151
1400	M	ME1	\$41,677
1350	F	FA3	\$46,040
1401	M	TA3	\$54,351
1499	M	ME3	\$60,235
1101	M	SCP	\$26,212
1333	M	PT2	\$124,048
1402	M	TA2	\$45,661
1479	F	TA3	\$54,299
1403	M	ME1	\$39,301
1739	M	PT1	\$93,124
1658	M	SCP	\$25,120
1428	F	PT1	\$96,274
1782	M	ME2	\$49,483
1244	M	ME2	\$51,695
1383	M	BCK	\$36,152

Avg(Salary) gives a mean for all rows within this column only.

70

This is comparable to a SAS procedure (for example, the MEANS procedure), which computes statistics on table columns.

Summary Functions

Output

The SAS System
MeanSalary
54079.65

71

The SQL procedure supports numerous functions for calculating statistics. Some functions have more than one name to accommodate both SAS and SQL conventions.

Summary Functions

The following are selected functions:

AVG, MEAN	mean or average value
COUNT, FREQ, N	number of nonmissing values
MAX	largest value
MIN	smallest value
NMISS	number of missing values
STD	standard deviation
SUM	sum of values
VAR	variance

72

Summary Functions

Example: Add the JobCode column to the summarized query.

```
proc sql;  
  select JobCode, avg(Salary) as average  
  from airline.payrollmaster;
```

73

c2s5d03

Summary Functions

Partial Output

The SAS System	
Job Code	average
TA2	54079.65
ME2	54079.65
ME1	54079.65
FA3	54079.65
TA3	54079.65
ME3	54079.65
SCP	54079.65
PT2	54079.65
TA2	54079.65
TA3	54079.65
ME1	54079.65

74



By default, summary functions calculate statistics based on the entire table. The average is calculated and then re-merged with the individual rows in the table.

How can you find the average salary for each job code?

Grouping Data

You can use the GROUP BY clause to

- classify the data into groups based on the values of one or more columns
- calculate statistics for each unique value of the grouping columns.

75

Grouping Data

Example: Display the average salary for each job code.

```
proc sql;
  select JobCode, avg(Salary) as
         average format=dollar11.2
  from airline.payrollmaster
  group by JobCode;
```

76

c2s5d04

Grouping Data

Partial Output

The SAS System

Job Code	average
BCK	\$36,111.91
FA1	\$32,255.11
FA2	\$39,181.63
FA3	\$46,107.40
ME1	\$39,900.35
ME2	\$49,807.60
ME3	\$59,375.00
NA1	\$58,845.08
NA2	\$73,336.20
PT1	\$95,071.20
PT2	\$122,253.60
PT3	\$154,706.30
SCP	\$25,632.40

77

Analyzing Groups of Data

The COUNT(*) summary function counts the number of rows.

Example: Determine the total number of employees.

```
proc sql;
  select count(*) as count
  from airline.payrollmaster;
```

Output

The SAS System	
count	
	148

78

c2s5d05

The COUNT function is the only function that allows an asterisk (*) as an argument.

Analyzing Groups of Data

Example: Determine the total number of employees within each job category.

```
proc sql;
  select substr(JobCode,1,2)
         label='Job Category',
         count(*) as count
  from airline.payrollmaster
  group by 1;
```

79

c2s5d06

Analyzing Groups of Data

Output

The SAS System	
Job Category	count
BC	9
FA	34
ME	29
NA	8
PT	20
SC	7
TA	41

80

Analyzing Groups of Data

Example: Calculate each navigator's salary as a percentage of all navigators' salaries.

```
proc sql;
  select EmpID, Salary,
         (Salary/sum(Salary)) as percent
         format=percent8.2
  from airline.payrollmaster
  where JobCode contains 'NA';
```

81

c2s5d07

Analyzing Groups of Data

Output

The SAS System		
Emp ID	Salary	percent
1269	\$58,366	11.35%
1935	\$71,513	13.91%
1417	\$73,178	14.23%
1839	\$60,806	11.82%
1111	\$56,820	11.05%
1352	\$75,317	14.65%
1332	\$59,049	11.48%
1443	\$59,184	11.51%

82

PROC SQL automatically re-merges the summary statistic with the table to calculate the percentage. This requires two passes through the data: one to compute the column sum and another to compute each row's percentage of the total. A note appears in the SAS log when re-merging occurs.

Partial Log

NOTE: The query requires remerging summary statistics back with the original data.

Selecting Groups of Data with the HAVING Clause

The WHERE clause selects data based on values for individual rows. To select entire groups of data, use the HAVING clause.

Example: Display all job codes with an average salary of more than \$56,000.

```
proc sql;
  select JobCode, avg(Salary) as average
         format=dollar11.2
  from airline.payrollmaster
  group by JobCode
  having avg(Salary) > 56000 ;
```

83

c2s5d08

Alternatively, you can code the HAVING clause as follows:

```
having average > 56000;
having calculated average > 56000;
```

Selecting Groups of Data with the HAVING Clause

Output

The SAS System	
Job Code	average
ME3	\$59,375.00
NA1	\$58,845.08
NA2	\$73,336.20
PT1	\$95,071.20
PT2	\$122,253.60
PT3	\$154,706.30

84

2.6 Subqueries

Objectives

- Describe how to subset data based on values returned from other queries.
- Explain the difference between a correlated and noncorrelated subquery.

86

Subqueries

Subqueries have the following characteristics:

- are inner queries that return values to be used by an outer query to complete a subsetting expression in a WHERE or HAVING clause
- return single or multiple values to be used by the outer query
- can return only a single column

87

Subqueries are also known as nested queries, inner queries, and sub-selects.

Subqueries: Noncorrelated

Example: Display job codes where the group's average salary exceeds the company's average salary.

```
proc sql;  
  select JobCode, avg(Salary) as MeanSalary  
  from airline.payrollmaster  
  group by JobCode  
  having avg(Salary) >   
    (select avg(Salary)  
     from airline.payrollmaster);
```

Evaluate first

Then pass results to outer query

88

c2s6d01
...

Subqueries: Noncorrelated

After the subquery executes, the outer query code resolves to:

```
proc sql;  
  select JobCode, avg(Salary) as MeanSalary  
  from airline.payrollmaster  
  group by JobCode  
  having avg(Salary) > (54079.65);
```

89

Subqueries: Noncorrelated

Output

The SAS System	
Job Code	MeanSalary
ME3	59375
NA1	58845.08
NA2	73336.2
PT1	95071.2
PT2	122253.6
PT3	154706.3
TA3	55551.42

90

Subqueries: Noncorrelated

Example: Send birthday cards to employees with February birthdays. Names and addresses are in `airline.staffmaster`, and birth dates are in `airline.payrollmaster`.

```
proc sql;
  select EmpID, LastName, FirstName,
         City, State
  from airline.staffmaster
  where EmpID in
         (select EmpID
          from airline.payrollmaster
          where month(DateOfBirth)=2);
```

91

c2s6d02

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
  select EmpID, LastName,
         FirstName, City, State
  from airline.staffmaster
  where EmpID in
    (select EmpID
     from airline.payrollmaster
     where month(DateOfBirth)=2);
```

airline.payrollmaster
Partial Listing

EmpID	DateOfBirth
...	...
1038	11/13/1967
1420	02/23/1963
1561	12/03/1961
1434	07/14/1960
1414	03/28/1970
1112	12/03/1962
1390	02/23/1963
1332	09/20/1968
...	...

Step 1: Evaluate the inner query and build a virtual table that satisfies the WHERE criteria.

92

c2s6d02

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
  select EmpID, LastName,
         FirstName, City, State
  from airline.staffmaster
  where EmpID in
    (select EmpID
     from airline.payrollmaster
     where month(DateOfBirth)=2);
```

airline.payrollmaster
Partial Listing

EmpID	DateOfBirth
...	...
1038	11/13/1967
1420	02/23/1963
1561	12/03/1961
1434	07/14/1960
1414	03/28/1970
1112	12/03/1962
1390	02/23/1963
1332	09/20/1968
...	...

Virtual table contains
'1420','1390','1403','1404','1834','1103'.

93

c2s6d02

Noncorrelated Subqueries: How Do They Work?

```
proc sql;
  select EmpID, LastName,
         FirstName, City, State
  from airline.staffmaster
  where EmpID in
    (select EmpID
     from airline.payrollmaster
     where month(DateOfBirth)=2);
```

airline.payrollmaster
Partial Listing

EmpID	DateOfBirth
...	...
1038	11/13/1967
1420	02/23/1963
1561	12/03/1961
1434	07/14/1960
1414	03/28/1970
1112	12/03/1962
1390	02/23/1963
1332	09/20/1968
...	...

Step 2: Pass
'1420','1390','1403','1404','1834','1103'
to the outer query.

c2s6d02

94

Noncorrelated Subqueries: Output

The SAS System				
Emp ID	LastName	FirstName	City	State
1403	BOWDEN	EARL	BRIDGEPORT	CT
1404	CARTER	DONALD	NEW YORK	NY
1834	LONG	RUSSELL	NEW YORK	NY
1103	MCDANIEL	RONDA	NEW YORK	NY
1420	ROUSE	JEREMY	PATERSON	NJ
1390	SMART	JONATHAN	NEW YORK	NY

Does this look
familiar?

95

Selecting Data

If you specify the ANY keyword before a subquery, the comparison is true if it is true for any of the values that the subquery returns.

Keyword ANY	Signifies...
> ANY(20,30,40) returned from inner query	>20
< ANY(20,30,40) returned from inner query	< 40
= ANY(20,30,40) returned from inner query	=20 or =30 or =40

96

The ANY Keyword

Example: Are any low-level flight attendants (FA1 or FA2) older than any of the high-level flight attendants (FA3)?

Think

<select max(DateOfBirth)

```
proc sql;
title "FA1's or FA2's Older Than ANY FA3's";
select EmpID, JobCode, DateOfBirth
from airline.payrollmaster
where JobCode in ('FA1','FA2')
and DateOfBirth < any
(select DateOfBirth
from airline.payrollmaster
where JobCode='FA3');
```

97

c2s6d03



This would be equivalent to asking, “Who is older than any single level-3 flight attendant?”
An alternative WHERE clause is

```
where JobCode in ('FA1','FA2') and DateOfBirth <
(select max(DateOfBirth) from...);
```

The ANY Keyword

Partial Output

FA1's or FA2's Older Than ANY FA3's

Emp ID	Job Code	DateOfBirth
1574	FA2	01MAY1958
1475	FA2	19DEC1959
1124	FA1	14JUL1956
1422	FA1	08JUN1962
1368	FA2	15JUN1959
1411	FA2	31MAY1959
1477	FA2	25MAR1962
1970	FA1	29SEP1962
1413	FA2	20SEP1963
1434	FA2	15JUL1960
1390	FA2	23FEB1963

98

The ALL Keyword

The ALL keyword is true only if the comparison is true for all values returned.

Keyword ALL	Signifies...
> ALL(20,30,40) returned from inner query	> 40
< ALL(20,30,40) returned from inner query	< 20

99

Selecting Data

Example: Are there FA1's or FA2's who are older than all of the FA3's?

Think

<select min(DateOfBirth)

```
proc sql;
title "FA1's or FA2's Older Than ALL
FA3's";
  select EmpID, JobCode, DateOfBirth
  from airline.payrollmaster
  where JobCode in('FA1','FA2')
  and DateOfBirth < all
    (select DateOfBirth
     from airline.payrollmaster
     where JobCode='FA3');
```

100

c2s6d04

An alternative WHERE clause is

```
where JobCode in('FA1','FA2') and DateOfBirth <
(select min(DateOfBirth) from ...);
```

Selecting Data

Output

FA1's or FA2's Older Than ALL FA3's

Emp ID	Job Code	DateOfBirth
1124	FA1	13JUL1956
1415	FA2	12MAR1956

101

Correlated Subqueries

Rules for correlated subqueries include the following:

- cannot be evaluated independently, but depend on the values returned by the outer query for their results
- are evaluated for each row in the outer query

102

Correlated Subqueries

Example: Display the names and states of all navigator managers.

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
           supervisors.EmpID) ;
```

You must qualify each column with a table name.

103

c2s6d05
...

What does it mean to *qualify* a column? When a column appears in more than one table, the column name is preceded with the table name or alias to avoid ambiguity. In this example you use the table names **staffmaster** and **supervisors** in front of the column name of **EmpID**. Although table aliases are not used in this example, they are merely table nicknames and are discussed further in Section 3.2.

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Step 1: The outer query takes the first row in `airline.staffmaster` and finds the `EmpID`, `LastName`, `FirstName`, and `State`.

airline.staffmaster

Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors

Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA

104

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Step 2: Match `staffmaster.EmpID` with `supervisors.EmpID` to find the qualifying row in `airline.supervisors`.

```
airline.staffmaster.EmpID =
airline.supervisors.EmpID?
```

NO MATCH

airline.staffmaster

Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors

Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA

105

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Steps 1 and 2 (repeated): Read the next row from `airline.staffmaster`, and identify the qualifying row in `airline.supervisors`.

`airline.staffmaster.EmpID =`
`airline.supervisors.EmpID?`

106

MATCH

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA
1126	NY	TA

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Resolves
to FALSE

Step 3: The inner query now passes the `JobCategory` of the selected row in `airline.supervisors` back to the outer query via the `=` operator, where `JobCategory` is matched for selection in the outer query.

107

`airline.staffmaster.EmpID =`
`airline.supervisors.EmpID?`

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA
1126	NY	TA

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Continue repeating steps 1 - 3 until all rows are read from **airline.staffmaster**.

airline.staffmaster.EmpID =
airline.supervisors.EmpID?

108

NO MATCH

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID =
       supervisors.EmpID);
```

Continue repeating steps 1 - 3 until all rows are read from **airline.staffmaster**.

airline.staffmaster.EmpID =
airline.supervisors.EmpID?

109

MATCH

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA
1126	NY	TA

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID=
       supervisors.EmpID);
```

Resolves
to TRUE

Pass JobCategory from
airline.supervisors to outer
query for comparison.

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA

110

Correlated Subqueries

```
proc sql;
  select LastName, FirstName, State
  from airline.staffmaster
  where 'NA' =
    (select JobCategory
     from airline.supervisors
     where staffmaster.EmpID=
       supervisors.EmpID);
```

Resolves
to TRUE

Write LastName, FirstName, and
State from airline.staffmaster
as the first row in a newly created
report.

airline.staffmaster Partial Listing

EmpID	LastName	FirstName
1919	ADAMS	GERALD
1401	AVERY	JERRY
1269	CASTON	FRANKLIN
1935	FERNANDEZ	KATRINA
1124	FIELDS	DIANA
1677	KRAMER	JACKSON
1442	NEWKIRK	SANDRA
1417	NEWKIRK	WILLIAM
1352	RIVERS	SIMON

airline.supervisors Partial Listing

EmpID	State	JobCategory
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1385	CT	ME
1420	NJ	ME
1882	NY	ME
1935	CT	NA
1417	NJ	NA
1352	NY	NA
1106	CT	PT
1442	NJ	PT
1405	NJ	SC
1564	NY	SC
1401	NJ	TA

111

Correlated Subqueries

Build first row of report:

LastName	FirstName	State
FERNANDEZ	KATRINA	CT

112

SAS continues this process until all rows are read from the table referred to in the outer query, **airline.staffmaster**. At that point the third and final row of the report is written, as noted in the following slide.

Correlated Subqueries

Build third (and final) row of report:

LastName	FirstName	State
FERNANDEZ	KATRINA	CT
NEWKIRK	WILLIAM	NJ
RIVERS	SIMON	NY

125

Correlated Subqueries

The EXISTS condition tests for the existence of a set of values returned by the subquery.

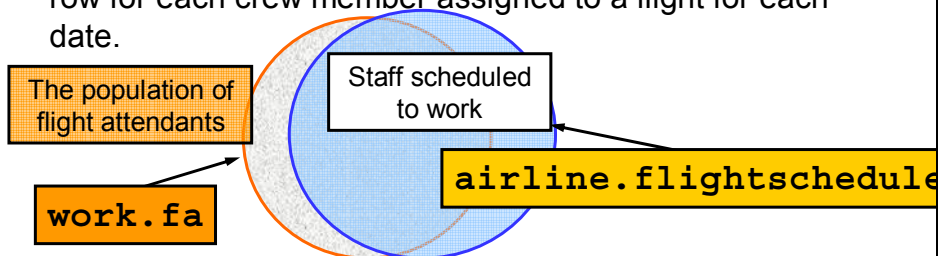
- The EXISTS condition is true if the subquery returns at least one row.
- The NOT EXISTS condition is true if the subquery returns no data.

126

Correlated Subqueries

Example: The temporary table **work.fa** is a subset of **airline.staffmaster** containing the names and IDs of all flight attendants.

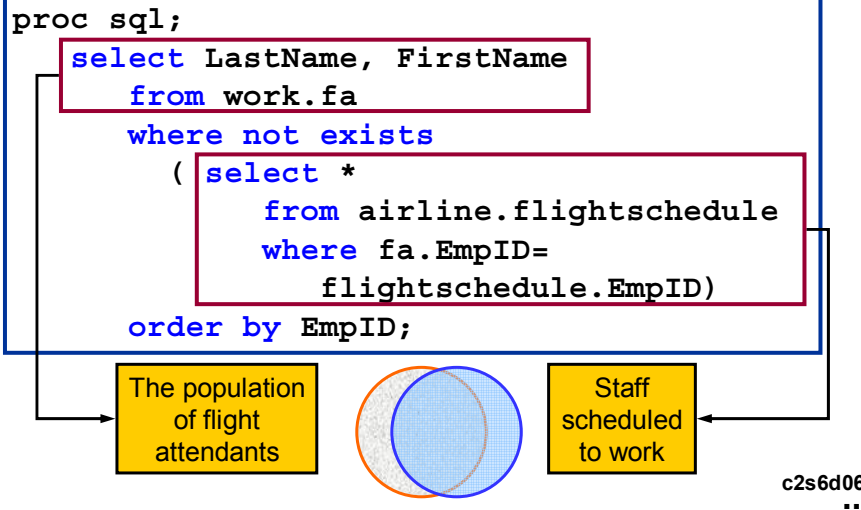
The **airline.flightschedule** table contains a row for each crew member assigned to a flight for each date.



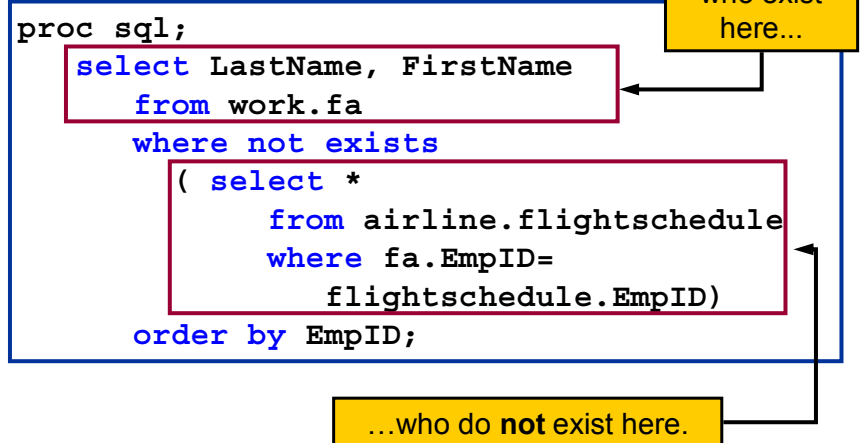
127

Correlated Subqueries

Example: Determine which flight attendants were not scheduled.



Correlated Subqueries



Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

What you want: Those
who are here...
but not here....

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule
Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

130

c2s6d06

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

MATCH

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule
Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

131

c2s6d06

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
    from work.fa
   where not exists
      (select *
        from airline.flightschedule
       where fa.EmpID=
             flightschedule.EmpID)
   order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

MATCH

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

132

c2s6d06

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
    from work.fa
   where not exists
      (select *
        from airline.flightschedule
       where fa.EmpID=
             flightschedule.EmpID)
   order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

NO MATCH

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

133

c2s6d06

...

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

Therefore, the NOT EXISTS clause is satisfied and the first row is written to the Output report.

LastName	FirstName
VEGA	FRANKLIN

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

c2s6d06
...

134

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

MATCH

LastName	FirstName
VEGA	FRANKLIN

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

c2s6d06
...

135

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

NO MATCH

LastName	FirstName
VEGA	FRANKLIN

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule
Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

c2s6d06
...

136

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

LastName	FirstName
VEGA	FRANKLIN
PATTERSON	RENEE

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule
Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

137

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

fa.EmpID = flightschedule.EmpID?

MATCH

LastName	FirstName
VEGA	FRANKLIN
PATTERSON	RENEE

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

c2s6d06

138

Correlated Subqueries

```
proc sql;
  select LastName, FirstName
  from work.fa
  where not exists
    (select *
     from airline.flightschedule
     where fa.EmpID=
           flightschedule.EmpID)
  order by EmpID;
quit;
```

PROC SQL terminated at EOF.

LastName	FirstName
VEGA	FRANKLIN
PATTERSON	RENEE

work.fa Partial Listing

FirstName	LastName	EmpID
..
RANDALL	PETERS	1413
NATHAN	SANDERSON	1414
FRANKLIN	VEGA	1415
MARIE	FLETCHER	1422
RENEE	PATTERSON	1424
JENNY	UNDERWOOD	1425
..

EOF

airline.flightschedule Partial Listing

Flight Number	Date	Destination	EmpID
..
182	02MAR2000	YYZ	1414
271	02MAR2000	CDG	1413
132	03MAR2000	YYZ	1425
821	03MAR2000	LHR	1422
..

← **Final Report**

c2s6d06

139

Correlated Subqueries

Final Report

LastName	FirstName
PATTERSON	RENEE
VEGA	FRANKLIN



Exercises

Submit a LIBNAME statement to assign the libref **airline** to the course SAS data library.
(TSO only: DISP=SHR)

TSO: libname airline '.sql.sasdata';

Directory-based systems: libname airline '.';

6. Summarizing Data

Use the **airline.flightdelays** table to answer the questions.

- What was the maximum delay experienced for any flight in the table?
- What was the maximum delay experienced by each of the destinations?
- What was the average delay for each destination?
- Which destinations have an average delay that is greater than the overall average delay?

Hint: Use a subquery to calculate the overall average delay.

7. Summarizing Data in Groups

Use the **airline.staffmaster** table to determine the number of people employed by the airline in each city. Order the report by state and city.

Output

Number of Employees in Each City		
State	City	Number Employed
CT	BRIDGEPORT	19
CT	STAMFORD	25
NJ	PATERSON	5
NJ	PRINCETON	10
NY	MT. VERNON	5
NY	NEW YORK	79
NY	WHITE PLAINS	5

8. Subqueries

Each month a memo is posted that lists the employees who have employment anniversaries for that month. Create the report for February and list the first and last names of all employees who were hired during the month of February of any year. You can find employee names in the **airline.staffmaster** table, and employee hire dates in the **airline.payrollmaster** table. Order the report by employee last name.

Output

Employees with February Anniversaries	
FirstName	LastName
JUSTIN	BLAIR
JASON	BROWN
GERALD	FOSTER
MICHAEL	HOWARD
MARILYN	REED
WAYNE	THOMPSON
KATHY	TRIPP
FRANKLIN	VEGA
CHIN	WANG
ELAINE	WARD
CAROLYN	WHALEY

9. Correlated Subqueries (Optional)

Create a report that shows the number of employees who are frequent flyers of the airline. Employees are listed in **airline.staffmaster** and frequent flyers are in **airline.frequentflyers**. (Hint: Names are stored differently in the two tables.)

Output

Number of Employees Listed in Frequent Flyer Table
count
136

10. Summarizing Data (Optional)

- a. Use the **airline.flightdelays** table to determine the number of times each flight was delayed. (Hint: There should be one row for each flight.) Order the report by flight number.

Partial Output

Delayed Arrivals		
FlightNumber	Destination	Times Delayed
114	LAX	19
132	YYZ	21
182	YYZ	15
183	WAS	16
202	ORD	16

- b. Use the **airline.flightdelays** table to determine the number of times each flight was on time (Times Delayed=0). Order your report by flight number.

Partial Output

On-Time Arrivals		
FlightNumber	Destination	On-time Count
114	LAX	3
132	YYZ	2
182	YYZ	3
183	WAS	3
202	ORD	2

2.7 Chapter Summary

The SQL procedure enables you to use SQL statements in a SAS program. When you use the SQL procedure, you do not need to repeat the PROC SQL statement with each query, and you do not need a RUN statement. Results of the query are displayed automatically and can be ordered. Queries contain statements that are composed of clauses.

A SELECT statement is used to query one or more SAS data sets. Use the SELECT statement to retrieve data from a table and to specify how to display a report.

You can use the VALIDATE keyword to verify the validity of the query's syntax. Messages are printed in the SAS log.

You can calculate new columns by using expressions or DATA step functions. You can subset rows by using a WHERE clause or eliminate duplicate rows by using the DISTINCT keyword. The CALCULATED keyword enables you to use a previously calculated value elsewhere in the query. Use a GROUP BY clause to apply summary functions to groups of values and include an ORDER BY clause to sort the output. You can customize output with SAS formats, labels, and titles.

Summary functions are available to summarize data for the entire table or for groups of data in the table. You can select groups of data to be processed by using a HAVING clause.

You can use a subquery to select data from a table based on the result returned by another query. Subqueries are typically used in a WHERE or HAVING clause and are evaluated before the outer query. A correlated subquery is a subquery that depends on values returned by the outer query.

General form of the SELECT statement:

```
SELECT column-1<, column-2> ...  
FROM table-1|view-1<, table-2|view-2> ...  
<WHERE expression>  
<GROUP BY column-1<, column-2> ...>  
<HAVING expression>  
<ORDER BY column-1<, column-2> ... <DESC>>;
```

2.8 Solutions to Exercises

1. Querying a Table

a.

```
proc sql;
  select *
    from airline.payrollmaster;
```

b.

```
select EmpID, Gender, JobCode, Salary
  from airline.payrollmaster;
```

c.

```
select EmpID, Gender, JobCode, Salary,
       Salary/3 as Tax
  from airline.payrollmaster;
```

d.

```
select EmpID, Gender, JobCode,
       Salary format=comma10.2,
       Salary/3 as Tax format=comma10.2
  from airline.payrollmaster;
```

e.

```
select EmpID, Gender, JobCode,
       Salary format=comma10.2,
       Salary/3 as Tax format=comma10.2
  from airline.payrollmaster
 where Gender='M';
```

f.

```
select EmpID, Gender, JobCode,
       Salary format=comma10.2,
       Salary/3 as Tax format=comma10.2
  from airline.payrollmaster
 where Gender='M' and JobCode contains 'FA';
quit;
```

2. Eliminating Duplicates

```
proc sql;
title 'Cities Where Employees Live';
  select distinct City
    from airline.staffmaster
   order by City;
quit;
title;
```

3. Subsetting Data

```
proc sql;
title 'Flights Less Than One Third Full';
  select FlightNumber, Date, Destination,
         Boarded+Transferred+Nonrevenue as Total,
         PassengerCapacity
  from airline.marchflights
  where calculated Total<(PassengerCapacity/3)
  order by 4 desc;
quit;
title;
```

4. Querying Data (Optional)

```
proc sql;
title "Frequent Fliers with First Names Beginning with an 'N'";
  select Name, ffid
  from airline.frequentflyers
  where Name like '%, N%';
quit;
title;
```

Alternate Solution

```
select Name, ffid
  from airline.frequentflyers
  where left(scan(Name,2,',')) like 'N%';
```

5. Using SAS Functions (Optional)

```
proc sql;
title 'Employee Age Information';
  select EmpID label='Employee ID',
         DateOfBirth format=mmddyy10.
         label='Birth Date',
         DateOfHire format=mmddyy10.
         label='Hire Date',
         int((DateOfHire-DateOfBirth)/365.25)
         label='Age At Employment'
  from airline.payrollmaster;
quit;
title;
```

6. Summarizing Data

a.

```
proc sql;
title 'Maximum Delay Experienced';
  select max(Delay) label='Max Delay'
  from airline.flightdelays;
```


b.

```

title 'Maximum Delay Experienced';
title2 'by Each Destination';
  select Destination,
         max(Delay) label='Max Delay'
  from airline.flightdelays
 group by Destination;

```

c.

```

title 'Average Delay for Each Destination';
  select Destination,
         avg(Delay) label='Average Delay'
  from airline.flightdelays
 group by Destination;

```

d.

```

title 'Destinations Having Average Delay';
title2 'Exceeding Overall Average';
  select Destination,
         avg(Delay) label='Average Delay'
  from airline.flightdelays
 group by Destination
 having avg(Delay) >
        (select avg(Delay)
         from airline.flightdelays);
quit;
title;

```

7. Summarizing Data in Groups

```

proc sql;
title 'Number of Employees in Each City';
  select State, City,
         count(*) label='Number Employed'
  from airline.staffmaster
 group by State, City
 order by State, City;
quit;
title;

```

8. Subqueries

```

proc sql;
title 'Employees with February Anniversaries';
  select FirstName, LastName
  from airline.staffmaster
  where EmpID in
        (select EmpID
         from airline.payrollmaster
         where month(DateOfHire)=2);
quit;
title;

```

9. Correlated Subqueries (Optional)

```
proc sql;
title 'Number of Employees Listed';
title2 'in Frequent Flyer Table';
  select count(*) as count
    from airline.frequentflyers
   where exists
      (select *
       from airline.staffmaster
      where Name=trim(LastName) || ', ' || FirstName);
quit;
title;
```

10. Summarizing Data (Optional)

a.

```
proc sql;
title 'Delayed Arrivals';
  select FlightNumber, Destination,
         count(*) label='Times Delayed'
    from airline.flightdelays
   where Delay>0
  group by FlightNumber, Destination
 order by FlightNumber;
```

b.

```
title 'On-time Arrivals';
  select FlightNumber, Destination,
         count(*) label='On-time Count'
    from airline.flightdelays
   where Delay=0
  group by FlightNumber, Destination
 order by FlightNumber;
quit;
title;
```

Chapter 3 Combining Tables

3.1	Overview.....	3-3
3.2	Joins	3-5
3.3	Complex Joins	3-27
3.4	Set Operators	3-46
3.5	Chapter Summary.....	3-77
3.6	Solutions to Exercises	3-79

3.1 Overview

Objectives

- Distinguish between joins and set operations.

3

Combining Data from Multiple Tables

Joins combine tables horizontally (side by side).

Table A	Table B
---------	---------

4

Combining Data from Multiple Tables

Set operations combine tables vertically (one on top of the other).



5

Which DATA step statements perform similar operations?

3.2 Joins

Objectives

- Describe the different joins available in PROC SQL.
- Use a table alias.
- Compare SQL joins to DATA step merges.

7

Types of Joins

PROC SQL supports the following two types of joins:

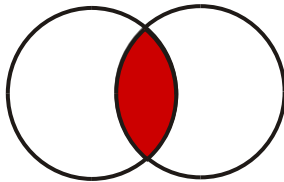
- inner joins
- outer joins

8

Types of Joins

Inner joins have the following characteristics:

- return only matching rows
- allow a maximum of 32 tables to be joined at the same time



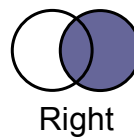
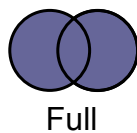
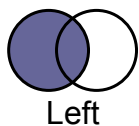
9



If the join involves views, the number of tables underlying the views, not the views themselves, counts toward the limit of 32.

Outer Joins

You can retrieve nonmatching rows, as well as matching rows, by using an outer join. Outer joins are limited to two tables at a time.



35

Cartesian Product


A query that lists multiple tables in the FROM clause, without row restrictions, results in all possible combinations of rows from all tables. This is called a *Cartesian product*.

```
select *
      from one, two;
```

11

c3s2d01

Cartesian Product



21

...

Cartesian Product

The number of rows in a Cartesian product is the product of the number of rows in the contributing tables.

$$3 \times 3 = 9$$

$$1,000 \times 1,000 = 1,000,000$$

22

...



A Cartesian product is rarely a desired query outcome. The SQL processor prints a warning in the log if a query involved a Cartesian product:

NOTE: The execution of this query involves performing one or more Cartesian product joins that cannot be optimized.

Inner Joins

Inner join syntax resembles Cartesian product syntax, but it has a WHERE clause that restricts how the rows can be combined.

General form of an inner join:

```
SELECT column-1, column-2, ...  
FROM table-1, table-2, ...  
WHERE join-condition(s)  
      <AND other subsetting conditions>  
      <other clauses>;
```

23



The distinguishing characteristics of inner join syntax are

- a list of two or more table names in the FROM clause
- one or more join conditions in the WHERE clause.

Inner Joins

Conceptually, PROC SQL performs the following tasks:

- first builds a Cartesian product
- then applies the specified restriction(s) and removes rows

24

Inner Joins

Table ONE

X	A
1	a
4	d
2	b

Table TWO

X	B
2	x
3	y
5	v

X	A	X	B
1	a	2	x
1	a	3	y
1	a	5	v
4	d	2	x
4	d	3	y
4	d	5	v
2	b	2	x
2	b	3	y
2	b	5	v

```
select *
  from one, two ...
```

25

c3s2d02
...

Inner Joins

Table ONE

X	A
1	a
4	d
2	b

Table TWO

X	B
2	x
3	y
5	v

X	A	X	B
1	a	2	x
1	a	3	y
1	a	5	v
4	d	2	x
4	d	3	y
4	d	5	v
2	b	2	x
2	b	3	y
2	b	5	v

```
select *
  from one, two
 where one.X=two.X;
```



26

c3s2d02
...

Inner Joins

Table ONE

X	A
1	a
4	d
2	b

Table TWO

X	B
2	x
3	y
5	v

X	A	X	B
2	b	2	x

```
select *
  from one, two
 where one.X=two.X;
```

27

c3s2d02

An inner join is sometimes called a conventional join, natural join, or *equijoin*.



Tables do not have to be sorted before they are joined.

Column X exists in both tables and occurs twice in the query result.

Inner Joins

Display the **X** column only once.

Table ONE

X	A
1	a
4	d
2	b

Table TWO

X	B
2	x
3	y
5	v

```
select one.X, a, b
  from one, two
 where one.X=two.X;
```

X	A	B
2	b	x

28

c3s2d03

Inner Joins

Display all combinations of rows with matching keys, including duplicates.

Table THREE

X	A
1	a1
1	a2
2	b1
2	b2
4	d

Table FOUR

X	B
2	x1
2	x2
3	y
5	v

```
select *
  from three, four
 where three.X=four.X;
```

X	A	X	B
2	b1	2	x1
2	b2	2	x1
2	b1	2	x2
2	b2	2	x2

31

c3s2d04

How many rows does a DATA step match-merge produce for $X = 2$?

Inner Joins

Example: Display the names, job codes, and ages of all New York employees.

- Employee names are found in the **airline.staffmaster** table.
- Employee job codes and birth dates are found in the **airline.payrollmaster** table.

32

Inner Joins

```

title 'New York Employees';
  select substr(FirstName,1,1) || '. ' ||
         LastName as Name,
         JobCode,
         int((today()-DateOfBirth)/365.25)
         as Age
  from airline.payrollmaster,
       airline.staffmaster
 where payrollmaster.EmpID=
        staffmaster.EmpID
        and State='NY'
 order by JobCode;

```

33

c3s2d05

Inner Joins

Partial Output (rows 49-58)

New York Employees		
Name	Job Code	Age
J. BOYCE	PT1	43
R. VENTER	PT1	35
D. CARTER	PT2	55
J. NEWTON	PT2	49
L. UPCHURCH	PT2	53
W. THOMPSON	PT2	54
R. STEPHENSON	PT2	57
R. LUFKIN	PT3	56
R. DENNIS	PT3	64
J. BRADLEY	SCP	41

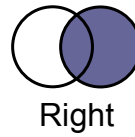
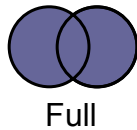
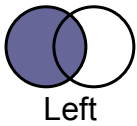
34



This program was run on September 12, 2006. Your results may differ.

Outer Joins

You can retrieve nonmatching rows, as well as matching rows, by using an outer join. Outer joins are limited to two tables at a time.



35



An outer join is an augmentation of an inner join. It returns all the rows generated by an inner join, plus others.

Outer Joins

General form of an outer join:

```
SELECT column <, column> ...
FROM table1
      LEFT|RIGHT|FULL JOIN
      table2
      ON join-condition(s)
      <other clauses>;
```

36

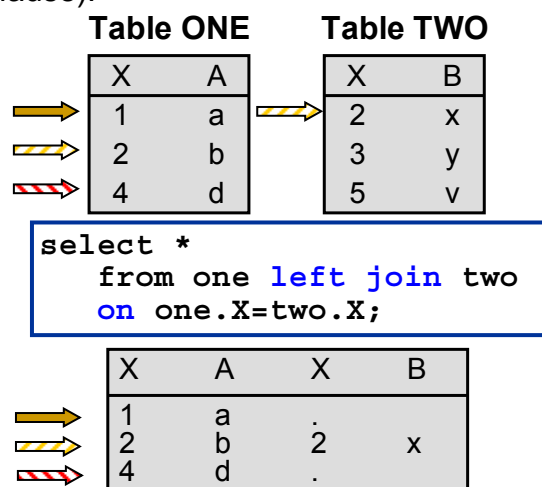
The distinguishing characteristics of outer join syntax are

- exactly two table names flanking one of the three JOIN operators in the FROM clause
- a special ON clause specifying the join condition(s).

A WHERE clause is permitted in order to specify general subsetting conditions.

Outer Joins

A left join retrieves matching rows from both tables, plus nonmatching rows from the left table (the first table in the FROM clause).

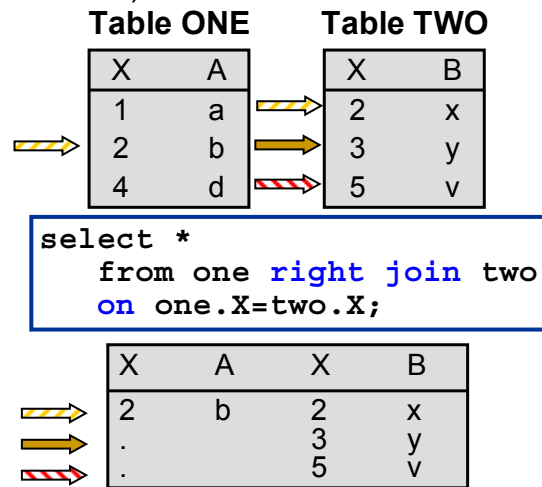


37

c3s2d06

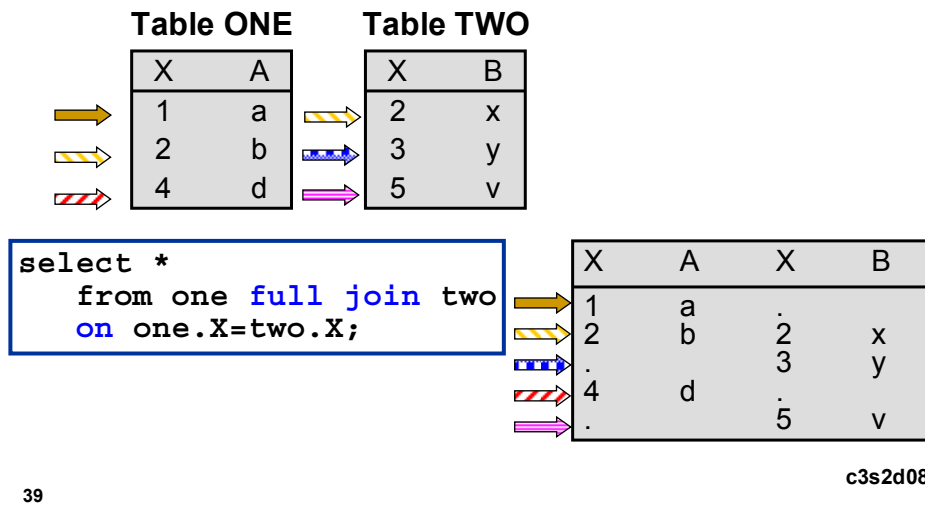
Outer Joins

A right join retrieves matching rows from both tables, plus nonmatching rows from the right table (the second table in the FROM clause).



Outer Joins

A full join retrieves matching rows and nonmatching rows from both tables.



Compare this result with the Cartesian product demonstrated earlier.


You can also write an inner join using this style of syntax:

```
SELECT *
  FROM table-1
  INNER JOIN table-two
  ON table-1.x = table2.x;
```

but the join is limited to two tables.

Outer Joins

Example: List all flights during March with corresponding delay information (if it exists).

 **airline.flightdelays** does not contain delay information for all of the March flights.

40

Outer Joins

```
title 'All March Flights';
proc sql;
    select marchflights.Date,
           marchflights.FlightNumber
              label='Flight Number',
           marchflights.Destination
              label='Left',
           flightdelays.Destination
              label='Right',
           Delay
    from airline.marchflights
    left join
        airline.flightdelays
    on marchflights.Date=flightdelays.Date
       and marchflights.FlightNumber=
           flightdelays.FlightNumber
    order by Delay;
```

41

c3s2d09

Outer Joins

Partial Output

All March Flights				
Date	Flight Number	Left	Right	DelayIn Minutes
16MAR2000	622	FRA		.
03MAR2000	416	WAS		.
17MAR2000	182	YYZ		.
14MAR2000	271	CDG		.
11MAR2000	290	WAS		.
08MAR2000	182	YYZ		.
.	132	YYZ		.
11MAR2000	202	ORD		.
29MAR2000	829	WAS		.
25MAR2000	872	LAX		.
22MAR2000	183	WAS		.
27MAR2000	982	DFW		.
25MAR2000	829	WAS	WAS	-10
18MAR2000	219	LHR	LHR	-10
09MAR2000	821	LHR	LHR	-10

Using a Table Alias

An *alias* is a table nickname. You can assign an alias to a table by following the table name in the FROM clause with the AS keyword and a nickname for the table. Then use the alias in other clauses of the QUERY statement.

43



A table alias is primarily used to reduce the amount of typing required to write a query. It is usually optional. There are, however, two situations that require a table alias:

- a self-join (a table is joined to itself), for example,

```
from airline.staffmaster as s1, airline.staffmaster as s2
```

- when referencing same-named columns from same-named tables in different libraries, for example,

```
from airline.flightdelays as ad,  
     work.flightdelays as wd  
where ad.delay > wd.delay
```

Using a Table Alias

```
select l.Date,
       l.FlightNumber
       label='Flight Number',
       l.Destination label='Left',
       r.Destination label='Right',
       Delay
from airline.marchflights as l
     left join
     airline.flightdelays as r
on l.Date=r.Date and
   l.FlightNumber=r.FlightNumber
order by Delay;
```

44

c3s2d10



The AS keyword is optional in a table alias. The alias can directly follow the table name in the FROM clause.

SQL Join versus DATA Step Merge

A DATA step with MERGE and BY statements combines rows differently from an outer join.

Table ONE

X	A
1	a
2	b
4	d

Table TWO

X	B
2	x
3	y
5	v

Table MERGED

X	A	B
1	a	
2	b	x
3		y
4	d	
5		v

```
data merged;
  merge one two;
  by X;
run;
```

45

c3s2d11

SQL Join versus DATA Step Merge

A DATA step with MERGE and BY statements combines rows differently from an outer join.

Table ONE

X	A
1	a
2	b
4	d

Table TWO

X	B
2	x
3	y
5	v

```
proc sql;
  select one.X, a, b
    from one full join two
    on one.X=two.X;
```

X	A	B
1	a	
2	b	x
.		y
4	d	
.		v

In the SQL procedure, the two X columns are not overlaid by default.

How can you achieve the same result using PROC SQL?

SQL Join versus DATA Step Merge

You can use the COALESCE function to overlay two columns.

Table ONE

X	A
1	a
2	b
4	d

Table TWO

X	B
2	x
3	y
5	v

```
select coalesce(one.X,two.X)
       label='X', a, b
  from one full join two
 on one.X=two.X;
```

X	A	B
1	a	
2	b	x
3		y
4	d	
5		v

47

c3s2d13

The COALESCE function

- returns the first value that is a SAS nonmissing value
- requires all arguments to have the same data type.

If you omit the LABEL= option or an alias in a coalesced column, it appears without a column heading.



If the tables being joined together had multiple matching keyfield values between the tables, the data step and the SQL procedure would not give the same results. SQL would be the only method that generates a Cartesian Product of all the matching rows

SQL Join versus DATA Step Merge

Joins do not require the following:

- sorted or indexed tables
- same-named columns in join expressions
- equality in join expressions

48

Tables can be joined on inequalities, for example,

```
select columns
  from table1 as a, table2 as b
 where  a.itemnumber=b.itemnumber
        and a.cost > b.price;
```

Internal Processing of Joins

Conceptually, during a join, SQL performs the following tasks:

- a Cartesian product is built internally
- WHERE processing selects the appropriate rows

In reality, however, the PROC SQL optimizer breaks the Cartesian product into smaller pieces.

49

SAS data sets are stored in pages that contain a certain number of observations. To reduce input/output, the SQL procedure optimizer uses these pages in its processing.

During a two-way join, the following tasks are completed:

1. The first page from table A is read into memory, with as many of the first pages from table B that can fit into available memory.
2. Valid rows are selected.
3. The first page of table A is kept in memory. All subsequent pages from table B that can fit into memory are read and step 2 is repeated.
4. All pages from table B are processed in combination with page 1 from table A. Steps 1 through 4 are repeated for page 2 from table A. The entire process stops when all rows in both tables are processed.

The SQL procedure optimizer can process an equijoin (a join on an equals condition, for example, **where x.idnum=y.idnum**) more efficiently than a join involving an inequality.

During a two-way equijoin, the following tasks are completed:

1. Both tables are sorted by the matching column (if necessary) and are grouped by the matching column's value into chunks.
2. The Cartesian product is only performed on matching portions of data.
3. After a section of data is processed, it is not processed again.



The SQL procedure optimizer has other algorithms from which to select when you optimize a join. For example, you can use a hashing algorithm when you join a small table with a large table.

In a multiway join (more than two tables), in order to minimize the Cartesian product, the SQL procedure optimizer

- splits the join into a number of two-way joins, and eliminates rows and columns from the intermediate tables as soon as they are no longer required
- decides the order in which the tables are processed
- processes the joins in the order that minimizes the intermediate Cartesian product..

3.3 Complex Joins

Objectives

- Understand techniques that simplify the coding of a complex query.
- Compare solving a problem using PROC SQL with traditional SAS programming.

51

In-Line Views

An *in-line view* has the following characteristics:

- a temporary table that exists only during query execution
- created when a FROM clause contains a query expression in place of a table name

52

In-Line Views

Example: Which destinations experience the worst delays?



How do you define worst delays?

53

In-Line Views

Output

Destination	Average Delay	Maximum Delay	Number of Delays	Number of Early Arrivals	Probability of Delay
WAS	1	15	76	75	0.50
YYZ	2	14	36	24	0.60
DFW	3	20	38	23	0.62
ORD	3	19	51	41	0.55
LAX	5	27	82	41	0.67
LHR	6	30	39	19	0.67
CPH	6	26	16	11	0.59
FRA	6	34	14	12	0.54
CDG	9	39	21	5	0.81

54

In-Line Views

```
select Destination,
       summarized columns,
       late / (late + early) as prob
       format=5.2
       label='Probability of Delay'
from summarized table
order by 2nd column;
```

55

In-Line Views

Boolean expressions can be used in the SELECT clause.

```
select Delay,
       (Delay > 0) as Late
from airline.flightdelays;
```

Partial Output

Delay	Late
0	0
8	1
-5	0
18	1

56

c3s3d01

A Boolean expression resolves either to 1 (true) or 0 (false).

In-Line Views

```
select *, Late/(Late+Early) as prob
       format=5.2 label='Probability of Delay'
from (select Destination,
             avg(Delay) as average
             format=3.0 label='Average Delay',
             max(Delay) as max
             format=3.0 label='Maximum Delay',
             sum(Delay > 0) as late
             format=3.0 label='Number of Delays',
             sum(Delay <= 0) as early
             format=3.0
             label='Number of Early Arrivals'
       from airline.flightdelays
       group by 1)
order by 2;
```

57

c3s3d02

When it is summed, a Boolean expression displays the number of rows that are true. A missing value will force a Boolean expression to resolve to true. Use a WHERE clause to omit missing value from the summarization.

You can use the calculated columns LATE and EARLY in the SELECT list because the in-line view is evaluated first.

In-Line Views

Output

Destination	Average Delay	Maximum Delay	Number of Delays	Number of Early Arrivals	Probability of Delay
WAS	1	15	76	75	0.50
YYZ	2	14	36	24	0.60
DFW	3	20	38	23	0.62
ORD	3	19	51	41	0.55
LAX	5	27	82	41	0.67
LHR	6	30	39	19	0.67
CPH	6	26	16	11	0.59
FRA	6	34	14	12	0.54
CDG	9	39	21	5	0.81

58

Handling a Complex Query

What are the names of the supervisors for the crew on the flight to Copenhagen on March 4, 2000?

- Step 1: Identify the crew for the flight.
- Step 2: Find the states and job categories of the crew returned from the first query.
- Step 3: Find the employee numbers of the crew supervisors based on the states and job categories generated by the second query.
- Step 4: Find the names of the supervisors based on the employee numbers returned from the third query.

59

...

Because this query involves four tables,

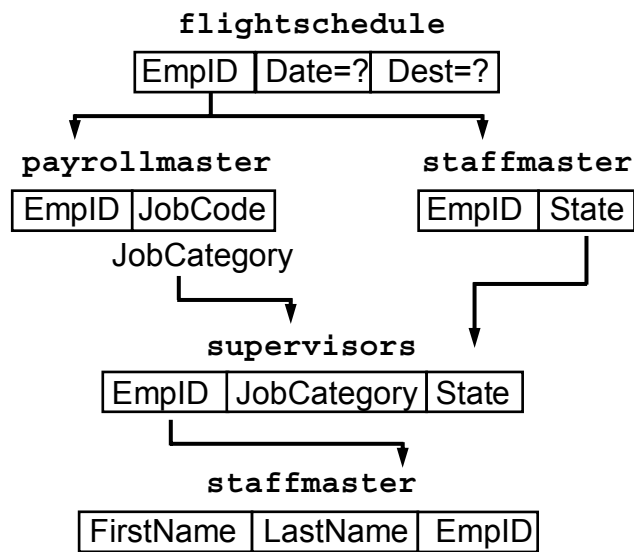
- **airline.flightschedule**
- **airline.staffmaster**
- **airline.payrollmaster**
- **airline.supervisors**

it may not be easy to code all at once. Split the query into small parts and test it each time that a new part is added.

The columns needed for this query are as follows:

- **EmpID**
- **FirstName**
- **LastName**
- **Date**
- **Destination**
- **JobCode**
- **JobCategory**
- **State**

Handling a Complex Query: Flow Diagram



Handling a Complex Query

Step 1: Identify the crew for the flight.

```
select EmpID
  from airline.flightschedule
 where Date='04mar2000'd
        and Destination='CPH';
```

61

c3s3d03

Handling a Complex Query

Step 1: Output

Emp ID
1556
1830
1124
1135
1437
1839

62

Handling a Complex Query

Step 2: Find the states and job categories of the crew returned from the first query.

```
select substr(JobCode,1,2) as JobCategory,
       State
  from airline.staffmaster as s,
       airline.payrollmaster as p
 where s.EmpID=p.EmpID and s.EmpID in
       (select EmpID
        from airline.flightschedule
        where Date='04mar2000'd
          and Destination='CPH');
```

63

c3s3d04

There is one supervisor for each state and job category.

Handling a Complex Query

Step 2: Output

JobCategory	State
FA	CT
FA	NY
NA	NY
PT	NY
PT	CT
FA	NY

64

Handling a Complex Query

Step 3: Find the employee numbers of the crew supervisors based on the states and job categories generated by the second query.

```
select EmpID
  from airline.supervisors as m,
       (select substr(JobCode,1,2) as JobCategory,
          State
         from airline.staffmaster as s,
              airline.payrollmaster as p
        where s.EmpID=p.EmpID and s.EmpID in
              (select EmpID
               from airline.flightschedule
               where Date='04mar2000'd and
                     Destination='CPH')) as c
 where m.JobCategory=c.JobCategory
        and m.State=c.State;
```

65

c3s3d05



You can assign an alias to an in-line view.

Handling a Complex Query

Step 3: Output

Supervisor Id
1431
1983
1352
1118
1106
1983

66



Note that ID number 1983 appears twice in the query result.

Handling a Complex Query

Step 4: Find the names of the supervisors.

```
select FirstName, LastName
  from airline.staffmaster where EmpID in
    (select EmpID
      from airline.supervisors as m,
      (select substr(JobCode,1,2) as
        JobCategory, State
      from airline.staffmaster as s,
      airline.payrollmaster as p
      where s.EmpID=p.EmpID and s.EmpID in
        (select EmpID
          from airline.flightschedule
          where Date='04mar2000'd and
              Destination='CPH')) as c
     where m.JobCategory=c.JobCategory
        and m.State=c.State);
```

67

c3s3d06

Handling a Complex Query

Step 4: Output

FirstName	LastName
SHARON	DEAN
ROGER	DENNIS
JASPER	MARSHBURN
SIMON	RIVERS
DEBORAH	YOUNG

68



Sharon Dean, ID 1983, appears once in this query result.

Handling a Complex Query

You can also solve this problem by using a multiway join.

```
select distinct e.FirstName, e.LastName
  from airline.flightschedule as a,
       airline.staffmaster as b,
       airline.payrollmaster as c,
       airline.supervisors as d,
       airline.staffmaster as e
 where a.Date='04mar2000'd and
       a.Destination='CPH' and
       a.EmpID=b.EmpID and
       a.EmpID=c.EmpID and
       d.JobCategory=substr(c.JobCode,1,2)
 and d.State=b.State and
       d.EmpID=e.EmpID;
```

69

c3s3d07

This code provides a more efficient solution to the query, but it is more difficult to build step-by-step.

You must have two copies of the **staffmaster** table: one to look up the states of the crew members and the other to look up the names of the supervisors. If you use a single copy of the table, it restricts the query to supervisors who were actually in the flight crew, if any.

Comparison with Traditional SAS Programs

Perform the same task using traditional SAS programming.

```
/* Find the crew for the flight. */
/* Program c3s3d08 */

proc sort data=airline.flightschedule (drop=flightnumber)
    out=crew (keep=empid);
    where destination='CPH' and date='04MAR2000'd;
    by empid;
run;

/* Find the State and job code for the crew. */

proc sort data=airline.payrollmaster (keep=empid jobcode)
    out=payroll;
    by empid;
run;

proc sort data=airline.staffmaster
    (keep=empid state firstname lastname)
    out=staff;
    by empid;
run;

data st_cat (keep=state jobcategory);
    merge crew (in=c)
          staff
          payroll;
    by empid;
    if c;
    jobcategory=substr(jobcode,1,2);
run;

/* Find the supervisor IDs. */

proc sort data=st_cat;
    by jobcategory state;
run;

proc sort data=airline.supervisors
    out=superv;
    by jobcategory state;
run;
```

(Continued on the next page.)


```

data super (keep=empid);
  merge st_cat(in=s)
        superv;
  by jobcategory state;
  if s;
run;

/* Find the names of the supervisors. */

proc sort data=super;
  by empid;
run;

data names(drop=empid);
  merge super (in=super)
        staff (keep=empid firstname lastname);
  by empid;
  if super;
run;

proc print data=names noobs uniform;
run;

```

Output

LastName	FirstName
MARSHBURN	JASPER
DENNIS	ROGER
RIVERS	SIMON
YOUNG	DEBORAH
DEAN	SHARON
DEAN	SHARON



The SQL query eliminated the duplicate names seen in this output.

In the example, the SQL query uses less CPU time, but more I/O operations than the non-SQL program (based on a mainframe benchmark in batch mode).

Choosing Between SQL Joins and DATA Step Merges

- DATA step merges are usually more efficient than SQL joins in combining small tables.
- SQL joins are usually more efficient than DATA step merges in combining large, unsorted tables.
- SQL joins are usually more efficient than DATA step merges in combining a large, indexed table with a small table.

70

A DATA step merge requires sorted data that calls for one or more SORT procedure steps. PROC SQL does not require sorted data.

Choosing Between SQL Joins and DATA Step Merges

- For ad hoc queries, select the method that you can code in the shortest time.
- For production jobs, experiment with different coding techniques and evaluate performance statistics.

71



Exercises

Submit a LIBNAME statement to assign the libref **airline** to the course SAS data library:

TSO: `libname airline '.sql.sasdata';`

Directory-based systems: `libname airline '.';`

1. Combining Data from Two Tables

Display the names of employees who have more than 20 years of service as of January 1, 2001. The **airline.staffmaster** table contains employee names, and the **airline.payrollmaster** table contains hire date information. Order the output by employee last name.

Employees with > 20 Years of Service as of 01JAN2001	
FirstName	LastName
JOSEPH	BAREFOOT
JUSTIN	BLAIR
DAVIS	CARAWAY
DONALD	CARTER
ROGER	DENNIS
KATRINA	FERNANDEZ
RAYMOND	HARTFORD
ANNE	KIRBY
ROY	LUFKIN
ALICE	MURPHY
JAMES	PEARSON
ROBERT	STEPHENSON
WAYNE	THOMPSON
ALAN	TUCKER
THERESA	UPDIKE
ELAINE	WARD
DARIUS	WELCH

2. Combining Data from Two Tables

Enhance the output from Exercise 1 by showing the number of years of service for each employee as of January 1, 2001.

Employees with 20 Years of Service as of 01JAN2001		
FirstName	LastName	YearsOf Service
JOSEPH	BAREFOOT	22
JUSTIN	BLAIR	21
DAVIS	CARAWAY	24
DONALD	CARTER	22
ROGER	DENNIS	22
KATRINA	FERNANDEZ	21
RAYMOND	HARTFORD	21
ANNE	KIRBY	22
ROY	LUFKIN	21
ALICE	MURPHY	22
JAMES	PEARSON	22
ROBERT	STEPHENSON	23
WAYNE	THOMPSON	23
ALAN	TUCKER	24
THERESA	UPDIKE	21
ELAINE	WARD	22
DARIUS	WELCH	21

3. Combining Data from Two Tables

Create a report that compares the number of passengers boarded with the capacity of the flight for all international flights. The **airline.internationalflights** table contains boarding information for **international** flights, and the **airline.marchflights** table contains capacity information for **all** flights. Order the output by flight number and date.

Hints:

- Use the PERCENT5. format for the column calculated as **Boarded/PassengerCapacity**.
- Ignore the Boarded column in **airline.marchflights**.

Partial Output

Capacity Figures for International Flights					
FlightNumber	Date	Boarded	Passenger Capacity	Percent	
132	.	98	178	55%	
132	01MAR2000	115	178	65%	
132	02MAR2000	106	178	60%	
132	03MAR2000	75	178	42%	
132	04MAR2000	117	178	66%	
132	05MAR2000	157	178	88%	
132	06MAR2000	150	178	84%	
132	07MAR2000	164	178	92%	
132	08MAR2000	104	178	58%	
132	09MAR2000	119	178	67%	
132	10MAR2000	98	178	55%	

4. Summarizing Data from Two Tables

Report the number of employees per job code for each state. Also display the average, maximum, and minimum salaries within the job code for each state. The **airline.staffmaster** table contains state data, and the **airline.payrollmaster** table contains job code and salary data. Order the report by state and job code.

Partial Output

Salary Statistics by State and Job Code						
State	Job Code	Total Employees	Average Salary	Maximum Salary	Minimum Salary	
CT	BCK	2	\$36,038.80	\$36,409.80	\$35,667.80	
CT	FA1	3	\$32,615.80	\$33,570.60	\$31,175.20	
CT	FA2	4	\$39,373.25	\$40,070.80	\$38,498.60	
CT	FA3	2	\$46,433.80	\$46,522.00	\$46,345.60	
CT	ME1	2	\$39,121.60	\$39,300.80	\$38,942.40	
CT	ME2	5	\$49,864.08	\$51,367.40	\$49,151.20	
CT	ME3	3	\$59,600.33	\$61,460.00	\$58,171.40	
CT	NA1	3	\$58,866.27	\$59,183.60	\$58,366.00	
CT	NA2	1	\$71,513.40	\$71,513.40	\$71,513.40	
CT	PT1	3	\$95,962.07	\$99,030.40	\$92,582.00	
CT	PT2	4	\$121,587.90	\$125,484.80	\$118,259.40	
CT	TA1	4	\$38,736.25	\$39,981.20	\$37,146.20	
CT	TA2	6	\$47,056.80	\$48,724.20	\$45,887.80	
CT	TA3	2	\$55,638.10	\$56,364.00	\$54,912.20	

5. Combining Data from Multiple Tables (Optional)

Create a flight and employee schedule that is ordered by flight number, date, last name, and first name. The data is in the tables **airline.staffmaster** (name information), **airline.flightschedule** (schedule information), and **airline.marchflights** (flight information).

Partial Output

Flight Schedule for Airline Employees							
Flt Num	Date	FirstName	LastName	Emp Num	Dep Time	Dest	
132	01MAR2000	JONATHAN	BOYCE	1739	15:35	YYZ	
132	01MAR2000	SHARON	DEAN	1983	15:35	YYZ	
132	01MAR2000	JAMES	NEWTON	1478	15:35	YYZ	
132	01MAR2000	JEREMY	RHODES	1111	15:35	YYZ	
132	01MAR2000	JONATHAN	SMART	1390	15:35	YYZ	
132	01MAR2000	DEBORAH	WOOD	1130	15:35	YYZ	
132	02MAR2000	MARSHALL	CAHILL	1574	15:35	YYZ	
132	02MAR2000	JACKSON	JOHNSON	1411	15:35	YYZ	
132	02MAR2000	LESLIE	JONES	1113	15:35	YYZ	
132	02MAR2000	JAMES	NEWTON	1478	15:35	YYZ	
132	02MAR2000	MICHAEL	PENNINGTON	1556	15:35	YYZ	
132	02MAR2000	JEREMY	RHODES	1111	15:35	YYZ	
132	03MAR2000	JONATHAN	BOYCE	1739	15:35	YYZ	
132	03MAR2000	DOROTHY	CARTER	1437	15:35	YYZ	
132	03MAR2000	JEREMY	RHODES	1111	15:35	YYZ	

6. Combining Data from Multiple Tables (Optional)

Display the flight attendants (job code of FA_) who are scheduled to fly to Copenhagen (CPH). Gather information from the tables **airline.staffmaster** (name information), **airline.payrollmaster** (job code information), and **airline.flightschedule** (schedule information). Order the report by employee last name.

Flight Attendants Scheduled for Copenhagen	
FirstName	LastName
BARBARA	ARTHUR
DOROTHY	CARTER
ANTHONY	COOPER
ALICIA	EATON
DIANA	FIELDS
MARIE	FLETCHER
LESLIE	JONES
KATHY	LAWRENCE
CAROL	PEARCE
EDITH	SANDERSON
JONATHAN	SMART
JENNY	UNDERWOOD
ANNA	VEGA
DIANE	WALTERS
DEBORAH	WOOD
DEBORAH	YOUNG

3.4 Set Operators

Objectives

- Use the SQL set operators.
- Compare the SQL set operators to traditional SAS programming tools.

74

Types of Set Operators

Set operators combine rows from two tables vertically.

The following are the four set operators:

- EXCEPT
- INTERSECT
- UNION
- OUTER UNION

General form of the set operators:

```
SELECT column <, column> ... FROM table1  
      set-operator <modifiers>  
SELECT column <, column> ... FROM table2 ;
```

75

Default Behavior of Set Operators

{ EXCEPT
INTERSECT
UNION

- Columns are matched by position and must be the same data type.
- Column names in the result set are determined by the first table.

- All columns from **both** tables are selected.

OUTER UNION

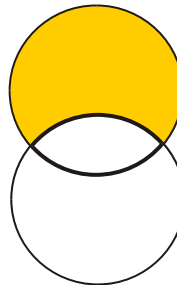
76

...

Types of Set Operators

EXCEPT

- Unique rows from the first table that are not found in the second table are selected.

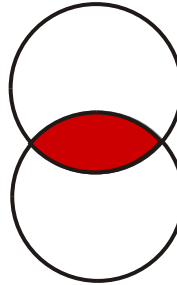


77

Types of Set Operators

INTERSECT

- Common unique rows from both tables are selected.




78

Types of Set Operators

UNION

- All unique rows from both tables are selected with columns overlaid.

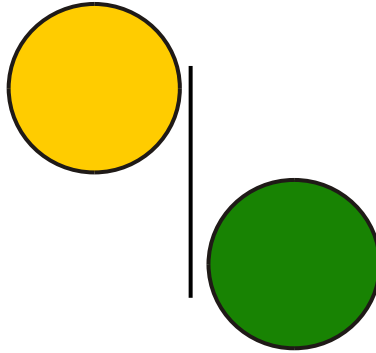


79

Types of Set Operators

OUTER UNION

- All rows from both tables, unique as well as non-unique, are selected.
- Columns are not overlaid.



80

Modifiers

You can use the following two keywords to modify the behavior of set operators:

- ALL
- CORRESPONDING

81

Modifiers

The following are characteristics of the **ALL** keyword:

- does not remove duplicate rows, and so avoids an extra pass through the data. Use the ALL keyword for better performance when it is possible.
- is not allowed in connection with an OUTER UNION operator. (It is implicit.)

82

Use the ALL keyword when

- you do not care if there are duplicates
- duplicates are not possible; for example, there is a unique or primary key constraint on the column.

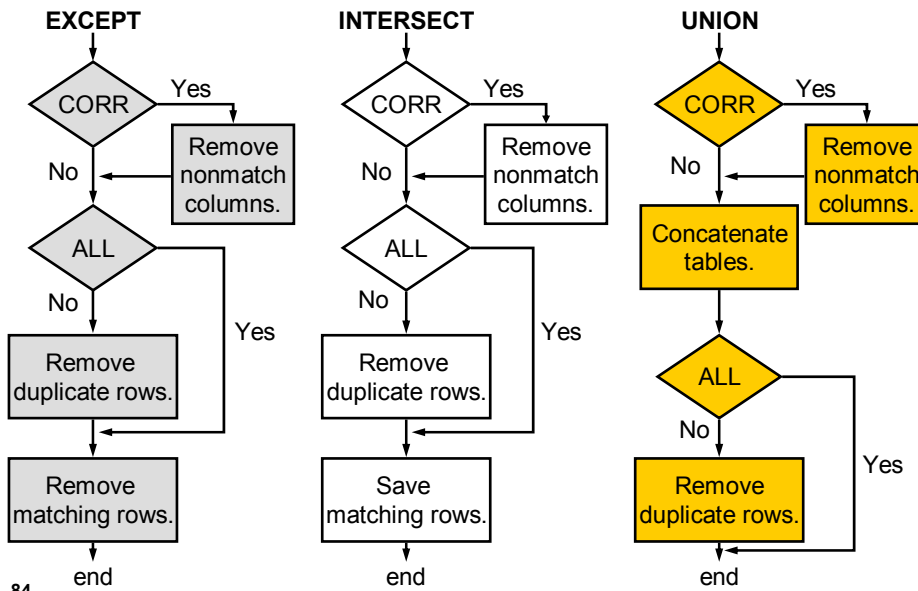
Modifiers

The following are characteristics of the **CORRESPONDING** keyword:

- overlays columns by name, instead of by position
- removes any columns not found in both tables when used in EXCEPT, INTERSECT, and UNION operations
- causes common columns to be overlaid when used in OUTER UNION operations
- can be abbreviated as CORR

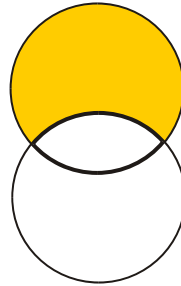
83

Set Operators and Modifiers: Flow Diagram



EXCEPT

- Unique rows from the first table that are not found in the second table are selected.



The EXCEPT Operator

Display the unique rows in table ONE that are not found in table TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except
select *
  from two;
```

86

c3s4d01

The EXCEPT Operator

The SQL processor removes duplicate rows within the tables.

Table ONE

X	A
1	a
→ 1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except
select *
  from two;
```

87

c3s4d01

The EXCEPT Operator

The SQL processor creates an intermediate result set by returning the rows that are found only in table ONE.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

①	②
1	a
1	b
2	c
4	e
6	g

```
select *
  from one
except
select *
  from two;
```

88

c3s4d01

The EXCEPT Operator

The column names are determined by table ONE in the final result set.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except
select *
  from two;
```

X	A
1	a
1	b
2	c
4	e
6	g

89

c3s4d01



Duplicate rows are omitted.

How can you include duplicate rows?

The EXCEPT Operator

Display the rows (duplicates included) that are found in table ONE but not in table TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except all
select *
  from two;
```

90

c3s4d02

The EXCEPT Operator

The SQL processor creates an intermediate result set by returning the rows that are found only in table ONE.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

①	②
1	a
1	a
1	b
2	c
4	e
6	g

```
select *
  from one
except all
select *
  from two;
```

91

c3s4d02

The EXCEPT Operator

The column names are determined by table ONE in the final result set.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except all
select *
  from two;
```

X	A
1	a
1	a
1	b
2	c
4	e
6	g

c3s4d02

92

The EXCEPT Operator

Display the unique rows that exist in table ONE and not in table TWO, based on same-named columns.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except corr
select *
  from two;
```

c3s4d03

93

The EXCEPT Operator

The SQL processor eliminates any columns not found in both tables.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except corr
select *
  from two;
```

94

c3s4d03

The EXCEPT Operator

Duplicate rows within each table are eliminated.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
except corr
select *
  from two;
```

95

c3s4d03

The EXCEPT Operator

The SQL processor creates an intermediate result set by returning the rows that are found only in table ONE.

Table ONE		Table TWO		
X	A	X	B	
1	a	1	x	↔
1	a	2	y	
1	b	3	z	
2	c	3	v	↔
3	v	5	w	
4	e			
6	g			

1
4
6


```

select *
  from one
except corr
select *
  from two;

```

96

c3s4d03

The EXCEPT Operator

Final result set.

Table ONE		Table TWO		
X	A	X	B	
1	a	1	x	
1	a	2	y	
1	b	3	z	
2	c	3	v	
3	v	5	w	
4	e			
6	g			


```

select *
  from one
except corr
select *
  from two;

```


X
4
6

97

c3s4d03

The EXCEPT Operator

airline.staffchanges and **airline.payrollchanges** contain information about

- current employees who have salary or job code changes
- new employees.

The new tables have the same layout as the **airline.staffmaster** and **airline.payrollmaster** tables.

98

The EXCEPT Operator

Example: Display the names of new employees.

```
select FirstName, LastName
   from airline.staffchanges
      except all
select FirstName, LastName
   from airline.staffmaster;
```

99

c3s4d04

The EXCEPT Operator

Output

FirstName	LastName
AMY	BRIDESTON
JIM	POWELL

100

The EXCEPT Operator

Example: How many employees have no changes in salary or job code?

```
select count(*) label='No. of Persons'
  from (select EmpID
        from airline.staffmaster
       except all
        select EmpID
        from airline.staffchanges);
```

101

c3s4d05

The EXCEPT Operator

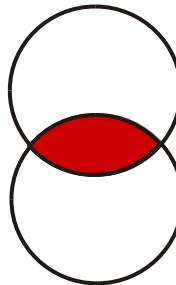
Output

No. of Persons
144

102

INTERSECT

- Common unique rows from both tables are selected.



103

The INTERSECT Operator

Display the unique rows common to table ONE and table TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect
select *
  from two;
```

104

c3s4d06

The INTERSECT Operator

The SQL processor removes duplicate rows within the tables.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g



Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect
select *
  from two;
```

105

c3s4d06

The INTERSECT Operator

The SQL processor creates an intermediate result set by returning the rows that are found in both tables.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

①	②
3	v

```
select *
  from one
intersect
select *
  from two;
```

106

c3s4d06

The INTERSECT Operator

The column names are determined by table ONE in the final result set.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect
select *
  from two;
```

X	A
3	v

107

c3s4d06

Would the addition of the ALL keyword have any effect in this example?

The INTERSECT Operator

Display the unique rows common to table ONE and table TWO, based on same-named columns.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect corr
select *
  from two;
```

108

c3s4d07

The INTERSECT Operator

Duplicate rows within each table are eliminated.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect corr
select *
  from two;
```

110

c3s4d07

The INTERSECT Operator

The SQL processor creates an intermediate result set by returning the rows that are found in both tables.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

1
1
2
3

```
select *
  from one
intersect corr
select *
  from two;
```

111

c3s4d07

The INTERSECT Operator

Final result set.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
intersect corr
select *
  from two;
```

X
1
2
3

112

c3s4d07

The INTERSECT Operator

Example: What are the names of the old employees who changed salary or job code?

```
select FirstName, LastName
  from airline.staffmaster
 intersect all
select FirstName, LastName
  from airline.staffchanges;
```

113

c3s4d08

The INTERSECT Operator

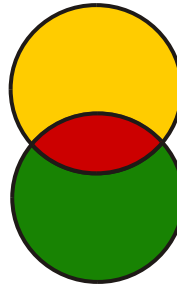
Output

FirstName	LastName
DIANE	WALTERS
KAREN	CARTER
NEIL	CHAPMAN
RAYMOND	SANDERS

114

UNION

- All unique rows from both tables are selected with columns overlaid.



115

The UNION Operator

Display the unique rows that table ONE and table TWO have all together.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union
select *
  from two;
```

116

c3s4d09

The UNION Operator

The SQL processor creates an intermediate result by concatenating and sorting ONE and TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union
select *
  from two;
```

①	②
1	a
1	a
1	b
1	x
2	c
2	y
3	v
3	v
3	z
4	e
5	w
6	g

c3s4d09

117

The UNION Operator

The SQL processor removes duplicate rows from the intermediate result.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union
select *
  from two;
```

①	②
1	a
1	a
1	b
1	x
2	c
2	y
3	v
3	v
3	z
4	e
5	w
6	g

c3s4d09

118

The UNION Operator

Final Result Set

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union
select *
  from two;
```

X	A
1	a
1	b
1	x
2	c
2	y
3	v
3	z
4	e
5	w
6	g

119

c3s4d09

Would the addition of the ALL keyword make any difference in this example?

Notice the overlay of columns A and B.

The UNION Operator

Display all of the unique rows of same-named columns in table ONE and table TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union corr
select *
  from two;
```

120

c3s4d10

The UNION Operator

The SQL processor creates an intermediate result by concatenating and sorting data from the first column.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union corr
select *
  from two;
```

X
1
1
1
1
2
2
3
3
3
4
5
6

121

c3s4d10

The UNION Operator

The SQL processor removes duplicate rows from the intermediate result to generate the final result.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
union corr
select *
  from two;
```

X
1
2
3
4
5
6

122

c3s4d10

The UNION Operator

Example: Display the total miles traveled, total bonus points earned, and total bonus points used by frequent flyers.

Output

Points and Miles Traveled by Frequent Flyers

Total Points Earned :	11,083,463
Total Points Used :	4,429,670
Total Miles Traveled:	10,477,963

123

The UNION Operator

```

title 'Points and Miles Traveled '
      'by Frequent Flyers';
select 'Total Points Earned :',
       sum(PointsEarned) format=comma12.
  from airline.frequentflyers
      union all
select 'Total Points Used   :',
       sum(PointsUsed) format=comma12.
  from airline.frequentflyers
      union all
select 'Total Miles Traveled:',
       sum(MilesTraveled) format=comma12.
  from airline.frequentflyers;

```


124

c3s4d11

In this example, the keyword ALL is used to control the sorting by specifying that PROC SQL make one pass only.

OUTER UNION

- All rows from both tables, unique as well as non-unique, are selected.
- Columns are not overlaid.



125

The OUTER UNION Operator

Display all data values from table ONE and table TWO.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
outer union
select *
  from two;
```

X	A	X	B
1	a	.	.
1	a	.	.
1	b	.	.
2	c	.	.
3	v	.	.
4	e	.	.
6	g	.	.
.	.	1	x
.	.	2	y
.	.	3	z
.	.	3	v
.	.	5	w

126

c3s4d12

With the OUTER UNION operator, the ALL keyword is implied.

The OUTER UNION Operator

Display all data values from table ONE and table TWO, but overlay common columns.

Table ONE

X	A
1	a
1	a
1	b
2	c
3	v
4	e
6	g

Table TWO

X	B
1	x
2	y
3	z
3	v
5	w

```
select *
  from one
outer union corr
select *
  from two;
```

X	A	B
1	a	
1	a	
1	b	
2	c	
3	v	
4	e	
6	g	
1		x
2		y
3		z
3		v
5		w

127

c3s4d13

Common columns can be overlaid using the CORR keyword.

The same result is obtained by using a DATA step with a SET statement.

The OUTER UNION Operator

Example: Display the employee numbers, job codes, and salaries of all mechanics.

```
select *
  from airline.mechanicslevel1
outer union corr
select *
  from airline.mechanicslevel2
outer union corr
select *
  from airline.mechanicslevel3;
```

128

c3s4d14

The OUTER UNION Operator

Partial Output

Employee Number	Job Code	Salary
1400	ME1	\$41,677
1403	ME1	\$39,301
1120	ME1	\$40,067
1121	ME1	\$40,757
1412	ME1	\$38,919
1200	ME1	\$38,942
1995	ME1	\$40,334
1418	ME1	\$39,207
1653	ME2	\$49,151
1782	ME2	\$49,483

129

SQL versus Traditional SAS Programming

The following programs produce the same report:

```
data allmechanics;
    set mechanicslevel1 mechanicslevel2;
run;
proc print data=allmechanics noobs;
run;
```

```
proc sql;
    select * from mechanicslevel1
        outer union corr
    select * from mechanicslevel2;
quit;
```

```
proc append base=mechanicslevel1
            data=mechanicslevel2;
run;
proc print data=one noobs;
run;
```

c3s4d15

130

Comparing Methods of Combining Tables Vertically

- PROC APPEND is the fastest method of performing a simple concatenation of two tables. The BASE= table is not completely read; only the DATA= table is completely read.
- When logical conditions are involved, you can choose either the DATA step or PROC SQL.
- SQL set operators generally require more computer resources, but the other operators are more convenient and flexible.

continued...

131

Comparing Methods of Combining Tables Vertically

- With the DATA step, you can process an unlimited number of tables at one time.
- With SQL set operators, you can work on only two tables at a time.
- If multiple DATA steps are required to perform the task, consider using PROC SQL.
- If you are unsure which method is best, use the techniques discussed in Chapter 5 to benchmark.

132



Although set operators work on only two tables at a time, you can chain multiple operators together, as in the mechanics example demonstrated earlier in this section.



Exercises

Submit a LIBNAME statement to assign the libref **airline** to the course SAS data library.

TSO: `libname airline '.sql.sasdata';`

Directory-based systems: `libname airline '.';`

7. Using Set Operators

Create a report that displays only the domestic (noninternational) flight numbers and destinations. **airline.marchflights** contains data on **all** flights. **airline.internationalflights** contains data on only the **international** flights.

FlightNumber	Destination
114	LAX
183	WAS
202	ORD
290	WAS
302	WAS
308	ORD
416	WAS
431	LAX
439	
439	LAX
523	ORD
829	WAS
872	LAX
921	DFW
982	DFW

8. Handling a Complex Query

airline.payrollchanges and **airline.staffchanges** contain data on employees with changes in job code or salary, as well as data on new employees. Create a report that displays information on new employees only, as shown below.

EmpID	FirstName	LastName	State	Job Code	DateOfHire
1447	AMY	BRIDESTON	NY	FA1	01NOV2000
1998	JIM	POWELL	NY	SCP	05NOV2000

To produce this report, break the problem into several steps.

- Find the **EmpID** values of the new employees. Data on long-standing employees is stored in **airline.staffmaster**. **airline.staffchanges** contains data on existing employees with status changes, plus new employees.
- In a separate query, display the **EmpID**, **FirstName**, **LastName**, and **State** columns from **airline.staffchanges**, with the **JobCode** and **DateOfHire** columns from **airline.payrollchanges**. (Six rows are displayed.)
- Combine the two queries in parts **a** and **b**, so that the results of **b** (displaying six employees) are subset to display only employees returned from **a**.

3.5 Chapter Summary

PROC SQL provides many ways to combine data from multiple tables. Join operations enable you to combine tables horizontally using a key value. You can use an inner join to retrieve rows from up to 32 tables. Conceptually, PROC SQL forms a Cartesian product (all possible combinations of rows) and then selects the rows that satisfy the WHERE expression(s).

Outer joins enable you to select matching rows as well as nonmatching rows. A left join selects matching rows plus nonmatching rows from the left table. A right join selects matching rows plus nonmatching rows from the right table. A full join selects matching rows plus nonmatching rows from both tables (similar to a DATA step merge). The COALESCE function is available to overlay columns in the output. You can assign an alias to a table to simplify qualified column references in the query.

An in-line view is created when the FROM clause contains a query expression instead of actual table names, and exists only during the execution of the query.

You can use set operators to combine two tables vertically, that is, one table displayed immediately above the other. The EXCEPT operator selects unique rows from the first table that are not found in the second table. The INTERSECT operator selects unique rows found in both tables. The UNION operator selects all unique rows from both tables. The OUTER UNION operator concatenates the two tables. You can use the ALL keyword to prevent duplicate rows from being eliminated. The CORRESPONDING keyword forces PROC SQL to compare columns by name rather than by position.

General form of an inner join:

```
SELECT column-1, column-2, ...
FROM table-1, table-2, ...
WHERE join-condition(s)
      <AND other subsetting conditions>
      <other clauses>;
```

General form of an outer join:

```
SELECT column-1, column-2, ...
FROM table-1
      LEFT|RIGHT|FULL JOIN
      table-2
      ON join-condition(s)
      <other clauses>;
```

General form of an inner join that uses the ON clause:

```
SELECT *
FROM table-1 INNER JOIN table-2
      ON table-1.x = table-2.x;
```

General form of a left join:

```
SELECT column-1, column-2, ...
FROM table-1 LEFT JOIN table-2
      ON expression;
```

General form of a right join:

```
SELECT column-1, column-2, ...  
FROM table-1 RIGHT JOIN table-2  
ON expression;
```

General form of a full join:

```
SELECT column-1, column-2, ...  
FROM table-1 FULL JOIN table-2  
ON expression;
```

General form of a set operation:

```
SELECT column-1, column-2, ...  
FROM table-1  
set-operator  
SELECT column-1, column-2, ...  
FROM table-2;
```

Set operators:

```
EXCEPT  
INTERSECT  
UNION  
OUTER UNION
```


3.6 Solutions to Exercises

1. Combining Data from Two Tables

```
proc sql;
title 'Employees with > 20 Years of Service';
title2 'as of 01JAN2001';
select FirstName,
       LastName
  from airline.staffmaster as s,
       airline.payrollmaster as p
 where s.EmpID=p.EmpID
       and int(('01jan2001'd - DateOfHire)/365.25) > 20
 order by LastName;
quit;
title;
```

2. Combining Data from Two Tables

```
proc sql;
title 'Employees with > 20 Years of Service';
title2 'as of 01JAN2001';
select FirstName,
       LastName,
       int(('01jan2001'd - DateOfHire) / 365.25)
         as YearsOfService
  from airline.staffmaster as s,
       airline.payrollmaster as p
 where s.EmpID=p.EmpID
       and calculated YearsOfService > 20
 order by LastName;
quit;
title;
```

3. Combining Data from Two Tables

```
proc sql;
title 'Capacity Figures for International Flights';
select i.FlightNumber,
       i.Date,
       i.Boarded,
       PassengerCapacity,
       i.Boarded / PassengerCapacity as Percent
         format=percent5.
  from airline.internationalflights as i,
       airline.marchflights as m
 where i.FlightNumber=m.FlightNumber
       and i.Date=m.Date
 order by 1, 2;
quit;
title;
```

4. Summarizing Data from Two Tables

```
proc sql;
title 'Salary Statistics by State and Job Code';
select State,
       JobCode,
       count(*) as TotalEmployees,
       avg(Salary) as AverageSalary
         format=dollar11.2,
       max(Salary) as MaximumSalary
         format=dollar11.2,
       min(Salary) as MinimumSalary
         format=dollar11.2
  from airline.staffmaster as s,
       airline.payrollmaster as p
 where s.EmpID=p.EmpID
 group by State, JobCode
 order by State, JobCode;
quit;
```

5. Combining Data from Multiple Tables (Optional)

```
proc sql;
title 'Flight Schedule for Airline Employees';
select f.FlightNumber,
       f.Date,
       FirstName format=$10.,
       LastName format=$10.,
       f.EmpID,
       DepartureTime as DepTime,
       f.Destination as Dest
  from airline.staffmaster as s,
       airline.flightschedule as f,
       airline.marchflights as m
 where s.EmpID=f.EmpID
       and f.FlightNumber=m.FlightNumber
       and f.Date=m.Date
 order by 1, 2, 4, 3;
quit;
title;
```

6. Combining Data from Multiple Tables (Optional)

```
proc sql;
title 'Flight Attendants Schedule for Copenhagen';
select distinct FirstName, LastName
  from airline.staffmaster as s,
       airline.payrollmaster as p,
       airline.flightschedule as f
 where s.EmpID=p.EmpID
       and s.EmpID=f.EmpID
       and JobCode like 'FA_'
       and Destination='CPH'
 order by LastName;
quit;
title;
```

7. Using Set Operators

```
proc sql;
select FlightNumber, Destination
  from airline.marchflights
except
select FlightNumber, Destination
  from airline.internationalflights;
quit;
```

8. Handling a Complex Query

a.

```
proc sql;
select EmpID
  from airline.staffchanges
except all
select EmpID
  from airline.staffmaster;
quit;
```

b.

```
proc sql;
select s.EmpID,
       FirstName,
       LastName,
       State,
       JobCode,
       DateOfHire
  from airline.staffchanges as s,
       airline.payrollchanges as p
 where s.EmpID = p.EmpID;
quit;
```

c.

```
proc sql;
select s.EmpID,
       FirstName,
       LastName,
       State,
       JobCode,
       DateOfHire
  from airline.staffchanges as s,
       airline.payrollchanges as p
 where s.EmpID = p.EmpID
       and s.EmpID in
       (select EmpID
        from airline.staffchanges
       except all
       select EmpID
        from airline.staffmaster);
quit;
```

Chapter 4 Creating and Modifying Tables and Views

4.1	Creating Tables	4-3
4.2	Creating Views	4-20
4.3	Creating Indexes.....	4-31
4.4	Maintaining Tables.....	4-45
4.5	Chapter Summary.....	4-60
4.6	Solutions to Exercises	4-63

4.1 Creating Tables

Objectives

- Define the column structure of a new table or use the column definitions from an existing table.
- Load data into a new table.
- Create a new table from the results of a query.

3

Creating Tables

Use the CREATE TABLE statement in three ways.

Creates an empty table
(Methods 1A and 1B).

CREATE TABLE *table-name* (*column-name*
type(length), <*column-name type(length)*>,...);
CREATE TABLE *table-name* **LIKE** *table-name*;

CREATE TABLE *table-name* **AS** *query-expression*;

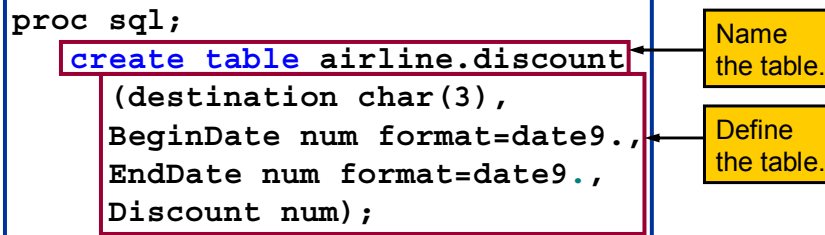
Populates table with a
query result (Method 2).

4

Creating Tables

Method 1A: Define the columns and fill in the data rows later.

```
proc sql;  
  create table airline.discount  
    (destination char(3),  
     BeginDate num format=date9.,  
     EndDate num format=date9.,  
     Discount num);
```



5

c4s1d01

Defining Columns

PROC SQL accepts

- types of CHARACTER or VARCHAR, but interprets both as SAS CHARACTER. Default length is 8 bytes.
- types of INTEGER, SMALLINT, DECIMAL, NUMERIC, FLOAT, REAL, and DOUBLE PRECISION, interpreting all as SAS NUMERIC with a length of 8 bytes.
- a type of DATE, interpreted as a SAS NUMERIC, with a length of 8 bytes and a DATE. informat and format.

6

Although SAS reads all of the above-mentioned data types, only CHARACTER and NUMERIC are used in SAS tables.

Defining Columns: More Examples

Method 1A:

```
proc sql;
  create table x
    (Name char(20),
     BirthDate date,
     Salary num format=comma10.2);
```

#	Variable	Type	Len	Format
1	Name	Char	20	
2	BirthDate	Num	8	DATE.
3	Salary	Num	8	COMMA10.2

7

c4s1d02

The table created above does not contain any rows. Use this method when the table you want to create is unlike any other existing table.

Defining Columns: More Examples

Method 1A:

```
proc sql;
  create table y
    (Dept varchar,
     Code integer label='Dept Code');
```

#	Variable	Type	Len	Label
1	Dept	Char	8	
2	Code	Num	8	Dept Code

8

c4s1d03

Defining Columns

Example: Create a table to store discounts for certain destinations and time periods in March. Define columns for destination, discount, and beginning and ending dates of the discount.

```
proc sql;
  create table discount
    (Destination char(3),
     BeginDate date label='BEGINS',
     EndDate date label='ENDS',
     Discount num);
```

Partial Log

```
NOTE:Table WORK.DISCOUNT created, with 0 rows
and 4 columns.
```

9

c4s1d04

Creating Tables

Method 1B: Copy a table. Use column definitions from another table and fill in the rows of data later.

```
proc sql;
  create table airline.delaycat
    (drop=DelayCategory DestinationType)
  like airline.flightdelays;
```

The columns
in this table
are copied
to the new table.

10

c4s1d05

...

Use Method 1A to create tables containing columns that do not already exist in other tables. In other words, you define your own columns.

Creating Tables

Method 2: Store a query result in a table that defines both columns and rows.

```
proc sql;
  create table airline.fa as
    select LastName, FirstName, Salary
      from airline.payrollmaster,
           airline.staffmaster
     where payrollmaster.EmpID
          =staffmaster.EmpID
          and JobCode contains 'FA' ;
  select *
    from airline.fa;
```

11

c4s1d06
...

This method is particularly helpful when you create subsets or supersets of tables.

Use of the CREATE TABLE statement shuts off the automatic report generation. Also, this is the only method of the three that **both** creates **and** populates a table at the same time.



Use this method when the table you want to create is similar or identical to another existing table.

Loading Data into a Table

Partial Output

The SAS System		
LastName	FirstName	Salary
ARTHUR	BARBARA	\$46,040
CAHILL	MARSHALL	\$40,001
CARTER	DOROTHY	\$46,346
COOPER	ANTHONY	\$45,104
DEAN	SHARON	\$46,787
DUNLAP	DONNA	\$40,443
EATON	ALICIA	\$38,902
FIELDS	DIANA	\$32,448
FLETCHER	MARIE	\$31,436
GOMEZ	ALAN	\$31,175

12

Loading Data into a Table

Method A: The SET Clause

```
INSERT INTO table-name
SET column-name=value,column-name=value,...;
```

Method B: The VALUES Clause

```
INSERT INTO table-name <(column list)>
VALUES (value,value,value, ...;)
```

Method C: A Query-expression

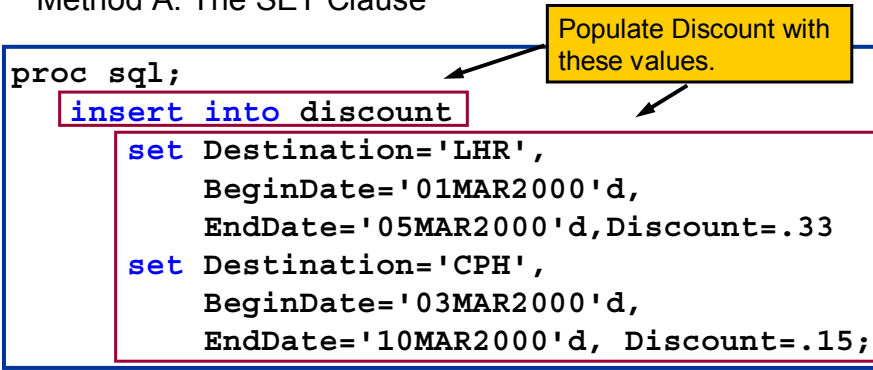
```
INSERT INTO table-name <(column list)>
SELECT columns FROM table-name
```

13

Loading Data into a Table

After the table is created, you can enter rows of data and use the INSERT statement with one of three methods.

Method A: The SET Clause



```
proc sql;
  insert into discount
    set Destination='LHR',
      BeginDate='01MAR2000'd,
      EndDate='05MAR2000'd, Discount=.33
    set Destination='CPH',
      BeginDate='03MAR2000'd,
      EndDate='10MAR2000'd, Discount=.15;
```

14

c4s1d07

You can nest a SELECT statement within a SET statement, as follows:

```
proc sql;
  insert into discount
    set Destination='LHR', BeginDate=(select max(Date)
    from airline.flightdelays);
```

Loading Data into a Table

Method B: The VALUES Clause

```
proc sql;
  insert into discount
    values ('LHR', '01MAR2000'd,
           '05MAR2000'd, .33)
    values ('CPH', '03MAR2000'd,
           '10MAR2000'd, .15);
```

15

c4s1d08

Loading Data into a Table

Method C: A Query-expression

```
proc sql;
  insert into discount(Destination,Discount)
  select Destination, Rate*.25
  from airline.fares
  where Type='special';
```

16

c4s1d09

Loading Data into a Table

Example: Create the **discount** table, insert four rows of data, and display the table.

```

proc sql;
  create table discount
    (Destination char(3),
     BeginDate date label='BEGINS',
     EndDate date label='ENDS',
     Discount num);
  insert into discount
  values('LHR','01MAR2000'd,'05MAR2000'd,.33)
  values('CPH','03MAR2000'd,'10MAR2000'd,.15)
  values('CDG','03MAR2000'd,'10MAR2000'd,.15)
  values('LHR','10MAR2000'd,'12MAR2000'd,.05);
  
```

17

c4s1d10
...

Loading Data into a Table

```

select *
  from discount;
  
```

Output

The SAS System			
Destination	BEGINS	ENDS	Discount
LHR	01MAR00	05MAR00	0.33
CPH	03MAR00	10MAR00	0.15
CDG	03MAR00	10MAR00	0.15
LHR	10MAR00	12MAR00	0.05

18

c4s1d10

You can use PROC PRINT in place of the final SELECT statement.

Integrity Constraints

- Integrity constraints are rules that table modifications must follow to guarantee validity of data.
- You can preserve the consistency and correctness of data by specifying integrity constraints for a SAS data file.
- SAS uses the integrity constraints to validate data when you insert or update the values of a variable for which you have defined integrity constraints.

19

Integrity Constraints

Integrity constraints

- are part of Version 8 of Base SAS software
- follow ANSI standards
- cannot be defined for views
- can be specified when a table is created or later when a table contains data.

20



Integrity constraints can be defined for historical versions of generation data sets using the DATASETS procedure but cannot be added using the SQL procedure.

Five Integrity Constraints

General:

- NOT NULL
- CHECK
- UNIQUE

Referential:

- PRIMARY KEY
- FOREIGN KEY

21

NOT NULL	means that data is required and ensures that corresponding columns have non-missing values in each row.
CHECK	specifies what values may be entered in a column. If a user attempts to enter data that violates this constraint, SAS rejects the value.
UNIQUE	ensures that every value in a column is unique. The same column can be defined as NULL, but only a single null value is allowed per UNIQUE column.
PRIMARY KEY	identifies the column as the table's primary key. Only unique values are permitted and the primary key cannot contain missing values.
FOREIGN KEY	links one or more columns in a table to a specific column in another table by matching a foreign key in one table with the primary key in another table. This parent/child relationship limits modifications made to both primary and foreign keys. The only acceptable values for a foreign key are values of the primary key or missing values.

Using PROC SQL to Create Integrity Constraints

General form of PROC SQL using integrity constraints:

```
PROC SQL;
  CREATE TABLE table
    (column-specification,...
    <constraint-specification,...>);
```

22

Using PROC SQL to Create Integrity Constraints

Example: Re-create the **discount** table with an integrity constraint to limit ticket discounting.

```
proc sql;
  create table discount
    (Destination char(3),
     BeginDate date label='BEGINS',
     EndDate date label='ENDS',
     Discount num,
     CONSTRAINT ok_discount check
       (Discount le .5));
```

23

c4s1d11

Using PROC SQL to Create Integrity Constraints

Example: Insert two rows using the default
UNDO_POLICY option (required).

```
proc sql;
  insert into discount
    values('CDG','03MAR2000'd,'10MAR2000'd,.15)
    values('LHR','10MAR2000'd,'12MAR2000'd,.55);
```

Stockholders might not
tolerate excessive airline
generosity!

24

c4s1d12

Using PROC SQL to Create Integrity Constraints

Partial Log

```
proc sql;
  insert into discount
    values('CDG','03MAR2000'd,'10MAR2000'd,.15)
    values('LHR','10MAR2000'd,'12MAR2000'd,.55);
ERROR: Add/Update failed for data set WORK.DISCOUNT because data
value(s) do not comply with integrity constraint ok_discount.
NOTE: This insert failed while attempting to add data from VALUES
clause 2 to the data set.
NOTE: Deleting the successful inserts before error noted above to
restore table to a consistent state.
```

0 rows inserted.

25

Rollbacks

If an INSERT or UPDATE statement experiences an error while it processes the statement, then the inserts or updates that were completed up to the point of the error by that statement can be undone by use of the UNDO_POLICY option.

Rollbacks with the UNDO_POLICY Option

- **UNDO_POLICY=REQUIRED** (the default)
undoes all inserts or updates that have been done to the point of the error. Sometimes the UNDO operation cannot be done reliably.
- **UNDO_POLICY=NONE**
prevents any updates or inserts from violating a constraint.
- **UNDO_POLICY=OPTIONAL**
reverses any updates or inserts that it can reverse reliably.

27



The ROLLBACK statement, although an ANSI standard, is not currently supported in the SQL procedure.

UNDO_POLICY=REQUIRED

PROC SQL performs UNDO processing for INSERT and UPDATE statements.

If the UNDO operation cannot be done reliably, PROC SQL does not execute the statement and issues an ERROR message.

UNDO cannot be attempted reliably in the following situations:

1. A SAS data set opened with CNTLLEV=RECORD can enable other users to update newly inserted records. An error during the insert deletes the record that the other user inserted.
2. A SAS/ACCESS view is not able to rollback the changes made by this statement without rolling back other changes at the same time.

Default: UNDO_POLICY=REQUIRED

UNDO_POLICY=NONE

PROC SQL skips records that cannot be inserted or updated, and writes to the SAS log a warning message similar to that written by PROC APPEND.

UNDO_POLICY=OPTIONAL

PROC SQL performs UNDO processing if it can be done reliably. If the UNDO cannot be done reliably, then no UNDO processing is attempted.

This option is a combination of the first two. If UNDO can be done reliably, then it is done. PROC SQL proceeds as if UNDO_POLICY=REQUIRED is in effect. Otherwise, it proceeds as if UNDO_POLICY=NONE was specified.

Using PROC SQL to Create Integrity Constraints

Example: Insert two rows using UNDO_POLICY=NONE.

```
proc sql undo_policy=none ;
  insert into discount
    values('CDG','03MAR2000'd,'10MAR2000'd,.15)
    values('LHR','10MAR2000'd,'12MAR2000'd,.55);
```

28

c4s1d13



An alternative is to create constraints using the DATASETS procedure.

Using PROC SQL to Create Integrity Constraints

Partial Log

```
WARNING: The SQL option UNDO_POLICY=REQUIRED is not in effect.
If an error is detected when processing this INSERT statement,
that error will not cause the entire statement to fail.
ERROR: Add/Update failed for data set WORK.DISCOUNT because data
value(s) do not comply with integrity constraint ok_discount.
NOTE: This insert failed while attempting to add data from VALUES
clause 2 to the data set.
NOTE: 2 rows were inserted into WORK.DISCOUNT. Of these 1 row was
rejected as an ERROR, leaving 1 row that was inserted successfully
```

1 of 2 rows inserted successfully.

29

Documenting Table and View Definitions and Integrity Constraints

The DESCRIBE statement displays the definition of the view or CREATE TABLE statement of a table.

General form of the DESCRIBE statement:

```
PROC SQL;  
  DESCRIBE TABLE table-name<,<table-name>>...;  
  DESCRIBE VIEW proc-sql-view <,<proc-sql-view>>...;  
  DESCRIBE TABLE CONSTRAINTS table-name  
    <,<table-name>> ...;
```

30

The DESCRIBE TABLE statement (without the CONSTRAINTS keyword) writes a CREATE TABLE statement to the SAS log for the specified table regardless of how the table was originally created (for example, with a DATA step).

If the table contains an index, CREATE INDEX statements for those indexes are also written to the SAS log. (A discussion of indexes is in Section 4.3.)

Documenting Table Definitions and Integrity Constraints

Example: Show the constraints for the **discount** table.

```
proc sql;  
  describe table discount;
```

31

c4s1d14

Documenting Table Definitions and Integrity Constraints

NOTE: SQL table WORK.DISCOUNT was created like:

```
create table WORK.DISCOUNT( bufsize=4096 )
(
  Destination char(3),
  BeginDate num format=DATE. informat=DATE. label='BEGINS',
  EndDate num format=DATE. informat=DATE. label='ENDS',
  Discount num
);
```

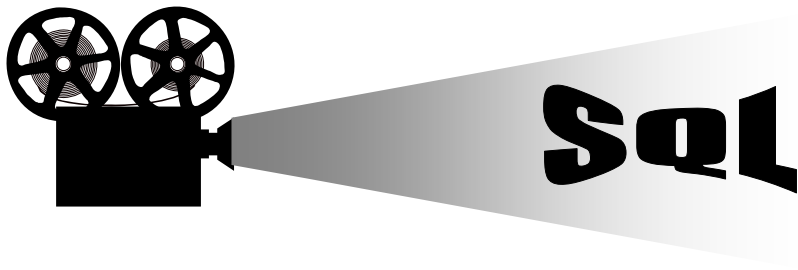
-----Alphabetic List of Integrity Constraints-----

#	Integrity Constraint	Type	Where Clause
1	ok_discount	Check	discount<=0.5

4.2 Creating Views

Objectives

- Create an SQL view and understand how it is best used.



Creating a View

A PROC SQL view

- is a stored query. It contains no rows of data.
- can be used in SAS programs in place of an actual SAS data file.
- can be derived from one or more tables, PROC SQL views, DATA step views, or SAS/ACCESS views.
- extracts underlying data when used, thus accessing the most current data.

35

Views are not separate copies of the data and are referred to as *virtual tables* because they do not exist as independent entities as do real tables. It may be helpful to think of a view as a movable frame or window through which you can see the data.

Thus, when the view is referenced by a SAS procedure or in a DATA step, it is executed, and conceptually, an internal table is built. PROC SQL processes this internal table as if it were any other table.

Creating a View

General form of the CREATE VIEW statement:

```
CREATE VIEW view-name AS  
query-expression;
```

Creating a View

Example: Create a view containing personal information for flight attendants. Have the view always return the employee's age as of the current date.

```
proc sql;
  create view airline.faview as
    select LastName, FirstName, Gender,
           int((today() - DateOfBirth) / 365.25)
             as Age,
           substr(JobCode, 3, 1) as Level,
           Salary
    from airline.payrollmaster,
         airline.staffmaster
   where JobCode contains 'FA' and
         staffmaster.EmpID=
         payrollmaster.EmpID;
```

37

c4s2d01

In this example, the view **airline.faview** creates a virtual table from the accompanying SELECT statement. Although the underlying tables, **airline.payrollmaster** and **airline.staffmaster**, can change, the instructions, which comprise the view, stay constant. Further, when this PROC SQL step is executed, SAS does not actually execute the SELECT statement following the AS keyword, but instead partially compiles and stores the SELECT statement in a data file with a member type of VIEW.

If the above example is modified to a CREATE TABLE statement and the alias **Age** is omitted, SAS creates a sequentially suffixed variable, starting with **_TEMA001**. The librefs for the tables in the FROM clause are optional in this case. It is assumed that the contributing tables are stored in the same library as the view itself, unless otherwise specified.

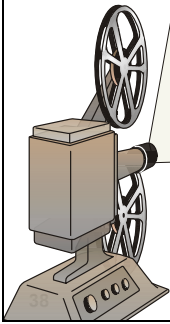
Using a View

```
proc sql;
  select *
  from airline.faview;
```

Your view
to the data

Partial Output

The SAS System					
LastName	FirstName	Gender	Age	Level	Salary
ARTHUR	BARBARA	F	36	3	\$46,040
CAHILL	MARSHALL	M	41	2	\$40,001
CARTER	DOROTHY	F	41	3	\$46,346
COOPER	ANTHONY	M	42	3	\$45,104
DEAN	SHARON	F	39	3	\$46,787
DUNLAP	DONNA	F	33	2	\$40,443
EATON	ALICIA	F	40	2	\$38,902



c4s2d02

Using a View

Example: Calculate the flight attendants' mean age, by level, using the **airline.faview** view.

```
proc tabulate data=airline.faview;
  class Level;
  var Age;
  table Level*Age*mean;
run;
```

Your view

39

c4s2d03

In both of the above examples, it only appears that the PROC SQL view, **airline.faview**, is a table because the view name itself is used in the same way as a SAS table name. However, it is **not** a table, but a stored query-expression only. Both tables and views are considered SAS data sets.

Using a View

PROC TABULATE Output

Level		
1	2	3
Age	Age	Age
Mean	Mean	Mean
39.82	43.81	45.43

40

Administering Views

Example: Write the view definition for **airline.fa** to the SAS log.

```
proc sql;
  describe view airline.faview;
```

NOTE: SQL view AIRLINE.FAVIEW is defined as:

```
select LastName, FirstName, Gender,
       INT((TODAY()-DateOfBirth)/365.25) as Age,
       SUBSTR(JobCode, 3, 1) as Level, Salary
from AIRLINE.PAYROLLMASTER, AIRLINE.STAFFMASTER
where JobCode contains 'FA' and
      (staffmaster.EmpID=payrollmaster.EmpID);
```

41

c4s2d04

Why Use Views?

You can

- access the most current data in changing tables, DATA step views, or SAS/ACCESS views
- pull together data from multiple database tables or even different databases
- simplify complex query-expressions and prevent users from altering code
- avoid storing a SAS copy of a large table.

42

Administering Views: Some General Guidelines

- Avoid the ORDER BY clause in a view definition. Otherwise, the data must be sorted each time the view is referenced.
- If the same data is used many times in one program, create a table rather than a view.
- Avoid specifying two-level names in the FROM clause when you create a permanent view that resides in the same library as the contributing table(s).

43

Administering Views

Example:

```
proc sql;  
  create view sasdata.master as  
    select *  
      from sasdata.payrollmaster;
```

It is better to omit the libref.

44

c4s2d05

Administering Views: Omitting the Libref

Example:

```
proc sql;  
  create view sasdata.master as  
    select *  
      from payrollmaster;
```

This looks like
work.payrollmaster,...

... but is **in reality**
sasdata.payrollmaster.

45

c4s2d06

Administering Views: Why Is It Better to Omit the Libref?

```
/* contents of master view */
select *
  from sasdata.payrollmaster;
```

Step 1: Assign libref.

```
libname airline 'SAS-data-library';
```

Step 2: Execute instructions in view.

```
proc print data=airline.master;
run;
```

Step 3: See
what went
wrong!

```
ERROR: Libname SASDATA is not
assigned.
ERROR: SQL View AIRLINE.MASTER
could not be processed.
```

'SAS-data-library'

libname=airline	
Table Name	Table Type
master	view
payrollmaster	data

46

c4s2d07
...

Administering Views: Why Is It Better to Omit the Libref?

```
select *
  from airline.payrollmaster
```

Step 1: Assign libref.

```
libname airline 'SAS-data-library';
```

Step 2: Execute instructions in view.

```
proc print data=airline.master;
run;
```

Step 3: Relax and be happy.

'SAS-data-library'

libname=airline	
Table Name	Table Type
master	view
payrollmaster	data

47

c4s2d08
...

Creating Views

An alternative: Embed the LIBNAME statement within a USING clause.

```
CREATE VIEW proc-sql-view AS query-expression
  <USING statement<, libname-clause> ... > ;
```

This enables you to store a SAS libref in the view and does not conflict with an identically named libref in the SAS session.

48

Administering Views: Using the Embedded LIBNAME Statement

```
libname sasdata 'SAS-data-library-one';
libname airline 'SAS-data-library-two';

proc sql;
  create view sasdata.journeymen as
  select *
    from airline.payrollmaster
   where JobCode like ' 2'
   using libname airline 'SAS-data-library-three';
quit;
proc print data = sasdata.journeymen ;
run;
```

1) While the view `sasdata.journeymen` is executing ...

2) ... the libref `airline` becomes active ...

3) ... overriding any earlier assignment for the duration of the view's execution.

4) After view executes, original libref assignment (3) is re-established and embedded assignment (2) is cleared.

49

c4s2d09

...

Administering Views

Example: Create a view to maintain table security.

fa1.info

fa2.info

fa3.info

```
create view manager.info as
  select *
    from fa1.info
  outer union corr
  select *
    from fa2.info
  outer union corr
  select *
    from fa3.info;
```

The **manager** data library can be assigned access privileges at the operating system level. The access privilege prevents non-managerial flight attendants from reading the library, but permits managers (who are authorized to access all SAS data libraries) to view all information.

4.3 Creating Indexes

Objectives

- Create an index on a table.
- Understand how an index is best used.

52

Creating Indexes

An *index* is an auxiliary data structure that specifies the location of rows based on the values of one or more **key** columns.

The SQL procedure can utilize an available index to optimize subsetting or joining tasks.

53

The index is a structure that boosts program performance by serving as a logical pointer to a physical location of a given value.

Creating Indexes

Indexed SAS Data Set				Index File			
Row	EmpID	Gender	JobCode	Key Column=JobCode			
1	1001	F	FA1	Key	Location		
2	1012	F	FA3	Value	Page(row,row...)		
3	1015	M	FA2	FA1	1(1,4, ...)	2(...)	...
...	FA2	1(3,6, ...)	2(...)	...
11	1104	M	FA3	FA3	1(2,11,...)	2(...)	...
...	Data Processed			
...	ROW	EmpID	Gender	JobCode
...	2	1012	F	FA3
...	11	1104	M	FA3
...
...

DATA or PROC Step

where JobCode='FA3';

54

Creating Indexes: Overview

Indexes provide fast access to small subsets of data...

```
proc sql;
  select *
    from airline.payrollmaster
   where JobCode='NA1';
```

One of many values
of the variable JobCode

55

c4s3d01



A small subset is $\leq 15\%$.

Creating Indexes: Overview

... and also enhance join performance.

```
proc sql;  
  select *  
    from airline.payrollmaster,  
         airline.staffmaster  
   where staffmaster.EmpID=  
         payrollmaster.EmpID;
```

When you subset data, you can select an index to optimize not only a WHERE clause with an equals comparison, but also a WHERE clause with the TRIM or SUBSTR function or the CONTAINS or LIKE operator.

Index Terminology

Two types of indexes are

- simple
based on values of only one column
- composite
based on values of more than one column
concatenated to form a single value, for example,
Date and **FlightNumber**.

57



Index naming rules are the same as the rules for other SAS data files. Start with a letter or underscore, and continue with a combination of letters, characters, or numbers, with a 32-character maximum.

Index Terminology

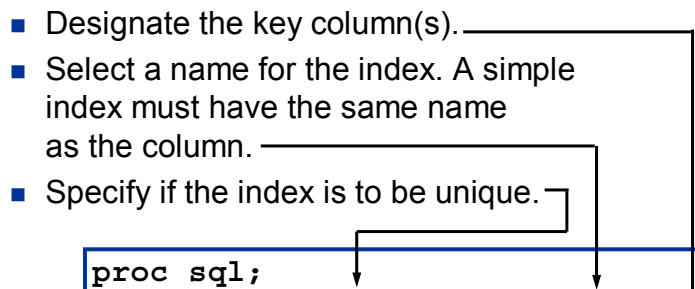
A table can have the following:

- multiple simple and composite indexes
- character and numeric key columns

58

Creating an Index

- Designate the key column(s).
- Select a name for the index. A simple index must have the same name as the column.
- Specify if the index is to be unique.



```
proc sql;  
  create unique index EmpID  
  on airline.payrollmaster (EmpID);
```

Creating an Index

General form of the CREATE INDEX statement:

```
CREATE <UNIQUE> INDEX index-name  
ON table-name(column-name, column-name);
```

Precede the INDEX keyword with the UNIQUE keyword to define a unique index.

60

Use of the optional UNIQUE keyword ensures that values in the row are unique. If a table contains multiple occurrences of the same value, the UNIQUE keyword is not accepted and the index is not defined on that column. Similarly, if you already have a uniquely defined index on a column and attempt to add a duplicate value to the table, the row is not inserted. For example, an index can be created on a column containing driver license or social security numbers, and thereby can prevent duplicate additions.

Additional notes:

1. Indexes can be based on either a character or numeric variable.
2. You do not want to create two indexes on the same variable.
3. You can achieve improved index performance if you create the index on a pre-sorted data set.
4. A composite index cannot have the same name as a variable.

Creating an Index: Examples

This simple index is based on **EmpID** and allows no duplicate ID numbers in the table.

```
proc sql;  
  create unique index EmpID  
  on airline.payrollmaster(EmpID);
```

Names must match for a simple index.

61

c4s3d03
...

Creating an Index: Examples

The composite index named DAILY is based on **FlightNumber** and **Date**.

```
proc sql;  
  create unique index daily  
  on airline.marchflights(FlightNumber,Date);
```

Cannot be a column name.

62

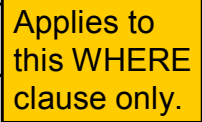
c4s3d04

To determine if an index is used, specify the SAS system option MSGLEVEL=I. A note appears in the SAS log when an index is selected for processing.

Indexing and Performance

Example: An index was created for the **JobCode** column of **airline.payrollmaster**.
Use the MSGLEVEL=I system option to determine which queries used the index.

```
options msglevel = i;  
proc sql;  
  select *  
    from airline.payrollmaster  
   where JobCode = 'NA1';  
INFO:Index JobCode selected for WHERE clause optimization.  
  
  select *  
    from airline.payrollmaster  
   where Salary gt 100000;
```



Controlling Index Usage in a WHERE Expression

Two data set options control the use of indexes:

- `IDXWHERE=` YES | NO
- `IDXNAME=<name>`.



64

When the `IDXTYPE=` option is

- | | |
|-----|--|
| YES | SAS uses the best available index to process the WHERE expression, even if SAS estimates that sequential processing is faster. |
| NO | SAS processes the data sequentially even if SAS estimates that processing with an index is better. |

When the `IDXNAME=` option is

- | | |
|---------------------------|---|
| <code><name></code> | SAS uses the named index regardless of performance estimates. |
|---------------------------|---|

If you do not use the `IDXTYPE=` option, SAS chooses whether to use an index. You can use either the `IDXTYPE=` or the `IDXNAME=` data set option, but not both.

Indexing and Performance

Suggested guidelines for using indexes:

- Keep the number of indexes to a minimum to reduce disk storage and update costs.
- Do not create an index for small tables; sequential access is faster on small tables.
- Do not create an index based on columns with a small number of distinct values, for example, Male and Female.
- An index performs best when it retrieves a relatively small number of rows, that is, <15%.

65

Indexing and Performance: Tradeoffs

Benefits

- Fast access to a small subset of data (<15%).
- Equijoins can be performed without internal sorts.
- Can enforce uniqueness.
- BY group processing without sorting.

Costs

- Extra CPU cycles and I/O operations to create an index.
- Extra disk space to store the index file.
- Extra memory to load index pages and code for use.
- Extra CPU cycles and I/O operations to maintain the index.

66



Exercises

Submit a LIBNAME statement to assign the libref **airline** to the course SAS data library.
(TSO only: DISP=SHR)

TSO: `libname airline '.sql.sasdata';`

Directory-based systems: `libname airline '.';`

1. Creating a Table

A frequent flyer earns points for each mile traveled with the airline. After accumulating a certain number of points, the frequent flyer is eligible for an award. You can claim a better award with more accumulated points.

- a.** Create a temporary table named **awards** to store award data. The table's columns must have the following attributes:

Name	Type	Length	Format	Label
ptsreqd	NUM	8		Points Required
rank	NUM	8	3.	
award	CHAR	25		

- b.** Load the following data into the table:

ptsreqd	rank	award
2000	1	free night in hotel
10000	2	50% discount on flight
20000	3	free domestic flight
40000	4	free international flight

- c.** Display the new table.

- d. The **airline.frequentflyers** table contains the number of points each frequent flyer earned (**PointsEarned**) and used (**PointsUsed**). Determine all appropriate awards for each frequent flyer based on the number of remaining points for each frequent flyer. An individual can receive multiple awards. Award levels are found in the new **awards** table. Process only the frequent flyers who live in Arizona (**STATE='AZ'**). Order the report by **FFID**.

Output

Awards Available to AZ Frequent Flyers				
FFID	Name	Available Points	Award	
WD0227	FOSTER, GERALD	29079	50% discount on flight	
WD0227	FOSTER, GERALD	29079	free domestic flight	
WD0227	FOSTER, GERALD	29079	free night in hotel	
WD0646	BOSTIC, MARIE	64544	50% discount on flight	
WD0646	BOSTIC, MARIE	64544	free domestic flight	
WD0646	BOSTIC, MARIE	64544	free international flight	
WD0646	BOSTIC, MARIE	64544	free night in hotel	
WD3022	CAHILL, LEONARD	46386	50% discount on flight	
WD3022	CAHILL, LEONARD	46386	free night in hotel	
WD3022	CAHILL, LEONARD	46386	free international flight	
WD3022	CAHILL, LEONARD	46386	free domestic flight	
WD4382	O'NEAL, ALICE	35047	50% discount on flight	
WD4382	O'NEAL, ALICE	35047	free domestic flight	
WD4382	O'NEAL, ALICE	35047	free night in hotel	
WD6061	RODRIGUEZ, MARIA	20642	free night in hotel	
WD6061	RODRIGUEZ, MARIA	20642	50% discount on flight	
WD6061	RODRIGUEZ, MARIA	20642	free domestic flight	
WD6080	SMART, JONATHAN	16266	free night in hotel	
WD6080	SMART, JONATHAN	16266	50% discount on flight	
WD7208	LONG, CASEY	19443	free night in hotel	
WD7208	LONG, CASEY	19443	50% discount on flight	
WD8375	COOPER, ANTHONY	5507	free night in hotel	
WD9829	COOK, JENNIFER	4401	free night in hotel	

- e. (Optional) Determine which frequent flyers are not eligible for any award. Order the report by **FFID**. Include all states.

Output

Frequent Flyers Ineligible for Awards		
FFID	Name	Available Points
WD0023	JACKSON, LAURA	-5
WD0231	GORDON, ANNE	-13054
WD0632	BROWN, JASON	-19367
WD1218	GRAHAM, MARY	-441
WD1637	NELSON, FELICIA	-6047
WD1700	WOOD, ALAN	-12836
WD1883	PENNINGTON, MICHAEL	-3957
WD2118	JOHNSON, ANTHONY	609
WD2741	EDGERTON, WAYNE	-29012
WD3129	FLOWERS, ANNETTE	-17635
WD3521	FIELDS, DIANA	-6151
WD4065	DONALDSON, KAREN	-6733
WD4781	HUNTER, CLYDE	1931
WD5020	BOYCE, RANDALL	1922

2. Creating a View

- a. Create a temporary view named **vsched** that extracts schedule information for airline employees. **vsched** must join data from the tables **airline.staffmaster** and **airline.flightschedule**. The view must include the date, flight number, and destination (in **airline.flightschedule**), and the name and ID number of each crew member assigned to that flight (in **airline.staffmaster**). Display the view and order the report by date, flight number, and employee last name.

Partial Output

View VSCHED					
Date	FlightNumber	Destination	FirstName	LastName	EmpID
01MAR2000	132	YYZ	JONATHAN	BOYCE	1739
01MAR2000	132	YYZ	SHARON	DEAN	1983
01MAR2000	132	YYZ	JAMES	NEWTON	1478
01MAR2000	132	YYZ	JEREMY	RHODES	1111
01MAR2000	132	YYZ	JONATHAN	SMART	1390
01MAR2000	132	YYZ	DEBORAH	WOOD	1130
01MAR2000	182	YYZ	FRANKLIN	CASTON	1269
01MAR2000	182	YYZ	ROGER	DENNIS	1118
01MAR2000	182	YYZ	ALAN	GOMEZ	1094
01MAR2000	182	YYZ	ALICE	MURPHY	1115
01MAR2000	182	YYZ	RANDALL	VENTER	1076
01MAR2000	182	YYZ	JOANN	YOUNG	1122

- b. Use the **vsched** view to display the schedule of Deborah Young (**EmpID= '1431'**). Order the report by date and flight.

Schedule for Deborah Young					
Date	FlightNumber	Destination	FirstName	LastName	EmpID
01MAR2000	387	CPH	DEBORAH	YOUNG	1431
03MAR2000	622	FRA	DEBORAH	YOUNG	1431
04MAR2000	821	LHR	DEBORAH	YOUNG	1431
05MAR2000	132	YYZ	DEBORAH	YOUNG	1431
07MAR2000	821	LHR	DEBORAH	YOUNG	1431

- c. (Optional) Use the view **vsched** and the table **airline.flightdelays** to determine how many delayed flights (**Delay>0**) each crew member was on. Order the report by employee first name and last name.

Partial Output

Number of Delayed Flights Experienced by Each Crew Member		
FirstName	LastName	count
ADAM	STEPHENSON	7
AGNES	WELLS	6
ALAN	GOMEZ	2
ALICE	MURPHY	4
ALICIA	EATON	3
ALVIN	GRAHAM	4
ANNA	VEGA	3
ANNE	PARKER	2
ANTHONY	COOPER	3
BARBARA	ARTHUR	1
CAROL	PEARCE	3
CASEY	RICHARDS	1
CHARLES	HARRIS	4
CHRISTINE	BRADY	2
CLYDE	HUNTER	1
DANIEL	GRANT	5
DEBORAH	WOOD	3

4.4 Maintaining Tables

Objectives

- Update or delete data values in an existing table.
- Add, drop, or alter the attributes of columns in a table.
- Delete tables, views, and indexes.

69

Maintaining Tables: Overview

You can use PROC SQL to do the following:

- modify values in a table or view
- add rows to a table or view
- delete rows from a table or view
- alter column attributes of a table
- add new columns to a table
- drop columns from a table
- delete an entire table, SQL view, or index

70

Updating Data Values

Use the UPDATE statement to modify column values in existing rows of a table or SAS/ACCESS view.

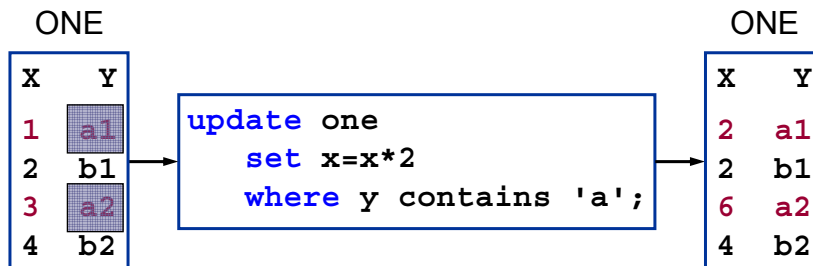
General form of the UPDATE statement:

```
UPDATE table-name
  SET column-name=expression,
      column-name=expression,...
  WHERE expression;
```

Careful! If you omit the WHERE expression, all rows are updated.

71

Updating Data Values



72

c4s4d01

Updating Data Values

Example: Give all level 1 employees a 5% raise.

```
proc sql;  
  update airline.payrollmaster  
    set Salary=Salary * 1.05  
    where JobCode like '__1';  
select *  
  from airline.payrollmaster;
```

73

c4s4d02

A SAS DATA step equivalent is as follows:

```
data airline.payrollmaster;  
  modify airline.payrollmaster;  
    if substr(JobCode,3)='1' then  
      Salary=Salary * 1.05;  
run;
```



You cannot create additional columns using the UPDATE statement.

Conditional Processing

Use a CASE expression to perform conditional processing. Assign new salaries based on job level. Two methods are available.

Method 1:

```
proc sql;
  update airline.payrollmaster
  set Salary=Salary *
      case substr(JobCode,3,1)
        when '1' then 1.05
        when '2' then 1.10
        when '3' then 1.15
        else 1.08
      end;
```

A CASE expression returns a single value. It is conditionally evaluated for each row of a table or view. Use multiple WHEN clauses when you want to execute the CASE expression for some but not all rows in the table. The optional ELSE expression provides an alternate action if none of the THEN expressions is executed.

Conditional Processing

Method 2:

```
proc sql;
  update airline.payrollmaster
    set Salary=Salary *
      case when substr(JobCode,3,1)='1'
        then 1.05
        when substr(JobCode,3,1)='2'
        then 1.10
        when substr(JobCode,3,1)='3'
        then 1.15
        else 1.08
      end;
end;
```

75

c4s4d04

Method 1 above is more efficient because the SUBSTR function is evaluated only once. This method also assumes an = comparison operator, which means that if you need a different operator, you must use Method 2.

If no ELSE expression is present and every WHEN condition is false, the result of the CASE expression is a missing value.

Conditional Processing

You can also use a CASE expression in other parts of a query, such as within a SELECT statement, to create new columns.

General form of the CASE expression within the SELECT statement:

```
SELECT column-1<, column-2> ...
  CASE <case-operand>
    WHEN when-condition THEN result-expression
    <WHEN when-condition THEN result-expression>
    <ELSE result-expression>
  END <as column>
FROM table;
```

76

Conditional Processing

Example: Display employee names, job codes, and job levels.

```
proc sql;
  select LastName, FirstName, JobCode,
         case substr(JobCode,3,1)
           when '1' then 'junior'
           when '2' then 'intermediate'
           when '3' then 'senior'
           else 'none'
         end as level
  from airline.payrollmaster,
       airline.staffmaster
  where staffmaster.EmpID=
        payrollmaster.EmpID;
```

77

c4s4d05

In traditional SAS programming language, you create a user-defined format with the FORMAT procedure to display a character string of your choice, in place of a stored value.

Conditional Processing

Partial Output

The SAS System			
LastName	FirstName	Job Code	level
ADAMS	GERALD	TA2	intermediate
ALEXANDER	SUSAN	ME2	intermediate
APPLE	TROY	ME1	junior
ARTHUR	BARBARA	FA3	senior
AVERY	JERRY	TA3	senior
BAREFOOT	JOSEPH	ME3	senior
BAUCOM	WALTER	SCP	none
BLAIR	JUSTIN	PT2	intermediate
BLALOCK	RALPH	TA2	intermediate
BOSTIC	MARIE	TA3	senior

78

Loading Data into a Table: Review

Method A: The SET Clause

```
INSERT INTO table  
  SET column-1=value,column-2=value,...;
```

Method B: The VALUES Clause

```
INSERT INTO table <(column-list)>  
  VALUES (value,value,value, ...);
```

Method C: A Query-expression

```
INSERT INTO table-1 <(column-list)>  
  SELECT columns FROM table-2
```

79

Deleting Rows

Use the DELETE statement to eliminate unwanted rows from a table or SAS/ACCESS view.

General form of the DELETE statement:

```
DELETE FROM table  
  WHERE expression;
```

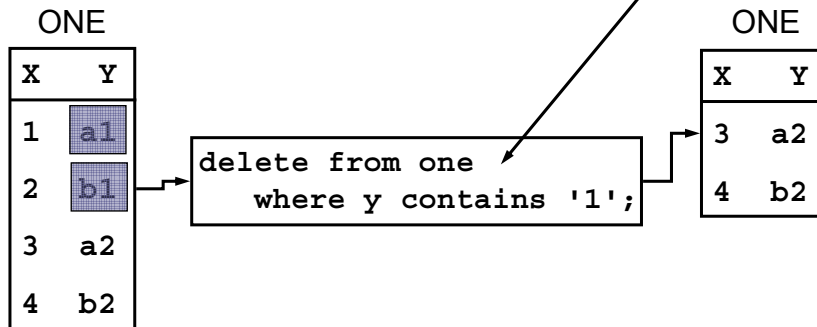
80



If you do not specify a WHERE clause, all rows are deleted.

Deleting Rows

What happens if you accidentally place a semicolon here?



81

c4s4d06
...

Deleting Rows

Example: From the **airline.frequentflyers** table, delete all frequent flyers who either used all their points or used more than they have.

```
proc sql;
  delete from airline.frequentflyers
    where PointsEarned-PointsUsed <= 0;
```

Partial Log

NOTE: 11 rows were deleted from AIRLINE.FREQUENTFLYERS.

82

c4s4d07



Compare this process with the subsetting IF statement used in traditional SAS programming language.

Altering Columns

Use the ALTER statement to manipulate columns in a table three different ways.

General form of the ALTER statement:

```
ALTER TABLE table  
  ADD column-definition, column-definition, ...  
  DROP column-1, column-2, ...  
  MODIFY column-definition, column-definition, ...;
```

83

Altering Columns

1. Add columns to a table.

```
proc sql;  
  alter table airline.payrollmaster  
    add Bonus num format=comma10.2,  
        Level char(3);
```

84

c4s4d08

After adding columns, use the UPDATE statement to assign values to those columns. These added columns initially contain missing values.

Altering Columns

2. Drop columns from a table.

```
proc sql;  
    alter table airline.flightdelays  
        drop DestinationType;
```

85

c4s4d09

An alternative is to use the DROP= data set option as follows:

```
create table airline.flightdelays  
    select *  
        from airline.flightdelays (drop=DestinationType);
```

Altering Columns

3. Modify attributes of existing columns in a table. You can alter a column's length, informat, format, and label.

```
proc sql;  
    alter table airline.payrollmaster  
        modify Bonus num format=comma8.2,  
            Level char(1)  
                label='Employee Level';
```

86

c4s4d10

Altering Columns

Example: Alter **airline.payrollmaster** as follows:

1. Add a new column named **Age**.
2. Change the **DateOfBirth** column to the MMDDYY10. format.
3. Drop the **DateOfHire** column.

Create the columns here.

Populate the rows here.

```
proc sql;
  alter table airline.payrollmaster
    add Age num
    modify DateOfBirth date format=mmddyy10.
    drop DateOfHire;
  update airline.payrollmaster
    set Age=int((today()-DateOfBirth)/365.25);
```

87

c4s4d11

Altering Columns

Before altering

The SAS System						
EmpID	Gender	Job Code	Salary	DateOfBirth	DateOfHire	
1919	M	TA2	\$48,126	16SEP1958	07JUN1985	
1653	F	ME2	\$49,151	19OCT1962	12AUG1988	
1400	M	ME1	\$41,677	08NOV1965	19OCT1988	

```
select *
  from airline.payrollmaster;
```

After altering

The SAS System					
EmpID	Gender	Job Code	Salary	DateOfBirth	Age
1919	M	TA2	\$48,126	09/16/1958	41
1653	F	ME2	\$49,151	10/19/1962	37
1400	M	ME1	\$41,677	11/08/1965	34

88

Deleting Tables, Indexes, and Views

Use the DROP statement to delete an entire table, SQL view, or index.

General form of the DROP statement:

```
DROP TABLE table-1, table-2, ...;  
DROP VIEW view-1, view-2, ...;  
DROP INDEX index-1, index-2, ...  
      FROM table;
```

89

Deleting Tables, Indexes, and Views

Example: Delete the index EmpID from the **airline.payrollmaster** table and delete the temporary table **Discount**.

Partial Log

```
proc sql;  
  drop index EmpID  
    from airline.payrollmaster;
```

NOTE: Index EmpID has been dropped.

```
drop table Discount;
```

NOTE: Table WORK.DISCOUNT has been dropped.

90

c4s4d12

When you delete a table, all indexes on that table are automatically deleted. If you copy a table, all indexes are copied.

In Summary

- UPDATE
 - SET
 - CASE
 } Modifies values in existing columns, that is, changes row values.
- ALTER
 - ADD
 - DROP
 - MODIFY
 } Adds or drops columns, or changes column attributes.
- INSERT
 - SET
 - VALUES
 } Inserts data rows at end of existing tables.
- DELETE
 } Removes rows as specified with a WHERE expression.
- DROP
 } Deletes an entire table, view, or index.

91

Updating Views

You can update the data underlying PROC SQL views using the INSERT, DELETE, and UPDATE statements, but

- you can only update a single table through a view. It cannot be joined or linked to another table, nor contain a subquery.
- you can update a column using the column's alias, but not a derived column.
- you cannot update the table through a summary query.
- you cannot update a view containing an ORDER BY clause.

92

Updating Views

Create a view...

```
proc sql;  
    create view airline.raise as  
        select EmpID, JobCode,  
               Salary, Salary/12  
               as MonthlySalary  
               format=dollar12.  
        from airline.payrollmaster;
```

... and then update the view.

```
proc sql;  
    update airline.raise  
        set Salary=Salary * 1.20  
        where JobCode='PT3';
```



Exercises

Submit a LIBNAME statement to assign **airline** to the course SAS data library.

TSO: `libname airline '.sql.sasdata';`

Directory-based systems: `libname airline '.';`

3. Modifying a Table

- a. Create a temporary table named **tdelay** that is a copy of the table **airline.flightdelays**, but contains only the data for March 1, 2000.
- b. Flight numbers must be modified to differentiate international flights from domestic flights. Change the **FlightNumber** column in the **tdelay** table from three characters to four characters wide.
- c. Modify the flight numbers so that international flights (**DestinationType**='International') have flight numbers beginning with 'I'. Domestic flight numbers remain the same. Display **tdelay**.
- d. Eliminate the **DestinationType** column from **tdelay**.
- e. Delay categories must be altered to reflect new standards. Change the values of the **DelayCategory** column as indicated below. Display the **tdelay** table.

Value of Delay	New Value of DelayCategory
0 and below	'No Delay'
1 to 15	'Acceptable'
16 and over	'Excessive'

- f. Delete the **tdelay** table.

4.5 Chapter Summary

You can use PROC SQL to create tables in several ways. You can define columns or borrow column definitions with the CREATE TABLE statement. Use an INSERT statement to enter rows of data into the table. Use the CREATE TABLE statement with an AS keyword to store the result of a query into a table.

A view is a stored query that contains no data but can be used as a table. You can create or update views using PROC SQL. You can use the DESCRIBE statement to display the definition of a PROC SQL view in the SAS log.

You can use PROC SQL to create indexes on tables. PROC SQL can use indexes to optimize the processing of WHERE clauses and joins.

PROC SQL enables you to alter or delete rows of data in existing tables or views using the UPDATE and DELETE statements, respectively. You can use the ALTER statement to add, delete, or modify the attributes of columns in an existing table. Use the DROP statement to delete tables, views, and indexes.

General form of PROC SQL using integrity constraints:

```
PROC SQL;  
  CREATE TABLE table  
    (column-specification,...  
    <constraint-specification,...>
```

General forms of the CREATE TABLE statement:

```
CREATE TABLE table  
  (column-1 type(length),  
  column-2 type(length), ...);
```

```
CREATE TABLE table-1  
  LIKE table-2;
```

```
CREATE TABLE table-1 AS  
  SELECT column-1, column-2, ...  
  FROM table-2 ...;
```


General forms of the INSERT statement:

```
INSERT INTO table
  SET column-1=value,
      column-2=value, ...;
```

```
INSERT INTO table
  VALUES (value,value, ...;)
```

```
INSERT INTO table-1
  SELECT column-1, column-2, ...
  FROM table-2 ...;
```

General form of the CREATE VIEW statement:

```
CREATE VIEW view-name AS
  query-expression;
```

General form of the DESCRIBE statement:

```
DESCRIBE VIEW view-name;
```

General form of the CREATE INDEX statement:

```
CREATE <UNIQUE> INDEX index-name
  ON table(column-1,column-2);
```

General form of the UPDATE statement:

```
UPDATE table | view
  SET column-1 expression,
      column-2=expression, ...
  WHERE expression;
```

General form of the CASE expression within the SELECT statement:

```
SELECT column-1<, column-2> ...
  CASE <case-operand>
    WHEN when-condition THEN result-expression
    <WHEN when-condition THEN result-expression>
    <ELSE result-expression>
  END;
```

General form of the DELETE statement:

```
DELETE FROM table
  WHERE expression;
```

General form of the ALTER statement:

```
ALTER TABLE table  
    ADD column-definition, column-definition, ...  
    DROP column-1, column-2, ...;  
    MODIFY column-definition, column-definition, ...
```

General forms of the DROP statement:

```
DROP TABLE table-1, table-2, ...;  
DROP view-1, view-2, ...;  
DROP INDEX index-1, index-2, ... FROM table;
```

4.6 Solutions to Exercises

1. Creating a Table

a.

```
proc sql;
  create table awards
    (ptsreqd num label='Points Required',
     rank num format=3.,
     award char(25));
```

b.

```
insert into awards
  values( 2000, 1, 'free night in hotel')
  values(10000, 2, '50% discount on flight')
  values(20000, 3, 'free domestic flight')
  values(40000, 4, 'free international flight');
```

Alternate Solution

```
insert into awards
  set ptsreqd=2000, rank=1,
    award='free night in hotel'
  set ptsreqd=10000, rank=2,
    award='50% discount on flight'
  set ptsreqd=20000, rank=3,
    award='free domestic flight'
  set ptsreqd=40000, rank=4,
    award='free international flight';
```

c.

```
select *
  from awards;
```

d.

```
title 'Awards Available to AZ Frequent Flyers';
select FFID,Name,
  PointsEarned-PointsUsed
  label='Available Points',
  award
  from airline.frequentflyers,awards
 where (PointsEarned-PointsUsed)>=ptsreqd
  and State='AZ'
 order by 1;
```

e. (Optional)

```
title 'Frequent Flyers Ineligible for Awards';
select distinct FFID, Name,
       PointsEarned-PointsUsed
       label='Available Points'
from airline.frequentflyers
where (PointsEarned-PointsUsed)<all
      (select ptsreqd
       from awards)
order by 1;
quit;
title;
```

2. Creating a View

a.

```
proc sql;
create view vsched as
select Date, FlightNumber,
       flightschedule.Destination,
       FirstName, LastName, staffmaster.EmpID
from airline.staffmaster, airline.flightschedule
where staffmaster.EmpID=flightschedule.EmpID;
title 'View VSCHED';
select *
from vsched
order by Date, FlightNumber, LastName;
```

b.

```
title 'Schedule for Deborah Young';
select *
from vsched
where EmpID='1431'
order by Date, FlightNumber;
```

c. (Optional)

```
title 'Number of Delayed Flights';
title2 'Experienced by Each Crew Member';
select FirstName, LastName,
       count(*) as count
from vsched, airline.flightdelays
where vsched.FlightNumber=flightdelays.FlightNumber
and vsched.Date=FlightDelays.Date
and Delay>0
group by FirstName, LastName
order by FirstName, LastName;
quit;
title;
```

3. Modifying a Table

a.

```
proc sql;  
  create table tdelay as  
    select *  
      from airline.flightdelays  
     where Date='01mar2000'd;
```

b.

```
alter table tdelay  
  modify FlightNumber char(4);
```

c.

```
update tdelay  
  set FlightNumber='I' || FlightNumber  
  where DestinationType='International';  
select *  
  from tdelay;
```

d.

```
alter table tdelay  
  drop DestinationType;
```

e.

```
update tdelay  
  set DelayCategory=  
    case  
      when Delay<=0 then 'No Delay'  
      when 0<Delay<=15 then 'Acceptable'  
      else 'Excessive'  
    end;  
select *  
  from tdelay;
```

f.

```
drop table tdelay;  
quit;
```


Chapter 5 Additional SQL Features

5.1	Setting SQL Procedure Options.....	5-3
5.2	Dictionary Tables and Views	5-9
5.3	Interfacing PROC SQL with Macro Language (Optional).....	5-18
5.4	Program Testing and Performance	5-29
5.5	Chapter Summary.....	5-34

5.1 Setting SQL Procedure Options

Objectives

- Use SQL procedure options to control processing details.
- Reset PROC SQL options without re-invoking the procedure.

3

Controlling Processing

The SQL procedure offers a variety of options and statements that affect processing.

General form of the PROC SQL statement:

```
PROC SQL options;
```

4

Controlling Processing

Selected options:

INOBS= n sets a limit of n rows from each source table that contributes to a query.

OUTOBS= n restricts the number of rows that a query outputs (displays or writes to a table).

continued...

5

Controlling Processing

PRINT|NOPRINT controls whether the results of a SELECT statement are displayed.

NONUMBER|NUMBER controls whether the row number is printed as the first column in the output.

NODOUBLE|DOUBLE double-spaces the report.

continued...

6

The default value appears first in the slides.

Controlling Processing

NOFLOW|FLOW|
FLOW=*n*|FLOW=*n m*

controls the appearance of wide character columns. The FLOW option causes text to be flowed in its column rather than wrapping the entire row. Specifying *n* determines the width of the flowed column. Specifying *n* and *m* floats the width of the column between the limits to achieve a balanced layout.

7

Controlling Processing

Example: Display the AWARDS table with flowed character columns and double-spacing.

```
proc sql flow=13 double;  
  select *  
    from awards;
```

8

c5s1d01

Controlling Processing

Output

Points Required	Rank	Award
2000	1	free night in hotel
10000	2	50% discount on flight
20000	3	free domestic flight
40000	4	free international flight

9

Controlling Processing

Example: Read ten rows from
`airline.marchflights`.

```
proc sql inobs=10;  
  select FlightNumber, Date  
    from airline.marchflights;
```

10

c5s1d02

Controlling Processing

Output

FlightNumber	Date
182	01MAR2000
114	01MAR2000
202	01MAR2000
219	01MAR2000
439	01MAR2000
387	01MAR2000
290	01MAR2000
523	01MAR2000
982	01MAR2000
622	01MAR2000

11



After you specify an option, it remains in effect until you change it or you re-invoke PROC SQL.

Resetting Options

You can use the RESET statement to add or change PROC SQL options without re-invoking the procedure.

General form of the RESET statement:

```
RESET options;
```

12

Resetting Options

Example: Display two rows from the payroll table and print the row number. Then display the rows without printing the row number.

```
proc sql outobs=2 number;
  select * from airline.payrollmaster;
```

Output

Row	Emp ID	Gender	Job Code	Salary	DateOfBirth	DateOfHire
1	1919	M	TA2	\$48,126	16SEP1958	07JUN1985
2	1653	F	ME2	\$49,151	19OCT1962	12AUG1988

13

c5s1d03

Resetting Options

```
reset nonumber;
select *
from airline.payrollmaster;
```

Output

Emp ID	Gender	Job Code	Salary	DateOfBirth	DateOfHire
1919	M	TA2	\$48,126	16SEP1958	07JUN1985
1653	F	ME2	\$49,151	19OCT1962	12AUG1988

14

c5s1d04

5.2 Dictionary Tables and Views

Objectives

- Use dictionary tables and views to obtain information about SAS files.

16

Overview

You can retrieve information about SAS session metadata by querying *dictionary tables* with PROC SQL. Dictionary tables follow these rules:

- created at initialization
- updated automatically
- limited to read-only access

17

“Metadata consist of information that characterizes data. Metadata are used to provide documentation for data products. In essence, metadata answer **who**, **what**, **when**, **where**, **why**, and **how** about every facet of the data that are being documented.”

<http://geology.usgs.gov/tools/metadata/tools/doc/faq.html#motivation>

Overview

The metadata available in dictionary tables includes the following:

- SAS files
- external files
- system options, macros, titles, and footnotes

18

Overview of SAS File Metadata

DICTIONARY.MEMBERS

- general information about data library members

DICTIONARY.TABLES

- detailed information about data sets

DICTIONARY.COLUMNS

- detailed information on variables and their attributes

DICTIONARY.CATALOGS

- information about catalog entries

DICTIONARY.VIEWS

- general information about data views

DICTIONARY.INDEXES

- information on indexes defined for data files

continued...

19

Overview of SAS File Metadata

DICTIONARY.CHECK_CONSTRAINTS

- information about known check constraints

DICTIONARY.CONSTRAINT_COLUMN_USAGE

- information about columns that are referred to by integrity constraints

DICTIONARY.CONSTRAINT_TABLE_USAGE

- information about tables that have integrity constraints defined on them

DICTIONARY.TABLE_CONSTRAINTS

- information about integrity constraints in all known tables

DICTIONARY.REFERENTIAL_CONSTRAINTS

- information about referential constraints

20

Overview of Other Metadata

DICTIONARY.EXTFILES

- information about currently assigned filerefs

DICTIONARY.OPTIONS

- information about current settings of SAS system options

DICTIONARY.MACROS

- information about macro variables

DICTIONARY.TITLES

- information about text assigned to titles and footnotes

21

SAS librefs are limited to eight characters. DICTIONARY is an automatically assigned, reserved word.

Exploring Dictionary Tables

```
describe table dictionary.tables;
```

Partial Log

NOTE: SQL table DICTIONARY.TABLES was created like:

```
create table DICTIONARY.TABLES
(
  libname char(8) label='Library Name',
  memname char(32) label='Member Name',
  memtype char(8) label='Member Type',
  dbms_memtype char(32) label='DBMS Member Type',
  memlabel char(256) label='Dataset Label',
  typemem char(8) label='Dataset Type',
  crdate num format=DATETIME informat=DATETIME label='Date Created',
  modate num format=DATETIME informat=DATETIME label='Date Modified',
  nobs num label='Number of Observations',
  obslen num label='Observation Length',
  nvar num label='Number of Variables', ...);
```

The DESCRIBE TABLE statement is a good tool for exploring dictionary tables. The complete log notes from the DESCRIBE statement are shown below:

```
create table DICTIONARY.TABLES
(
  libname char(8) label='Library Name',
  memname char(32) label='Member Name',
  memtype char(8) label='Member Type',
  dbms_memtype char(32) label='DBMS Member Type',
  memlabel char(256) label='Dataset Label',
  typemem char(8) label='Dataset Type',
  crdate num format=DATETIME informat=DATETIME label='Date Created',
  modate num format=DATETIME informat=DATETIME label='Date Modified',
  nobs num label='Number of Physical Observations',
  obslen num label='Observation Length',
  nvar num label='Number of Variables',
  protect char(3) label='Type of Password Protection',
  compress char(8) label='Compression Routine',
  encrypt char(8) label='Encryption',
  npage num label='Number of Pages',
  filesize num label='Size of File',
  pcompress num label='Percent Compression',
  reuse char(3) label='Reuse Space',
  bufsize num label='Bufsize',
  delobs num label='Number of Deleted Observations',
  nlobs num label='Number of Logical Observations',
  maxvar num label='Longest variable name',
  maxlabel num label='Longest label',
  maxgen num label='Maximum number of generations',
  gen num label='Generation number',
  attr char(3) label='Dataset Attributes',
  indxtype char(9) label='Type of Indexes',
  datarep char(32) label='Data Representation',
  sortname char(8) label='Name of Collating Sequence',
  sorttype char(4) label='Sorting Type',
  sortchar char(8) label='Charset Sorted By',
  reqvector char(24) format=$HEX48 informat=$HEX48 label='Requirements Vector',
  datarepname char(170) label='Data Representation Name',
  encoding char(256) label='Data Encoding',
  audit char(3) label='Audit Trail Active?',
  audit_before char(3) label='Audit Before Image?',
  audit_admin char(3) label='Audit Admin Image?',
  audit_error char(3) label='Audit Error Image?',
  audit_data char(3) label='Audit Data Image?'
);
```

Using Dictionary Information

Example: Display information about the files in the **airline** library.

```
options nolabel nocenter;
proc sql;
select memname format=$20.,nobs,nvar,crdate
   from dictionary.tables
   where libname='AIRLINE';
```

23

c5s2d02

Using Dictionary Information

Output

memname	nobs	nvar	crdate
FAVIEW	.	6	15SEP06:20:15:59
FLIGHTDELAYS	624	8	18MAR00:20:53:16
FLIGHTSCHEDULE	270	4	18MAR00:20:53:17
FREQUENTFLYERS	206	11	02NOV01:09:12:41
INTERNATIONALFLIGHTS	201	4	18MAR00:20:53:16
MARCHFLIGHTS	635	13	18MAR00:20:53:16
MECHANICSLEVEL1	8	3	18MAR00:20:53:17
MECHANICSLEVEL2	14	3	18MAR00:20:53:17
MECHANICSLEVEL3	7	3	18MAR00:20:53:17
PAYROLLCHANGES	6	6	18MAR00:20:53:17
PAYROLLMASTER	148	6	18MAR00:20:53:17
PAYROLLMASTER2	148	6	21JAN04:13:33:35
STAFFCHANGES	6	6	18MAR00:20:53:18
STAFFMASTER	148	6	18MAR00:20:53:17
SUPERVISORS	19	3	18MAR00:20:53:18

24

Using Dictionary Information

Example: Determine which tables contain the EmpID column.

```
select memname
  from dictionary.columns
 where libname='AIRLINE' and name='EmpID';
```

25

c5s2d03

Using Dictionary Information

Output

memname
FLIGHTSCHEDULE
MECHANICSLEVEL1
MECHANICSLEVEL2
MECHANICSLEVEL3
PAYROLLCHANGES
PAYROLLMASTER
STAFFCHANGES
STAFFMASTER
SUPERVISORS

26

Using Dictionary Information

To use session metadata in other procedures or in a DATA step, you can do the following:

- create a PROC SQL view based on a dictionary table
- use views provided in the SASHELP library that are based on the dictionary tables

27

Using Dictionary Information

Example: Use `sashelp.vmember` to extract information from `DICTIONARY.MEMBERS` in a PROC TABULATE step.

```
proc tabulate data=sashelp.vmember format=8.;  
  class libname memtype;  
  keylabel N=' '  
  table libname, memtype/rts=10  
         misstext='None';  
run;
```

28

c5s2d04

Using Dictionary Information

Output

	Member Type				
	CATALOG	DATA	ITEMSTOR	MDDb	VIEW
Library Name					
AIRLINE	None	14	None	None	1
SASHELP	101	91	2	2	30
SASUSER	5	106	1	None	4
SQL	None	14	None	None	1
WORK	None	2	None	None	1

5.3 Interfacing PROC SQL with Macro Language (Optional)

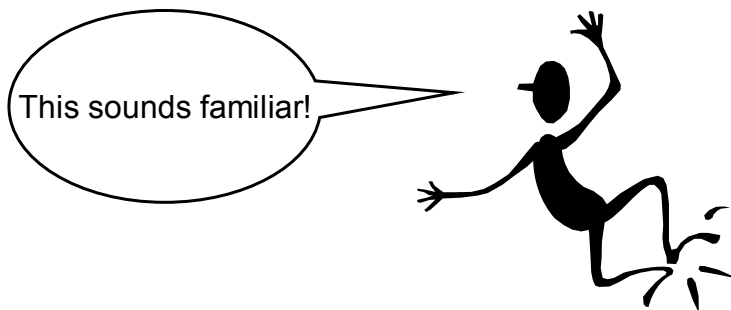
Objectives

- Create and use SAS macro variables in PROC SQL.
- Understand the use of SAS macros with SQL processing.
- Use the automatic SAS macro variables created by PROC SQL.

31

Resolving Symbolic References

Macro variable references embedded within PROC SQL code are resolved as the source code is tokenized.



32

Resolving Symbolic References

```
%let datasetname=payrollmaster;
%let bigsalary=100000;
```

SYMBOL TABLE	
<u>Name</u>	<u>Value</u>
datasetname	payrollmaster
bigsalary	100000

33

c5s3d01
...

Resolving Symbolic References

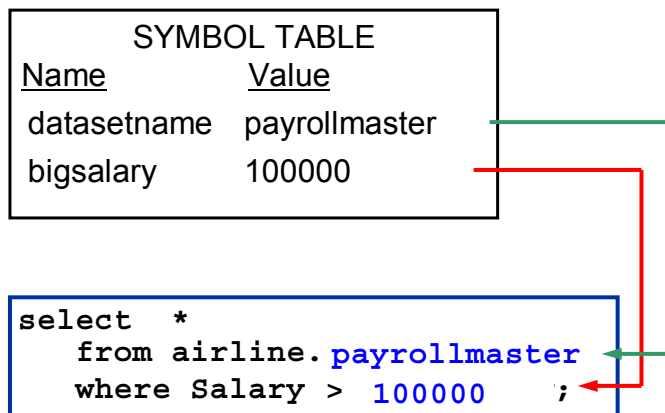
SYMBOL TABLE	
<u>Name</u>	<u>Value</u>
datasetname	payrollmaster
bigsalary	100000

```
select *
  from airline.&datasetname
 where Salary > &bigsalary;
```

34

c5s3d02
...

Resolving Symbolic References



35

c5s3d02
...

Creating Macro Variables

- SQL enables a query to pass data values to variables in the host software system. SAS chose to implement these host variables as macro variables.
- PROC SQL can create or update macro variables using an INTO clause. This clause can be used in three ways.

36

PROC SQL can create or update macro variables in either local or global symbol tables.

The INTO clause occurs between the SELECT and FROM clauses. It cannot be used in a CREATE TABLE or CREATE VIEW statement. Use the NOPRINT option if you do not need a display of the query result.

Creating Macro Variables: Method 1

General form of the SELECT statement with an INTO keyword:

```
SELECT column-1, column-2, ...
INTO :macro-var-1, :macro-var-2, ...
FROM ...
```

Method 1 extracts values **only** from the **first** row of the query result.

37

This method is often used with queries that return only one row.

Creating Macro Variables: Method 1

```
reset noprint;
select avg(Salary),
       min(Salary),
       max(Salary)
into :mean, :min, :max
from airline.payrollmaster;
%put &mean &min &max;
```

Partial Log

```
54079.65  25120.2 155930.6
```

38

c5s3d03

Creating Macro Variables: Method 1

Calculate the average salary of employees with a particular job code. Store the average in a macro variable and use the average to display all employees in that job code who have a salary above the average. Place the average in a title.

39

Creating Macro Variables: Method 1

```
%let code=NA1;
proc sql noprint;
select avg(Salary) into :mean
    from airline.payrollmaster
    where JobCode="&code";
reset print;
title1 "&code Employees Earning Above- "
      "Average Salaries";
title2 "Average Salary for &code Employees "
      "Is &mean";
select *
    from airline.payrollmaster
    where Salary > &mean and JobCode="&code";
```

40

c5s3d04

Creating Macro Variables: Method 1

Output

NA1 Employees Earning Above-Average Salaries					
Average Salary for NA1 Employees Is 58845.08					
Job					
EmpID	Gender	Code	Salary	DateOfBirth	DateOfHire
1839	F	NA1	\$60,806	02DEC1968	07JUL1991
1332	M	NA1	\$59,049	20SEP1968	07JUN1989
1443	F	NA1	\$59,184	21NOV1966	01SEP1989

41

Creating Macro Variables: Method 2

General form of the SELECT statement to create a macro variable:

```
SELECT a, b, ...
INTO :a1-:an, :b1-:bn
FROM ...
```

Method 2 extracts values from the first n rows of the query result and puts them into a series of n macro variables.

42

Creating Macro Variables: Method 2

How many frequent flyers are in each of the three member types (GOLD, SILVER, BRONZE)?

```
reset noprint;  
select MemberType,  
       count(*) as Frequency  
into :memtype1-:memtype3,:freq1-:freq3  
from airline.frequentflyers  
group by MemberType;
```

43

c5s3d05

Creating Macro Variables: Method 2

Example

```
%put Member types: &memtype1 &memtype2 &memtype3;  
%put Frequencies: &freq1 &freq2 &freq3;
```

Partial Log

```
Member types: BRONZE GOLD SILVER  
Frequencies: 61 60 85
```

44

c5s3d06

Creating Macro Variables: Method 3

General form of the SELECT statement to create a macro variable:

```
SELECT column-1, column-2, ...
INTO :macro-var-1 SEPARATED BY 'delimiter' ,
      :macro-var-2 SEPARATED BY 'delimiter' ...
FROM ...
```

Method 3 extracts values from all rows of the query result and puts them into a single macro variable, separated by the specified delimiter.

45

Creating Macro Variables: Method 3

Put the unique values of all international destinations into a single macro variable.

```
select distinct Destination
  into :airportcodes
      separated by ' '
  from airline.internationalflights;
%put &airportcodes;
```

Partial Log

```
CDG CPH FRA LHR YYZ
```

46

c5s3d07

The long string value in the macro variable can be parsed into the individual short values using %SCAN.

Automatic Macro Variables

Execution of a PROC SQL query or non-query statement updates the following automatic macro variables:

SQLLOBS	records the number of rows output or deleted.
SQLRC	contains the return code from each SQL statement.
SQLLOOPS	contains the number of iterations processed by the inner loop of PROC SQL.

47

Automatic Macro Variables

Write a macro that accepts a state code as a parameter and creates a table containing employees from that state. Display a maximum of ten rows from the table.

48

Automatic Macro Variables

```
%macro state(st);
proc sql;
create table &st as
select LastName, FirstName
  from airline.staffmaster
 where State="&st";
%put NOTE: The table &st has &sqlobs rows.;
title1 "&st Employees";
%if &sqlobs > 10 %then %do;
%put
  NOTE: Only the first 10 rows are displayed.;
  title2 "NOTE: Only 10 rows are displayed.";
  reset outobs=10;
%end;
select * from &st;
quit;
%mend state;
```

49

c5s3d08

Automatic Macro Variables

```
%state(NY)
```

Partial Log

```
NOTE: Table WORK.NY created, with 89 rows and 2 columns.
NOTE: The table NY has 89 rows.
NOTE: Only the first 10 rows are displayed.
WARNING: Statement terminated early due to OUTOBS=10 option.
```

50

c5s3d09

Automatic Macro Variables

Output

NY Employees	
NOTE: Only 10 rows are displayed.	
LastName	FirstName
APPLE	TROY
ARTHUR	BARBARA
BAUCOM	WALTER
BLALOCK	RALPH
BOSTIC	MARIE
BOYCE	JONATHAN
BRADLEY	JEREMY
BRYANT	LEONARD
BURNETTE	THOMAS
CAHILL	MARSHALL

5.4 Program Testing and Performance

Objectives

- Use PROC SQL options to test SQL code.
- Understand SAS log messages and accurately benchmark SAS code.

Testing and Performance Options

PROC SQL statement options are available to aid in testing programs and evaluating performance.

The following are selected options:

- EXEC|NOEXEC controls whether submitted SQL statements are executed.
- NOSTIMER|STIMER reports performance statistics in the SAS log for each SQL statement.
- NOERRORSTOP|ERRORSTOP is used in batch and noninteractive jobs to make PROC SQL enter syntax-check mode after an error occurs.

54



To use the STIMER SQL option, the system option STIMER or FULLSTIMER must also be in effect.

Other PROC SQL statement options that are useful in testing include

- INOBS=*n*
- OUTOBS=*n*.

Testing and Performance Options

Display the columns that are retrieved when you use SELECT * in a query and display any macro variable resolutions, but do not execute the query.

```
%let datasetname=payrollmaster;

proc sql feedback noexec;
  select *
    from airline.&datasetname;
```

55

c5s4d01

Testing and Performance Options

Partial Log

NOTE: Statement transforms to:

```
select PAYROLLMASTER.EmpID, PAYROLLMASTER.Gender,
PAYROLLMASTER.JobCode,
PAYROLLMASTER.Salary, PAYROLLMASTER.DateOfBirth,
PAYROLLMASTER.DateOfHire
      from AIRLINE.PAYROLLMASTER;
```

NOTE: Statement not executed due to NOEXEC option.

56

Testing and Performance Options

This is a log from a PROC SQL step with the STIMER statement option that executes a single query. The first note concerns the invocation of PROC SQL:

NOTE: The SQL statement used the following resources:

```
CPU      time -          00:00:00.01
Elapsed time -          00:00:00.68
EXCP count - 28
Task memory - 110K (20K data, 90K program)
Total memory - 864K (760K data, 104K program)
```

The second note concerns the query itself.

NOTE: The SQL statement used the following resources:

```
CPU      time -          00:00:00.23
Elapsed time -          00:00:03.61
EXCP count - 157
Task memory - 1213K (828K data, 385K program)
Total memory - 2258K (1840K data, 418K program)
```

57



This program was run in batch under z/OS. Performance measures, as well as the actual numbers, vary greatly across installations and operating systems. (The query used is the one about the supervisors of the crew on the Copenhagen flight, using subqueries and in-line views, from Section 3.3.)

Testing and Performance Options

Example

The third note reflects the totals for the procedure.

```
NOTE: The SQL procedure used the following resources:
      CPU      time -          00:00:00.25
      Elapsed time -          00:00:04.34
      EXCP count - 186
      Task memory - 1213K (828K data, 385K program)
      Total memory - 2258K (1840K data, 418K program)
```

58

General Guidelines for Benchmarking Programs

- Never use elapsed time for comparison because it might be affected by concurrent tasks.
- Benchmark two programs in separate SAS sessions. If benchmarking is done within one SAS session, statistics for the second program can be misleading because the SAS supervisor might have loaded modules into memory from prior steps.

59

General Guidelines for Benchmarking Programs

- Run each program multiple times and average the performance statistics.
- Use realistic data for tests. Program A could be better than program B on small tables and worse on large tables.

5.5 Chapter Summary

You can use options in the PROC SQL statement to affect SQL processing. You can limit the number of rows read or written during a query or limit the number of internal loops PROC SQL performs. PROC SQL can notify you when any of the processing limits that you set are reached.

Options are also available that affect the form of the output. You can flow character columns, number your rows, or double-space output. The RESET statement enables you to change options without having to re-invoke the procedure.

Dictionary tables can be queried to display SAS session metadata. The dictionary tables are generated at run time and are read-only. You can also use views stored in the SASHELP library that are based on the dictionary tables.

You can combine the SAS macro facility with PROC SQL in the same way as any other SAS step. PROC SQL, however, is capable of passing data from a query result into a macro variable. PROC SQL also updates several automatic macro variables that contain information about the last query executed.

There are PROC SQL statement options available to test and evaluate program performance. For example, the STIMER option in the PROC SQL statement can request resource usage information on each statement executed. The SAS log displays information on CPU usage, I/O counts, and other statistics.

General form of the PROC SQL statement:

PROC SQL *options*;

General form of the RESET statement:

RESET *options*;

General forms of the SELECT statement with an INTO keyword:

```
SELECT column-1, column-2, ...
      INTO :macro-var-1, :macro-var-2, ...
      FROM ...

SELECT a, b, ...
      INTO :a1-:an, :b1-:bn
      FROM    ...

SELECT column-1, column-2, ...
      INTO :macro-var-1 SEPARATED BY 'delimiter'
      FROM ...
```


Macro variables created by PROC SQL:

&SQLOBS

&SQLRC

&SQLOOPS

Selected PROC SQL statement options:

INOBS=*n*

OUTOBS= *n*

PRINT|NOPRINT

NONUMBER|NUMBER

NODOUBLE|DOUBLE

NOFLOW|FLOW|FLOW=*n*|FLOW=*n m*

EXEC|NOEXEC

NOSTIMER|STIMER

NOERRORSTOP|ERRORSTOP

Appendix A Overview of Table and Column Names

A.1	Table and Column Names Sorted by Column Names	A-3
------------	--	------------

A.1 Table and Column Names Sorted by Column Names

<u>Column Name</u>	<u>Table Names</u>
Address	FREQUENTFLYERS
Boarded	INTERNATIONALFLIGHTS MARCHFLIGHTS
City	FREQUENTFLYERS STAFFCHANGES STAFFMASTER
Date	FLIGHTDELAYS FLIGHTSCHEDULE INTERNATIONALFLIGHTS MARCHFLIGHTS
DateOfBirth	PAYROLLCHANGES PAYROLLMASTER
DateOfHire	PAYROLLCHANGES PAYROLLMASTER
DayOfWeek	FLIGHTDELAYS
Delay	FLIGHTDELAYS
DelayCategory	FLIGHTDELAYS
DepartureTime	MARCHFLIGHTS
Deplaned	MARCHFLIGHTS
Destination	FLIGHTDELAYS FLIGHTSCHEDULE INTERNATIONALFLIGHTS MARCHFLIGHTS
DestinationType	FLIGHTDELAYS
Distance	MARCHFLIGHTS
EmpID	FLIGHTSCHEDULE MECHANICSLEVEL1 MECHANICSLEVEL2 MECHANICSLEVEL3 PAYROLLCHANGES PAYROLLMASTER STAFFCHANGES STAFFMASTER SUPERVISORS

<u>Column Name</u>	<u>Table Names</u>
FFID	FREQUENTFLYERS
FirstName	STAFFCHANGES STAFFMASTER
FlightNumber	FLIGHTDELAYS FLIGHTSCHEDULE INTERNATIONALFLIGHTS MARCHFLIGHTS
Freight	MARCHFLIGHTS
Gender	PAYROLLCHANGES PAYROLLMASTER
JobCategory	SUPERVISORS
JobCode	MECHANICSLEVEL1 MECHANICSLEVEL2 MECHANICSLEVEL3 PAYROLLCHANGES PAYROLLMASTER
LastName	STAFFCHANGES STAFFMASTER
Mail	MARCHFLIGHTS
MemberType	FREQUENTFLYERS
MilesTraveled	FREQUENTFLYERS
Name	FREQUENTFLYERS
Nonrevenue	MARCHFLIGHTS
Origin	FLIGHTDELAYS MARCHFLIGHTS
PassengerCapacity	MARCHFLIGHTS
PhoneNumber	FREQUENTFLYERS STAFFCHANGES STAFFMASTER
PointsEarned	FREQUENTFLYERS
PointsUsed	FREQUENTFLYERS
Salary	MECHANICSLEVEL1 MECHANICSLEVEL2 MECHANICSLEVEL3 PAYROLLCHANGES PAYROLLMASTER

<u>Column Name</u>	<u>Table Names</u>
State	FREQUENTFLYERS STAFFCHANGES STAFFMASTER SUPERVISORS
Transferred	MARCHFLIGHTS
ZipCode	FREQUENTFLYERS

Appendix B Overview of Table and Column Names

B.1 Table and Column Names Sorted by Table Name	B-3
--	------------

B.1 Table and Column Names Sorted by Table Name

Table Name	Column Name	Column Type	Column Length	Column Format
FLIGHTDELAYS	FlightNumber	char	3	
	Date	num	8	DATE9.
	Origin	char	3	
	Destination	char	3	
	DelayCategory	char	15	
	DestinationType	char	15	
	DayOfWeek	num	8	
	Delay	num	8	
FLIGHTSCHEDULE	FlightNumber	char	3	\$3.
	Date	num	8	DATE9.
	Destination	char	3	\$3.
	EmpID	char	4	
FREQUENTFLYERS	FFID	char	6	
	MemberType	char	6	
	Name	char	25	\$18.
	Address	char	20	
	PhoneNumber	char	12	
	City	char	20	\$20.
	State	char	2	\$2.
	ZipCode	char	5	\$5.
	MilesTraveled	num	8	10.
	PointsEarned	num	8	10.
	PointsUsed	num	8	10.
INTERNATIONALFLIGHTS	FlightNumber	char	3	
	Date	num	8	DATE9.
	Destination	char	3	
	Boarded	num	8	

Table Name	Column Name	Column Type	Column Length	Column Format
MARCHFLIGHTS	FlightNumber	char	3	
	Date	num	8	DATE9.
	DepartureTime	num	8	TIME5.
	Origin	char	3	
	Destination	char	3	
	Distance	num	8	
	Mail	num	8	
	Freight	num	8	
	Boarded	num	8	
	Transferred	num	8	
	Nonrevenue	num	8	
	Deplaned	num	8	
	PassengerCapacity	num	8	
MECHANICSLEVEL1	EmpID	char	4	
	JobCode	char	3	
	Salary	num	8	DOLLAR9.
MECHANICSLEVEL2	EmpID	char	4	
	JobCode	char	3	
	Salary	num	8	DOLLAR9.
MECHANICSLEVEL3	EmpID	char	4	
	JobCode	char	3	
	Salary	num	8	DOLLAR9.
PAYROLLCHANGES	EmpID	char	4	
	Gender	char	1	
	JobCode	char	3	
	Salary	num	8	DOLLAR9.
	DateOfBirth	num	8	DATE9.
	DateOfHire	num	8	DATE9.
PAYROLLMASTER	EmpID	char	4	
	Gender	char	1	
	JobCode	char	3	
	Salary	num	8	DOLLAR9.
	DateOfBirth	num	8	DATE9.
	DateOfHire	num	8	DATE9.

Table Name	Column Name	Column Type	Column Length	Column Format
STAFFCHANGES	EmpID	char	4	
	LastName	char	15	
	FirstName	char	15	
	City	char	15	
	State	char	2	
	PhoneNumber	char	12	
STAFFMASTER	EmpID	char	4	
	LastName	char	15	
	FirstName	char	15	
	City	char	15	
	State	char	2	
	PhoneNumber	char	12	
SUPERVISORS	EmpID	char	4	
	State	char	2	
	JobCategory	char	2	

Appendix C Table Listings

C.1	Partial Table Listings.....	C-3
------------	------------------------------------	------------

C.1 Partial Table Listings

AIRLINE.FLIGHTDELAYS Table

Flight Number	Date	Origin	Destination	Delay Category	Destination Type	Day Of Week	Delay
182	01MAR2000	LGA	YYZ	No Delay	International	4	0
114	01MAR2000	LGA	LAX	1-10 Minutes	Domestic	4	8
202	01MAR2000	LGA	ORD	No Delay	Domestic	4	-5
219	01MAR2000	LGA	LHR	11+ Minutes	International	4	18
439	01MAR2000	LGA	LAX	No Delay	Domestic	4	-4

AIRLINE.FLIGHTSCHEDULE Table

Flight Number	Date	Destination	Emp ID
132	01MAR2000	YYZ	1739
132	01MAR2000	YYZ	1478
132	01MAR2000	YYZ	1130
132	01MAR2000	YYZ	1390
132	01MAR2000	YYZ	1983

AIRLINE.FREQUENTFLYERS Table

FFID	Member Type	Name	Address	PhoneNumber
WD7152	BRONZE	COOPER, LESLIE	66 DRIVING WAY	501/377-0703
WD8472	BRONZE	LONG, RUSSELL	9813 SUMTER SQUARE	501/367-1097
WD1576	GOLD	BRYANT, ALTON	763 THISTLE DRIVE	501/776-0631
WD3947	SILVER	NORRIS, DIANE	77 PARKWAY PLAZA	501/377-3739
WD9347	SILVER	PEARSON, BRYAN	9999 MARKUP MANOR	501/855-4780

City	State	Zip Code	Miles Traveled	Points Earned	PointsUsed
Little Rock	AR	72201	30833	31333	0
Monticello	AR	71655	25570	26070	0
Bauxite	AR	72011	56144	58644	27000
North Little Rock	AR	72119	40922	45922	23000
Bella Vista	AR	72714	4839	9839	0

AIRLINE.MECHANICSLLEVEL2 Table

Emp ID	Job Code	Salary
1653	ME2	\$49,151
1782	ME2	\$49,483
1244	ME2	\$51,695
1065	ME2	\$49,126
1129	ME2	\$48,901

AIRLINE.MECHANICSLLEVEL3 Table

Emp ID	Job Code	Salary
1499	ME3	\$60,235
1409	ME3	\$58,171
1379	ME3	\$59,170
1521	ME3	\$58,136
1385	ME3	\$61,460

AIRLINE.PAYROLLCHANGES Table

Emp ID	Gender	Job Code	Salary	DateOf Birth	DateOf Hire
1639	F	TA3	\$59,164	30JUN1955	31JAN1982
1065	M	ME3	\$53,326	29JAN1942	10JAN1985
1561	M	TA3	\$51,120	03DEC1961	10OCT1985
1221	F	FA3	\$41,854	25SEP1965	07OCT1989
1447	F	FA1	\$30,972	11AUG1970	01NOV2000

AIRLINE.PAYROLLMASTER Table

Emp ID	Gender	Job Code	Salary	DateOf Birth	DateOf Hire
1919	M	TA2	\$48,126	16SEP1958	07JUN1985
1653	F	ME2	\$49,151	19OCT1962	12AUG1988
1400	M	ME1	\$41,677	08NOV1965	19OCT1988
1350	F	FA3	\$46,040	04SEP1963	01AUG1988
1401	M	TA3	\$54,351	16DEC1948	21NOV1983

AIRLINE.STAFFCHANGES Table

Emp ID	LastName	First Name	City	State	PhoneNumber
1639	CARTER	KAREN	STAMFORD	CT	203/781-8839
1065	CHAPMAN	NEIL	NEW YORK	NY	718/384-5618
1561	SANDERS	RAYMOND	NEW YORK	NY	212/588-6615
1221	WALTERS	DIANE	NEW YORK	NY	718/384-1918
1447	BRIDESTON	AMY	NEW YORK	NY	718/384-1213

AIRLINE.STAFFMASTER Table

Emp ID	LastName	First Name	City	State	PhoneNumber
1919	ADAMS	GERALD	STAMFORD	CT	203/781-1255
1653	ALEXANDER	SUSAN	BRIDGEPORT	CT	203/675-7715
1400	APPLE	TROY	NEW YORK	NY	212/586-0808
1350	ARTHUR	BARBARA	NEW YORK	NY	718/383-1549
1401	AVERY	JERRY	PATERSON	NJ	201/732-8787

AIRLINE.SUPERVISORS Table

Emp ID	State	Job Category
1677	CT	BC
1834	NY	BC
1431	CT	FA
1433	NJ	FA
1983	NY	FA

Appendix D Index

=

= comparison operator, 4-49

A

administering a view, 4-25–4-30

alias, 2-14, 2-59, 3-20–3-21, 3-23, 3-35, 4-23

ALL keyword, 2-57, 3-49–3-51, 3-62, 3-68, 3-71

ALTER statement

SQL procedure, 2-9

altering columns, 4-53–4-55

ANY keyword, 2-56

APPEND procedure, 3-74

AS keyword, 2-14, 3-20–3-21

B

benchmarking programs

guidelines, 5-33

Boolean expression, 3-29–3-30

BY statement

DATA step, 3-21–3-22

C

CALCULATED keyword, 2-28–2-29

Cartesian product, 3-7–3-9, 3-17, 3-25–3-26

CASE expression, 4-48–4-49

character functions, 1-9–1-12

CHECK integrity constraint, 4-12

clauses

ESCAPE, 2-25

FROM, 3-7, 3-9, 3-15, 3-20–3-21, 3-27, 5-20

GROUP BY, 2-45

HAVING, 2-50

ON, 3-15

ORDER BY, 2-31–2-33

SELECT, 5-20

SET, 4-8

VALUES, 4-8–4-9

WHERE, 2-19–2-24, 2-26, 2-50, 3-9, 3-25, 4-51

COALESCE function, 3-23

concatenation, 3-74

concatenation operator, 1-11

conditional processing, 4-48–4-50

CONSTRAINTS keyword, 4-18

CONTAINS operator

indexes, 4-33

controlling processing, 5-3–5-7

conventional joins, 3-11

correlated subqueries, 2-59–2-71

CORRESPONDING keyword, 3-49–3-51, 3-72

COUNT function, 2-47

CREATE INDEX statement, 4-18, 4-36

CREATE statement

SQL procedure, 2-9

CREATE TABLE statement, 4-3, 4-7, 4-18

SQL procedure, 5-20

CREATE VIEW statement, 4-22

SQL procedure, 5-20

creating a view, 4-21–4-23, 4-29

creating indexes, 4-31–4-40

D

data

combining from multiple tables, 3-3

DATA step

BY statement, 3-21–3-22

LENGTH statement, 1-12

MERGE statement, 3-21–3-22

merges, 3-12, 3-40

SET statement, 3-72

DATA step merges, 3-12

data types, 4-4

DATASETS procedure, 4-11

integrity constraints, 4-17

defining columns, 4-3–4-5

DELETE statement, 4-51

SQL procedure, 2-9

deleting rows, 4-51–4-52

deleting tables, indexes, and views, 4-56

DESCRIBE statement

SQL procedure, 2-9

DESCRIBE TABLE statement, 4-18, 5-13

dictionary tables, 5-9–5-16

DESCRIBE TABLE statement, 5-13

SASHELP library, 5-16

DIF function, 2-16

DISTINCT keyword, 2-18

DROP statement

SQL procedure, 2-9, 4-56

E

ELSE expression, 4-49

equijoins, 3-11, 3-26

ESCAPE clause, 2-25
EXCEPT operator, 3-46–3-47, 3-50–3-59
EXISTS condition, 2-65
expressions
 CASE, 4-48–4-49
 THEN, 4-48

F
FEEDBACK option, 2-14
FOOTNOTE statement, 2-37
FOREIGN KEY integrity constraint, 4-12
FORMAT procedure, 4-50
FORMAT= option, 2-35
FROM clause, 3-7, 3-9, 3-15, 3-20–3-21, 3-27, 5-20
FULLSTIMER system option, 5-30
functions, 2-42–2-45
 character, 1-9–1-12
 COALESCE, 3-23
 COUNT, 2-47
 DIF, 2-16
 LAG, 2-16
 numeric, Error! Not a valid bookmark in entry on page 1-9
 SCAN, 1-11–1-12
 SUBSTR, 1-9–1-11, 4-49
 TRIM, 1-11

G
GROUP BY clause, 2-45

H
HAVING clause, 2-50

I
IDXNAME= data set option, 4-39
IDXWHERE= option, 4-39
INDEX keyword, 4-36
indexes, 4-31–4-40
 benefits and costs, 4-40
 composite, 4-34
 deleting, 4-56
 guidelines, 4-40
 IDXNAME= data set option, 4-39
 IDXWHERE= option, 4-39
 naming rules, 4-34
 simple, 4-34
INFORMAT= option, 2-35
in-line views, 3-27–3-30, 3-35
inner joins, 3-5–3-6, 3-8–3-14, 3-17, 3-21–3-22
 conventional joins, 3-11

 equijoins, 3-11
 natural joins, 3-11
INSERT statement
 SQL procedure, 2-10
integrity constraints, 4-11–4-14
 DATASETS procedure, 4-17
 SQL procedure, 4-13–4-14
INTERSECT operator, 3-46–3-48, 3-51, 3-60–3-65
INTO keyword, 5-21

J

JOIN operator, 3-15
joins, 3-3, 3-24–3-26
 comparing with DATA step merge, 3-40
 conventional, 3-11
 full, 3-17
 inner, 3-5–3-6, 3-8–3-14, 3-17, 3-21–3-22
 left, 3-15
 multiway, 3-26, 3-37
 natural, 3-11
 outer, 3-5–3-6, 3-14–3-18
 right, 3-15
 two-way, 3-25–3-26
 with views, 3-6

K

keywords
 ALL, 2-57, 3-49–3-51, 3-62, 3-68, 3-71
 ANY, 2-56
 AS, 2-14, 3-20–3-21
 CALCULATED, 2-28–2-29
 CONSTRAINTS, 4-18
 CORRESPONDING, 3-49–3-51, 3-72
 DISTINCT, 2-18
 INDEX, 4-36
 INTO, 5-21
 UNIQUE, 4-36
 VALIDATE, 2-7

L

LABEL= option, 2-35, 3-23
LAG function, 2-16
LENGTH statement
 DATA step, 1-12
LENGTH= option, 2-35
LIKE operator
 indexes, 4-33
loading data, 4-7–4-10

M

macro variables, 5-18–5-28
 automatic, 5-26–5-27
 SELECT statement, 5-21, 5-23, 5-25
 symbolic references, 5-18–5-19
 MERGE statement
 DATA step, 3-21–3-22
 merges
 comparing with SQL joins, 3-40
 metadata, 5-9–5-16
 modifiers
 set operators, 3-49–3-51
 multiway joins, 3-26, 3-37

N

natural joins, 3-11
 NOEXEC option, 2-8, 2-10
 NOLABEL system option, 2-36
 NOPRINT option, 5-20
 NOT EXISTS condition, 2-65
 NOT NULL integrity constraint, 4-12

O

ON clause, 3-15
 operators
 CONTAINS, 4-33
 EXCEPT, 3-46–3-47, 3-50–3-59
 INTERSECT, 3-46–3-48, 3-51, 3-60–3-65
 JOIN, 3-15
 LIKE, 4-33
 OUTER UNION, 3-46, 3-49–3-50
 UNION, 3-46–3-48, 3-51, 3-66–3-70
 options
 FEEDBACK, 2-14
 FORMAT=, 2-35
 FULLSTIMER system, 5-30
 IDXNAME=, 4-39
 IDXWHERE=, 4-39
 INFORMAT=, 2-35
 LABEL=, 2-35, 3-23
 LENGTH=, 2-35
 NOEXEC, 2-8, 2-10
 NOLABEL, 2-36
 NOPRINT, 5-20
 resetting, 5-7–5-8
 SORTSEQ=, 2-32
 SQL procedure, 5-3–5-4
 STIMER SQL, 5-30
 testing and performance, 5-30–5-31
 VALIDATE, 2-10
 ORDER BY clause, 2-31–2-33
 outer joins, 3-5–3-6, 3-14–3-18

 full joins, 3-17
 left joins, 3-15
 right joins, 3-15
 OUTER UNION operator, 3-46, 3-49–3-50

P

PRIMARY KEY integrity constraint, 4-12
 PRINT procedure
 SELECT statement, 4-10
 procedures
 APPEND, 3-74
 DATASETS, 4-11
 FORMAT, 4-50
 PRINT, 4-10
 SORT, 2-32, 3-40
 SQL, 1-4–1-5, 2-3, 3-22, 3-40, 4-11

Q

queries
 complex, 3-31–3-36
 QUERY statement, 3-20
 query-expression, 4-8–4-9

R

RESET statement
 SQL procedure, 2-10
 resetting options
 SQL procedure, 5-7–5-8
 right joins, 3-15
 ROLLBACK statement, 4-15–4-16

S

SAS data sets, 1-6
 SAS/ACCESS views, 1-6
 SASHELP library, 5-16
 SCAN function, 1-11–1-12
 SELECT clause, 5-20
 SELECT statement, 2-4–2-6, 3-29, 4-8
 PRINT procedure, 4-10
 SQL procedure, 2-10, 2-18, 5-21
 SET clause, 4-8
 set operations, 3-4
 set operators
 default behavior, 3-47
 modifiers, 3-49–3-51
 SET statement, 4-8
 DATA step, 3-72
 SORT procedure, 2-32, 3-40
 SORTSEQ= option
 SORT procedure, 2-32
 SQL. See Structured Query Language

- SQL procedure, 1-4-1-5, 2-3, 3-22, 3-40, 4-11
 - controlling processing, 5-3-5-7
 - CREATE TABLE statement, 5-20
 - CREATE VIEW statement, 5-20
 - creating integrity constraints, 4-13-4-14
 - data types, 4-4
 - indexes, 4-31
 - optimizer, 3-25-3-26
 - options, 5-3-5-4
 - re-merging, 2-49
 - resetting options, 5-7-5-8
 - SELECT statement, 2-18, 2-37
 - testing and performance options, 5-30-5-31
 - UNDO_POLICY= option, 4-16
 - views, 4-21
- statements
 - ALTER, 2-9
 - BY, 3-21-3-22
 - CREATE, 2-9
 - CREATE INDEX, 4-18, 4-36
 - CREATE TABLE, 4-3, 4-7, 4-18, 5-20
 - CREATE VIEW, 4-22, 5-20
 - DELETE, 2-9, 4-51
 - DESCRIBE, 2-9
 - DESCRIBE TABLE, 4-18, 5-13
 - DROP, 2-9, 4-56
 - FOOTNOTE, 2-37
 - INSERT, 2-10
 - MERGE, 3-21-3-22
 - QUERY, 3-20
 - RESET, 2-10
 - SELECT, 2-4-2-6, 2-10, 2-18, 2-37, 3-29, 4-8, 4-10, 5-21
 - SET, 3-72, 4-8
 - TITLE, 2-37
 - UPDATE, 2-10, 4-46-4-47
- STIMER SQL option, 5-30
- Structured Query Language (SQL), 1-3-1-4, 2-3
 - comparison with traditional SAS programs, 3-38-3-39
- subqueries, 2-51-2-55
 - correlated, 2-59-2-71
- subsetting IF statement, 4-52
- subsetting with the calculated values, 2-28-2-29
- subsetting with the WHERE clause, 2-19-2-24, 2-26
- SUBSTR function, 1-9-1-11, 4-49
 - indexes, 4-33
- summary functions, 2-42-2-45

- symbolic references, 5-18-5-19

T

tables

- altering columns, 4-53-4-55
- conditional processing, 4-48-4-50
- deleting, 4-56
- deleting rows, 4-51-4-52
- testing and performance options, 5-30-5-31
- THEN expression, 4-48
- TITLE statement, 2-37
- TRIM function, 1-11
 - indexes, 4-33
- two-way joins, 3-25-3-26

U

- UNDO_POLICY= option, 4-16
- UNION, 3-51
- UNION operator, 3-46-3-48, 3-66-3-70
- UNIQUE integrity constraint, 4-12
- UNIQUE keyword, 4-36
- UPDATE statement
 - PROC SQL, 4-46-4-47
 - SQL procedure, 2-10
- using a view, 4-24

V

- VALIDATE keyword, 2-7
- VALIDATE option, 2-10
- VALUES clause, 4-8-4-9
- views
 - administering, 4-25-4-30
 - creating, 4-21-4-23, 4-29
 - deleting, 4-56
 - in-line, 3-27-3-30, 3-35
 - PROC SQL, 4-21-4-23
 - SAS/ACCESS, 1-6
 - using, 4-24
- virtual tables, 4-21

W

- WHEN condition, 4-49
- WHERE clause, 2-19-2-24, 2-26, 2-50, 3-9, 3-15, 3-25, 4-51
 - indexes, 4-33